



The Scenario: The Peacetopia Bird-Watch System

The Mission: The people of Peacetopia are terrified of birds. You are the project lead building a smart camera system to sound an alarm whenever a bird enters the city.

The Resources: You have 10 million photos to learn from. Some have birds, some don't.

Challenge 1: The Performance Check

Goal: Identify if the "brain" of our system is working correctly.

You built an AI model to achieve the above goal. Before you launch this model, you check how many mistakes it makes (Error Rate). Here is what you find:

- **Training Set Error (The "Practice Exam"):** 15% Error
- **Validation Set Error (The "Real Exam"):** 16% Error

The Questions for the Audience:

1. **Diagnosis:** Is our system suffering from **High Bias** (it hasn't learned the patterns well enough) or **High Variance** (it's just memorizing and can't handle new data)?
2. **The Solution:** If you were the boss, what is the **one thing** you would tell your engineers to do to fix this?
 - A) Go get more data to train on.
 - B) Build a bigger, more complex AI model.

The "Cheat Sheet for the Answer:

- **Answer 1: High Bias.** Because it can't even get the "Practice Exam" right. It's failing to learn the basics.
- **Answer 2: B.** Since it's failing to learn, it needs a "bigger brain" (more layers/complexity) or better features. Adding more data won't help if the model isn't smart enough to understand the data it already has!

Challenge 2: Dividing the Map

Goal: Deciding how to use our 10,000,000 images.

In the "old days" of small data, people used to split data **70% for Training and 30% for Testing**. But you have **10 million images**. **The Questions for the Audience:**

1. **The Decision:** Would you use the old 70/30 split? Or would you give the AI **99%** for training and keep only **1%** for the test?
2. **The Justification:** Why is 1% (100,000 images) enough for the final test?

The Cheat Sheet for the Answer:

- **Answer 1:** Use the **99% / 1% split**.
- **Answer 2:** When you have "Big Data," 1% of 10 million is still **100,000 images**. In the real world, if you pass a test with 100,000 questions, we can be very confident you know your stuff! We want to give the AI as much data as possible to learn from (9.9 million images) because that's where the "learning" happens. Keeping 3 million images (30%) just for testing is a "waste" of data.

Answer:
