



**PIF**  
صندوق الاستثمارات العامة  
Public Investment Fund

أكاديمية كاوت  
KAUST ACADEMY



جامعة الملك عبد الله  
للتكنولوجيا  
King Abdullah University of  
Science and Technology

# LEVEL 2: Introduction to Computer Vision and Convolutional Neural Networks

---

## Day 4

# Course Outline

---

- Computer Vision
- Computer Vision Tasks
- Image representation
- Convolutional Neural Networks
- Padding
- Pooling Layer
- Activation
- SOTA CNN Architectures
- Data Handling

# Learning Objectives

---

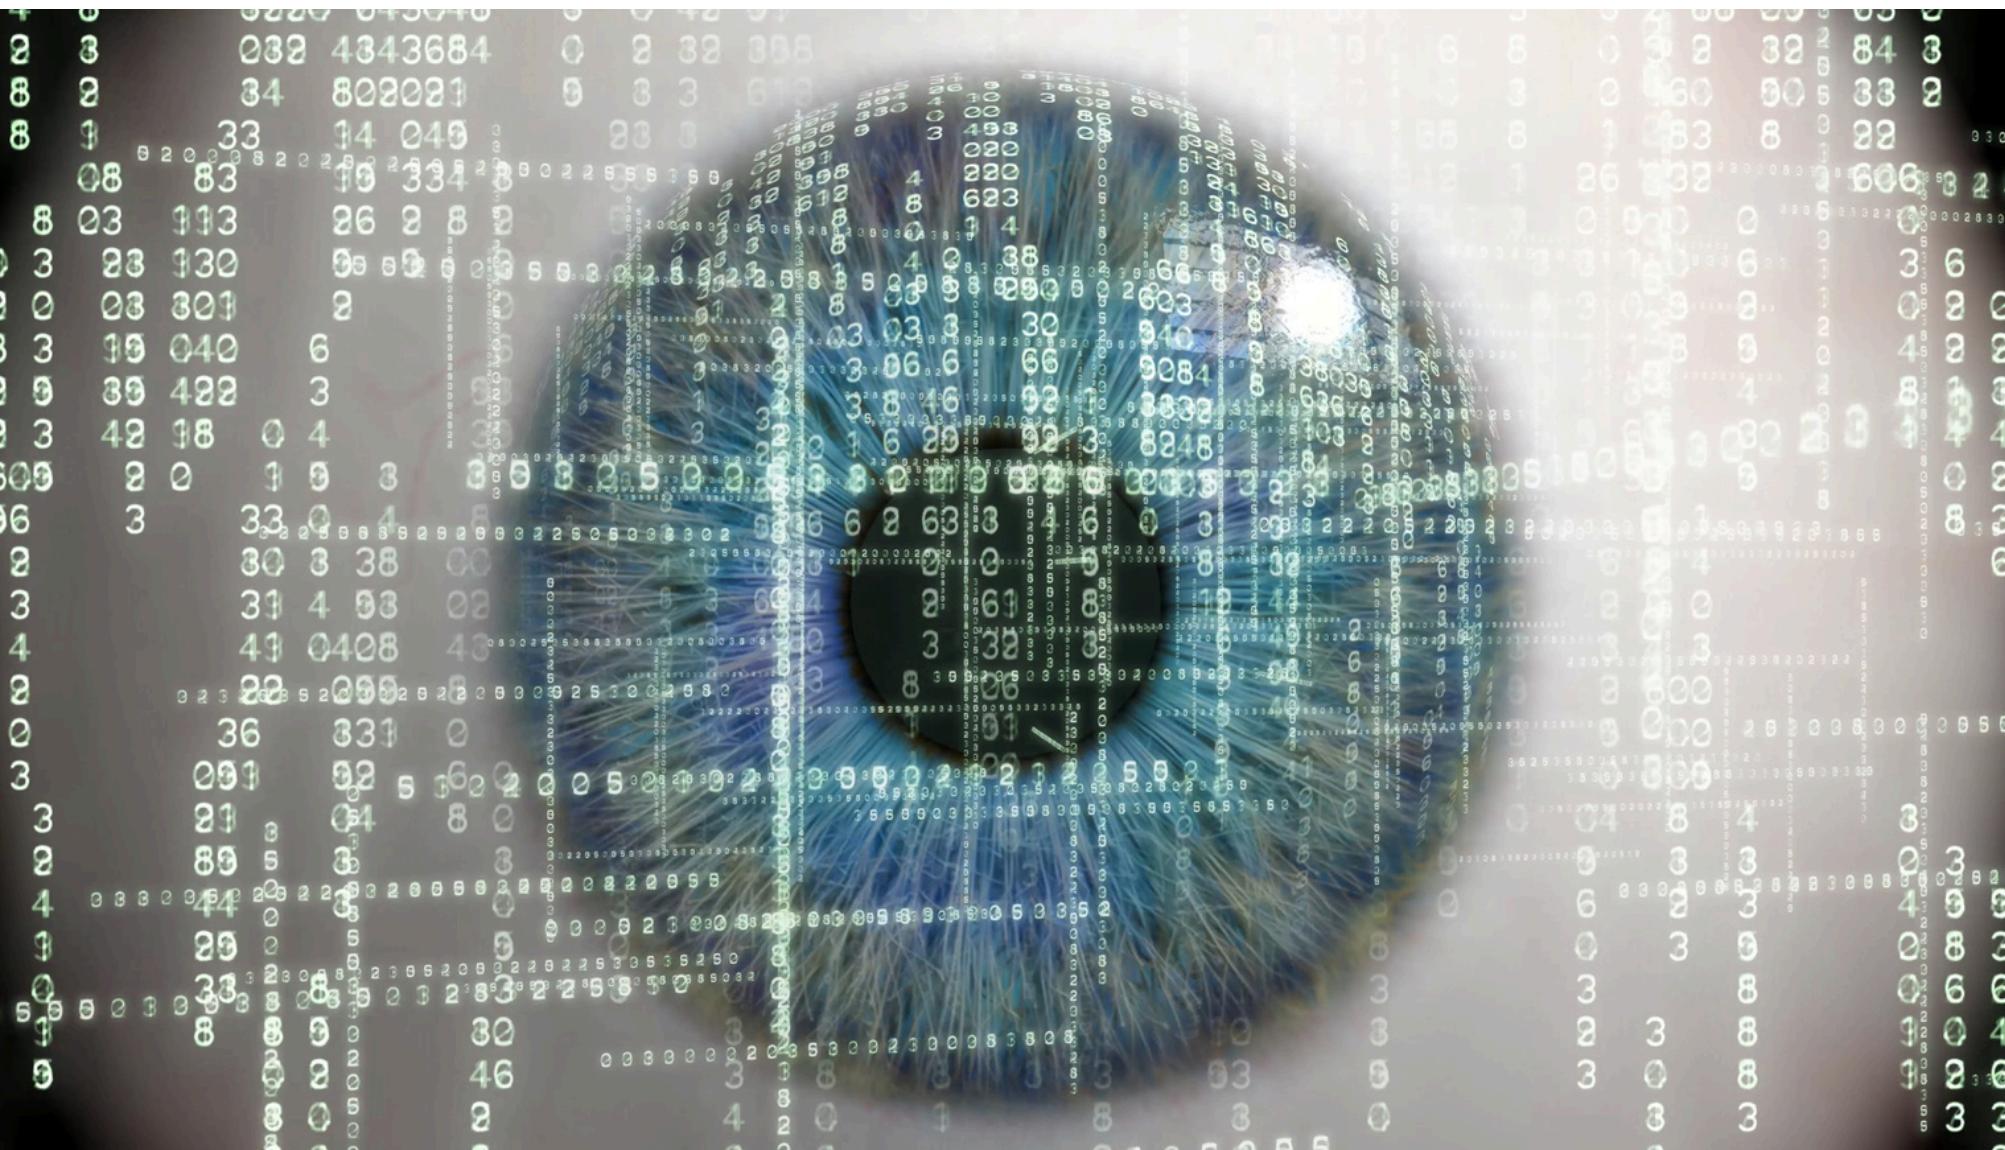
- Explain Computer Vision tasks such as image classification, object detection, and segmentation.
- Represent images numerically, including pixel values, channels, and feature maps.
- Describe and apply Convolutional Neural Networks (CNNs) for visual data processing.
- Analyze the role of padding and its impact on spatial dimensions and feature extraction.
- Differentiate pooling layers (max, average, global pooling) and explain their purpose.
- Select and apply activation functions and understand their effect on model learning.
- Compare SOTA CNN architectures (e.g., VGG, ResNet, EfficientNet) in terms of depth, efficiency, and performance.

# Computer Vision

# Computer Vision

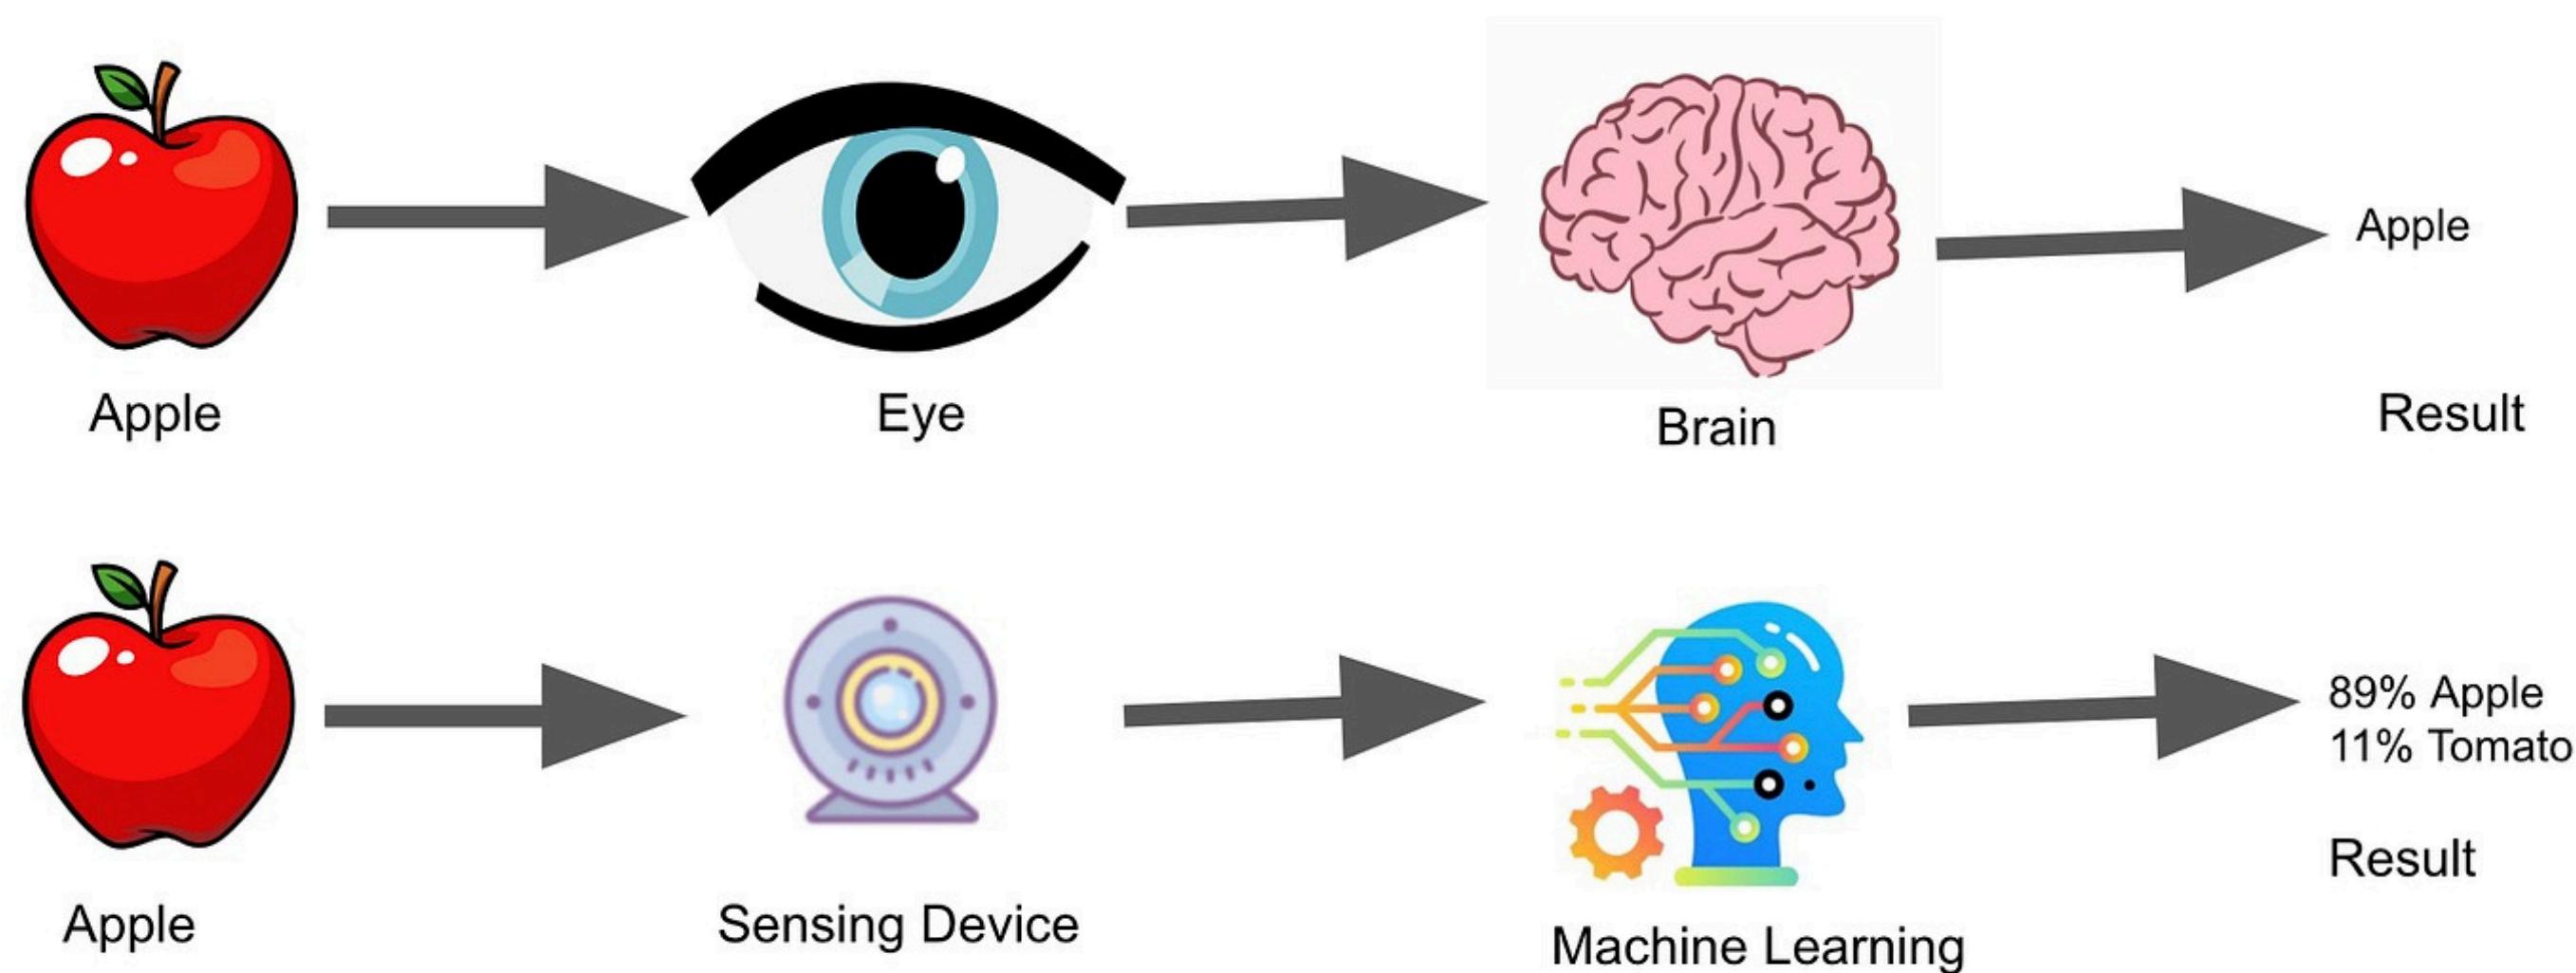
---

- Computer Vision (CV) enables machines to see, understand, and interpret the visual world.



# Computer Vision

- Computer Vision (CV) is inspired by how humans perceive and process visual information.



\*

# Computer Vision Tasks

# Computer Vision Tasks

---

## › Image Classification



*The class in this case is a cat*

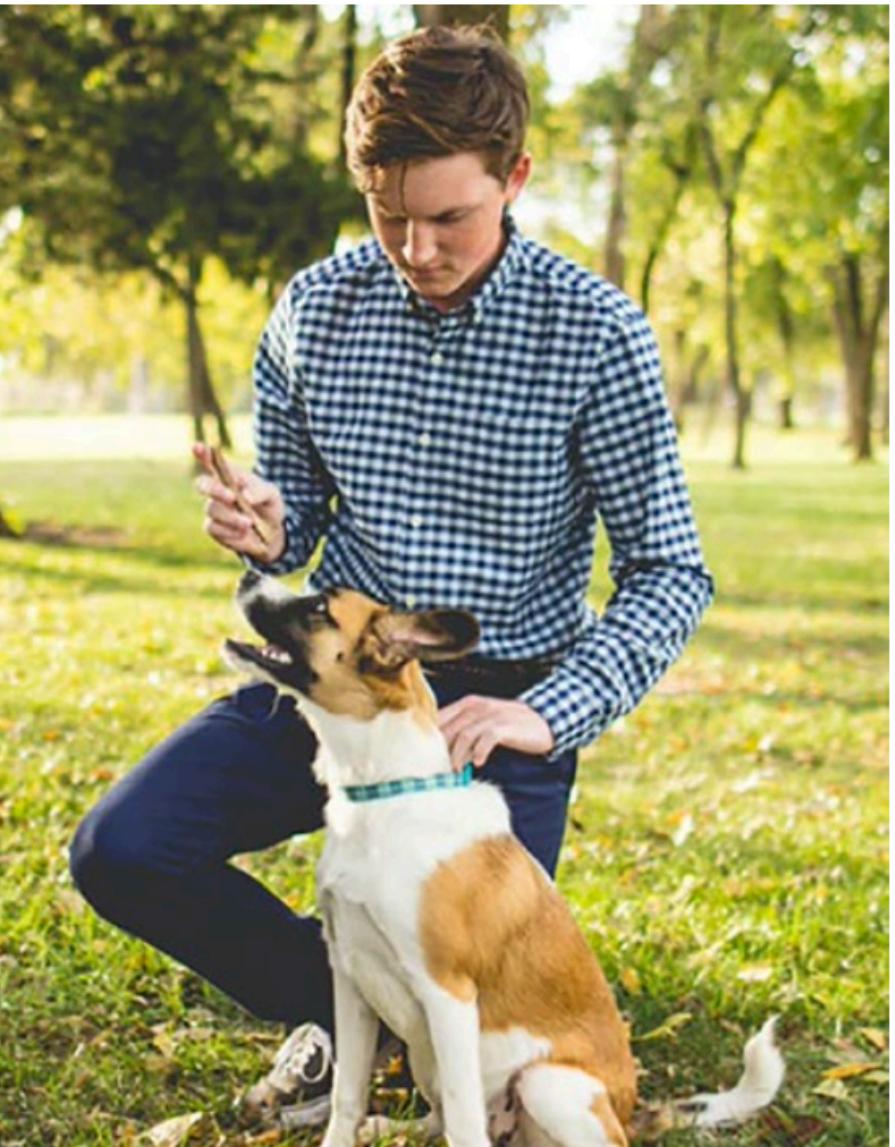


*The class in this case is a dog*

# Computer Vision Tasks

---

## ➤ Object Detection and Localization



*Class: No mask detected*



*Class: A mask is detected in the image*

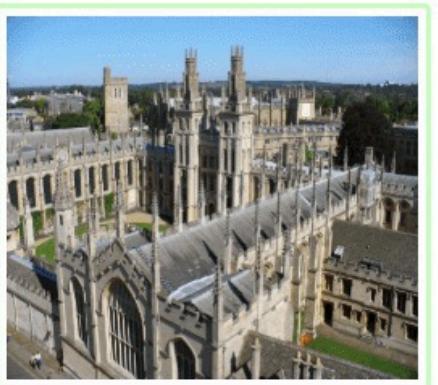
# Computer Vision Tasks

## ➤ Image Retrieval

Original Image



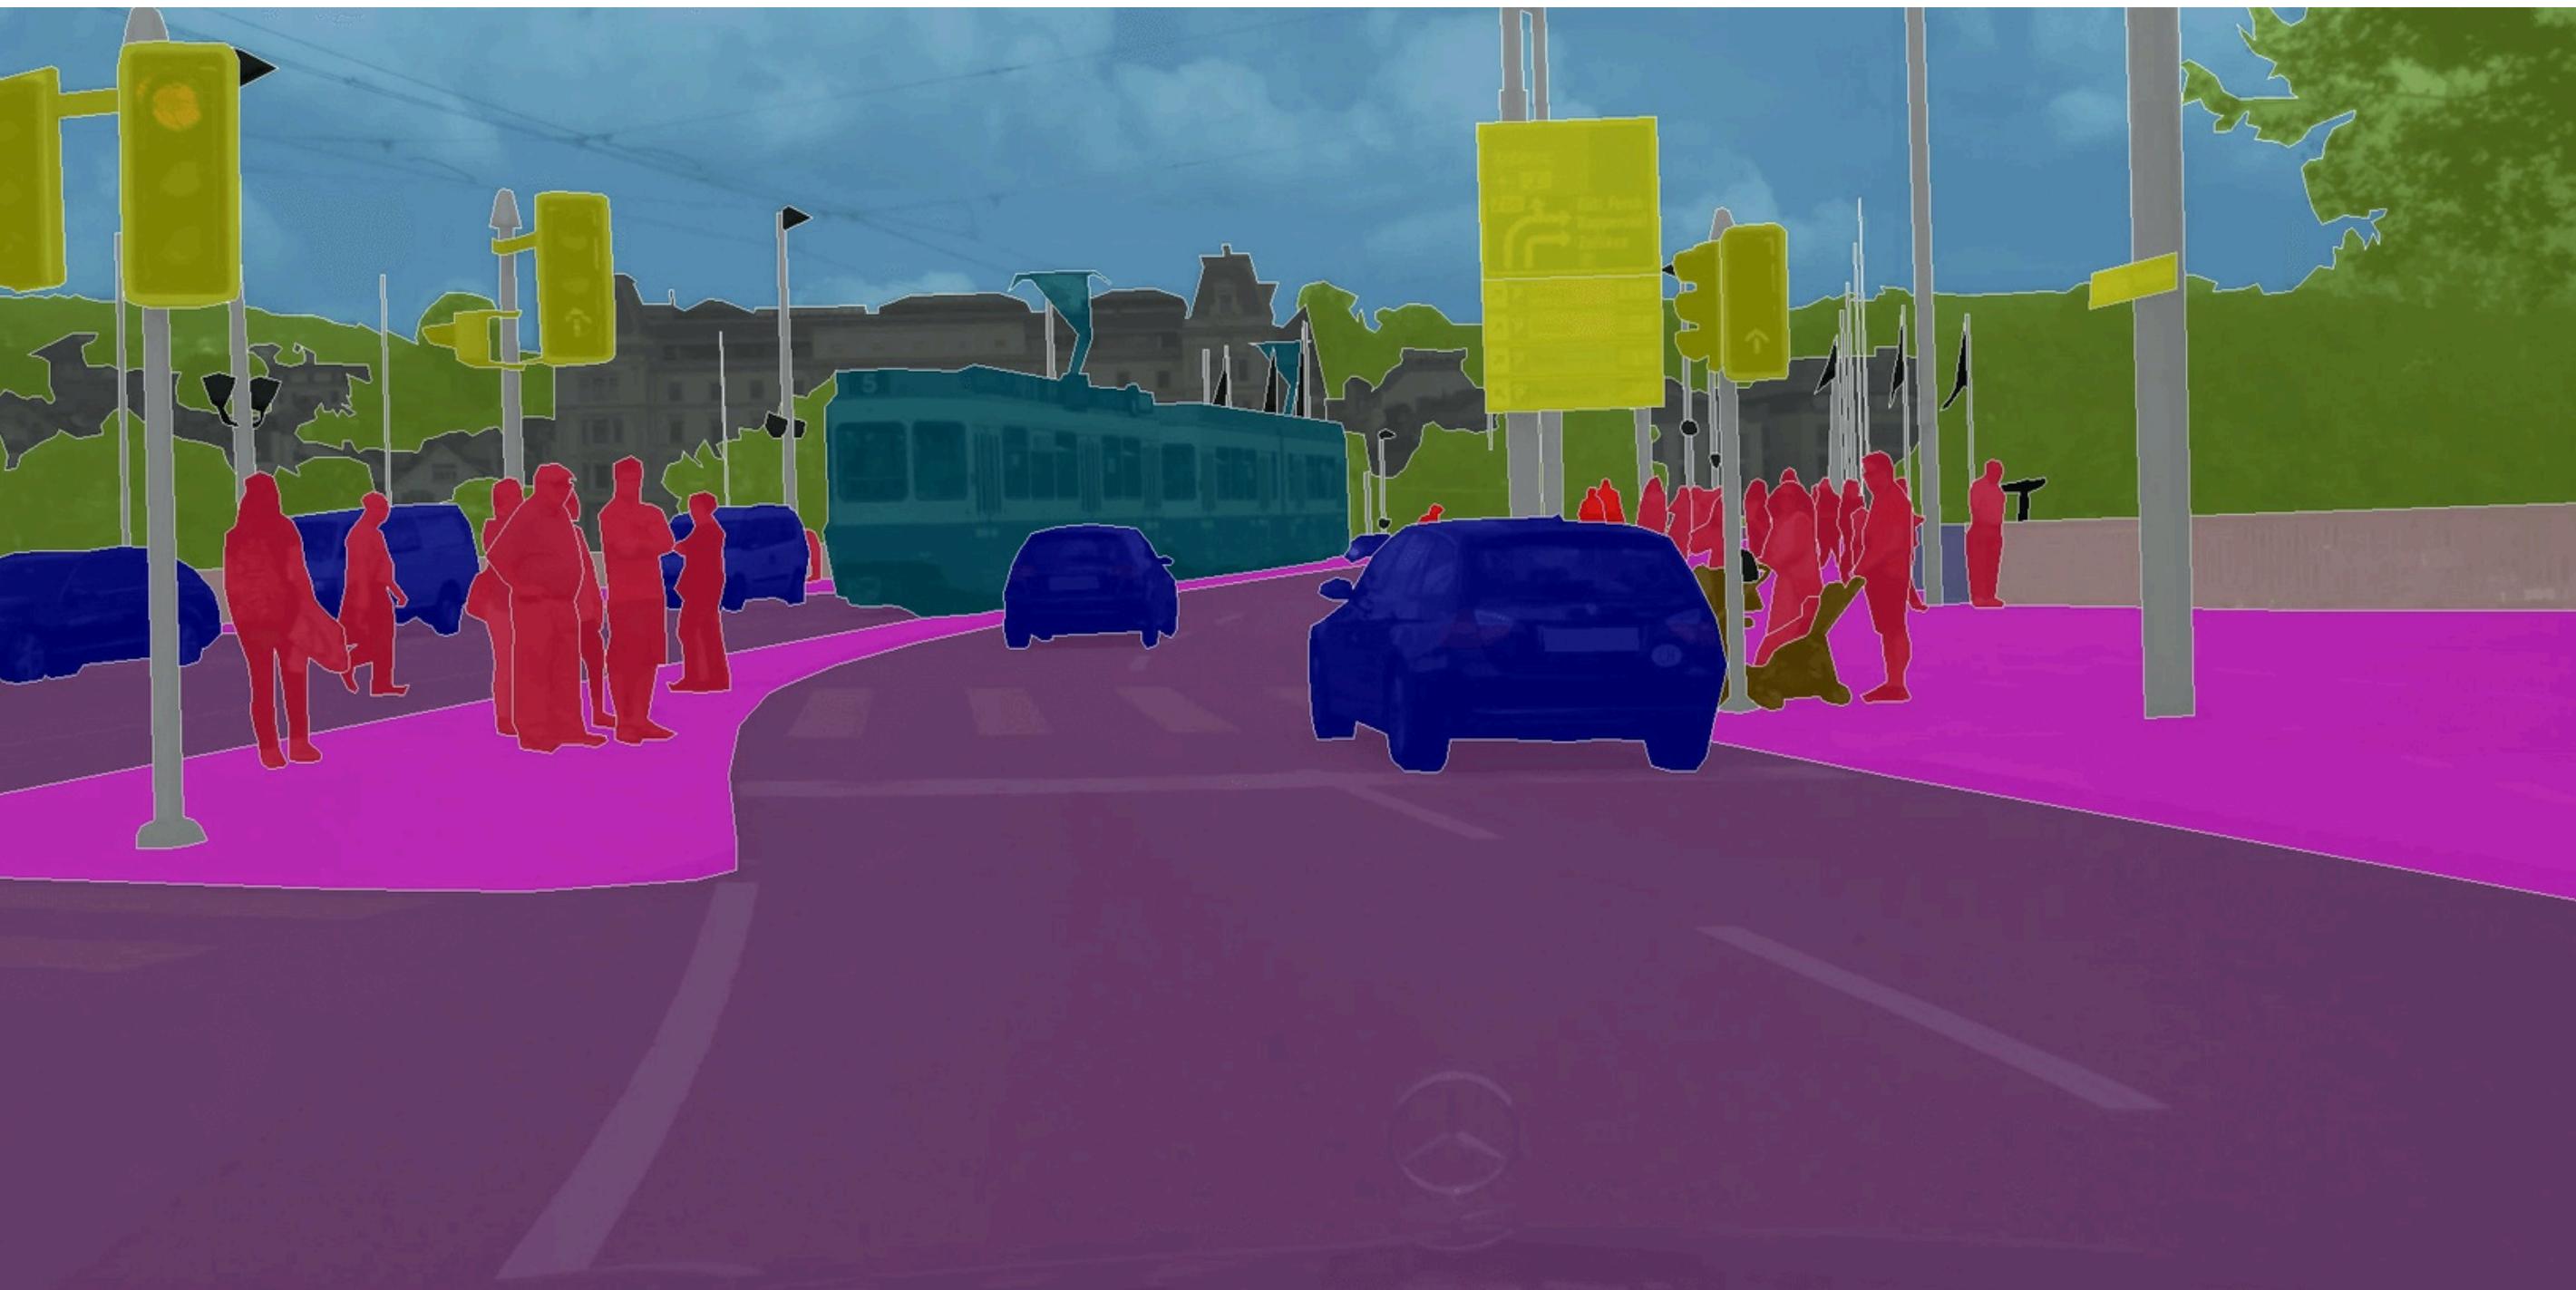
Visually Similar Images



# Computer Vision Tasks

---

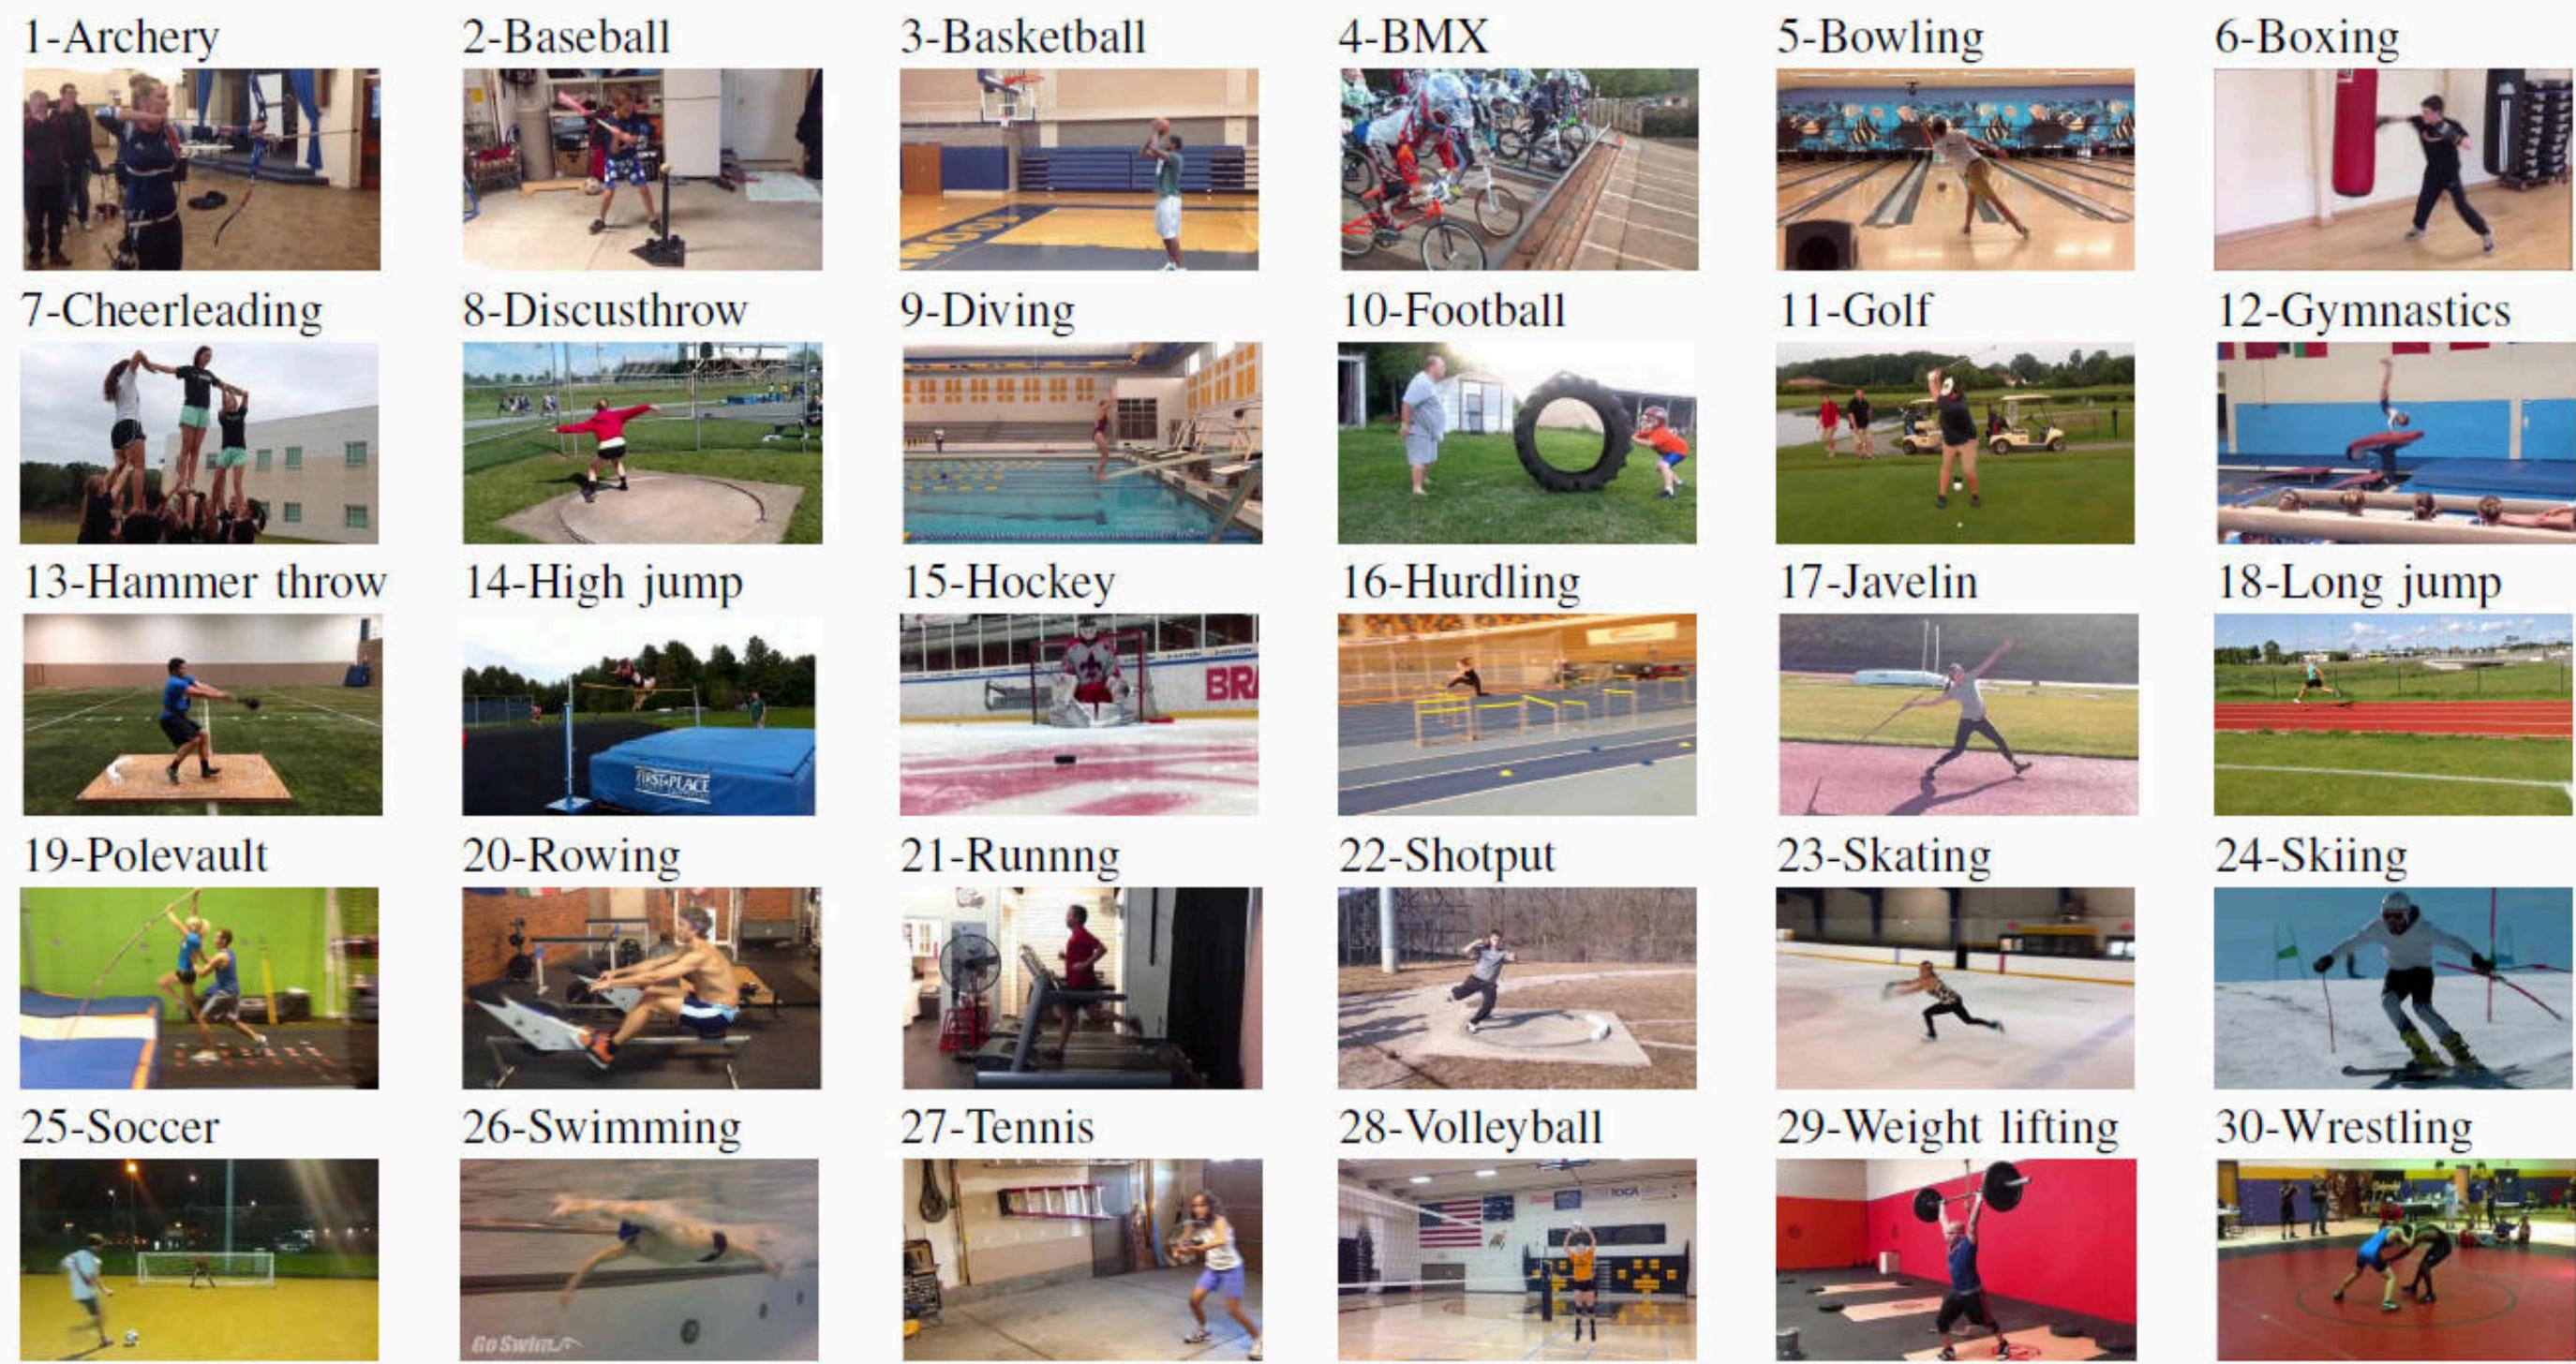
## › Image Segmentation



\*

# Computer Vision Tasks

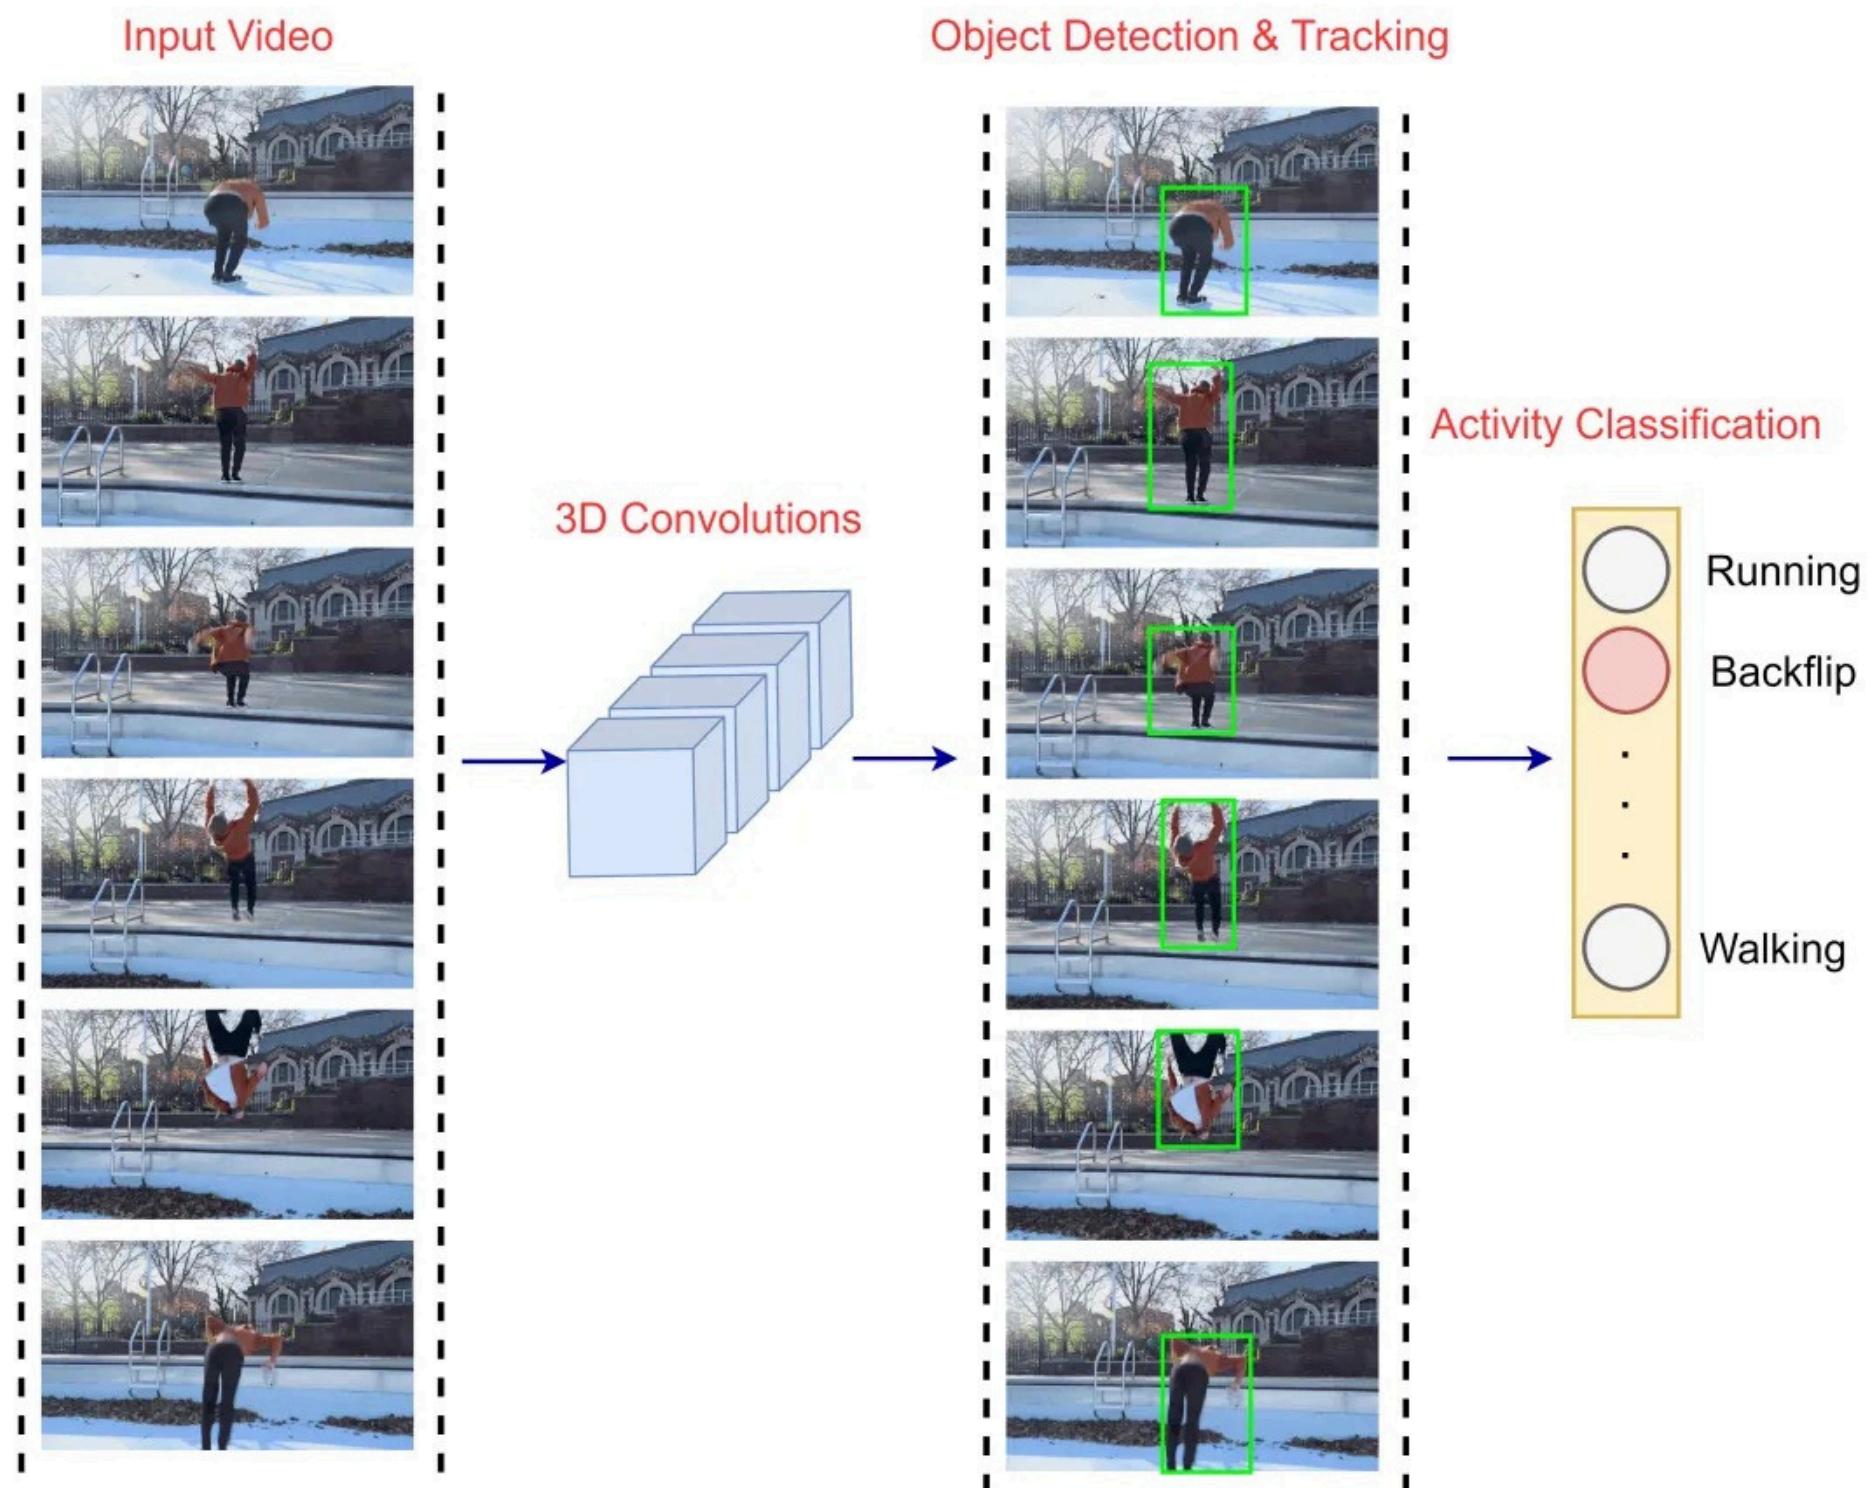
## ➤ Video Classification



\*

# Computer Vision Tasks

## > Activity Recognition



# Computer Vision Tasks

---

## > Medical Imaging



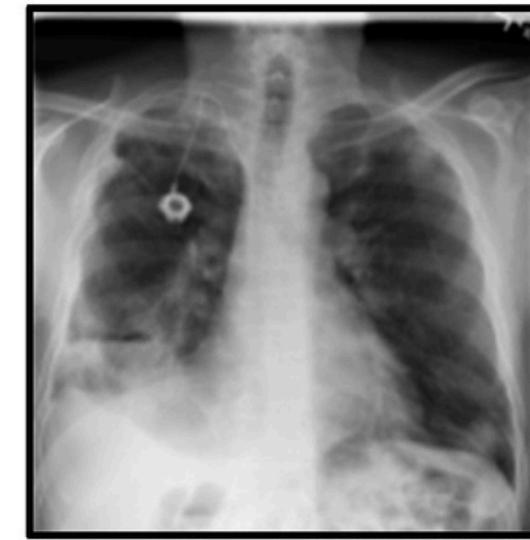
(a) Atelectasis



(b) Effusion



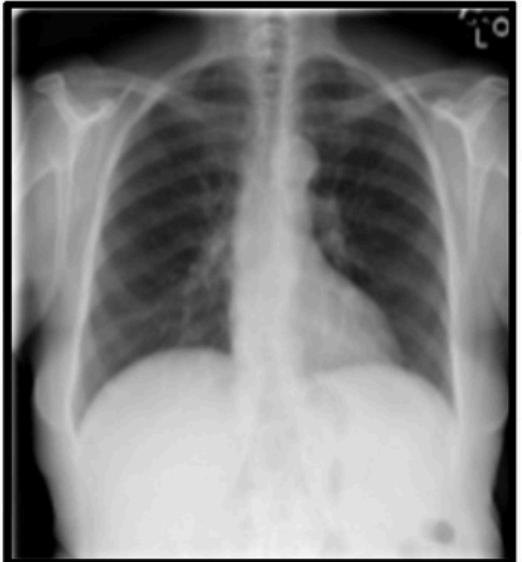
(c) Infiltration



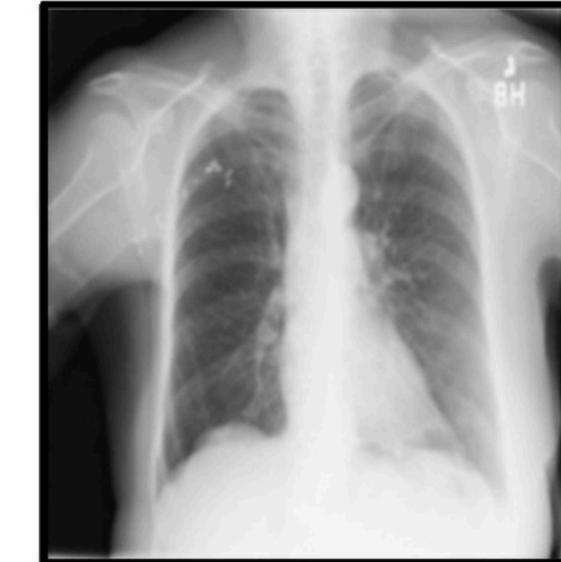
(d) Mass



(e) Nodule



(f) Normal



(g) Others



(h) Pneumothorax

# Computer Vision Tasks

---

## › Image Captioning

A young boy is playing basketball.



Two dogs play in the grass.



A dog swims in the water.



A little girl in a pink shirt is swinging.



# Computer Vision Tasks

---

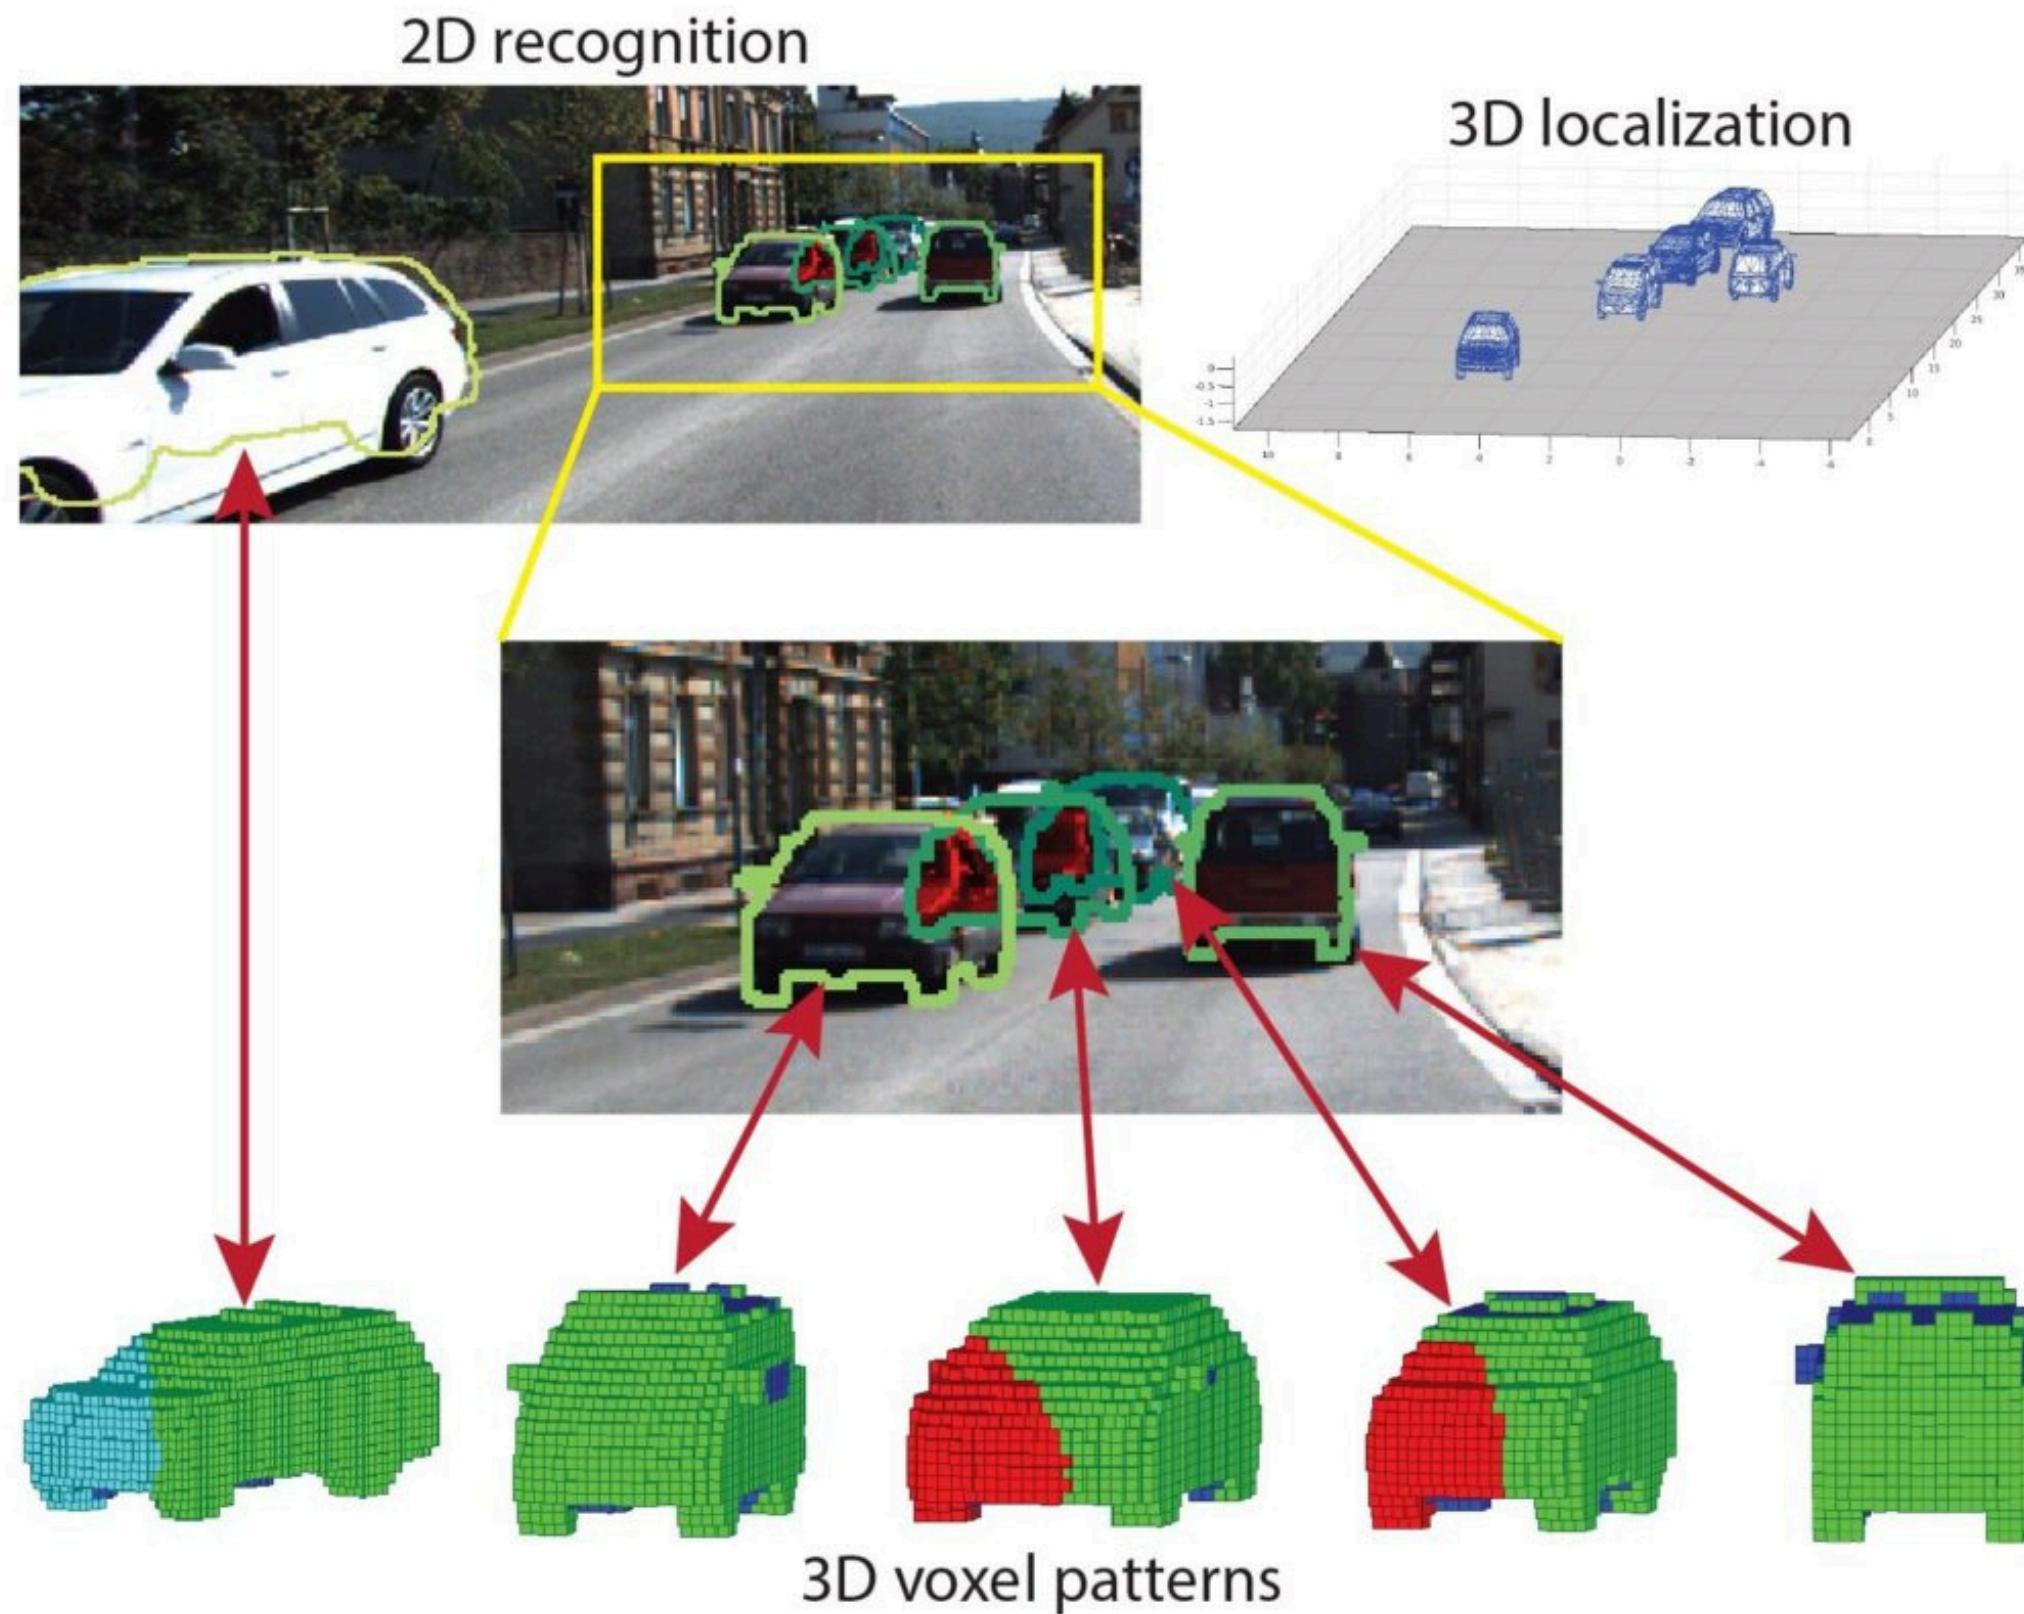
## › Image Captioning



An image generated by DALL-E 2, from the prompt "Teddy bears working on new AI research underwater with 1990s technology"

# Computer Vision Tasks

## > 3D Vision



\*

# Image Representation

# Image Representation

- Black\white
- Grey images



197	153	174	188	160	162	128	151	172	161	158	156
188	182	163	74	75	62	93	17	110	210	180	154
180	180	50	14	34	6	10	33	48	126	199	181
206	188	5	134	101	131	120	204	166	15	56	180
194	16	187	251	237	239	239	228	227	67	71	201
172	136	207	233	233	214	220	239	228	98	74	206
188	16	179	209	188	215	211	168	138	78	20	169
189	87	168	84	10	168	138	11	31	62	22	148
199	168	191	193	158	227	178	149	182	98	36	180
205	174	158	252	236	231	149	178	238	43	95	234
190	216	116	149	236	187	85	150	79	38	218	241
190	224	147	108	227	210	127	198	36	149	255	224
190	214	173	66	108	143	96	90	2	108	249	216
187	196	238	75	1	81	47	0	6	217	258	211
183	202	237	148	0	0	12	136	200	138	243	236
195	206	129	297	177	131	133	200	178	11	96	218

157	153	174	168	150	152	129	151	172	161	155	156
195	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	189	181
206	108	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	106	207	233	233	214	220	239	228	98	74	206
188	88	179	209	188	215	211	168	138	78	20	169
189	97	165	84	10	168	138	11	31	62	22	148
199	168	191	193	158	227	178	149	182	106	36	190
205	174	156	252	236	231	149	178	229	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	198	36	103	255	224
190	214	173	66	108	143	96	90	2	109	249	216
187	196	238	75	1	81	47	0	6	217	255	211
183	202	237	148	0	0	12	136	200	138	243	236
195	206	129	297	177	131	133	200	178	11	96	218

\*

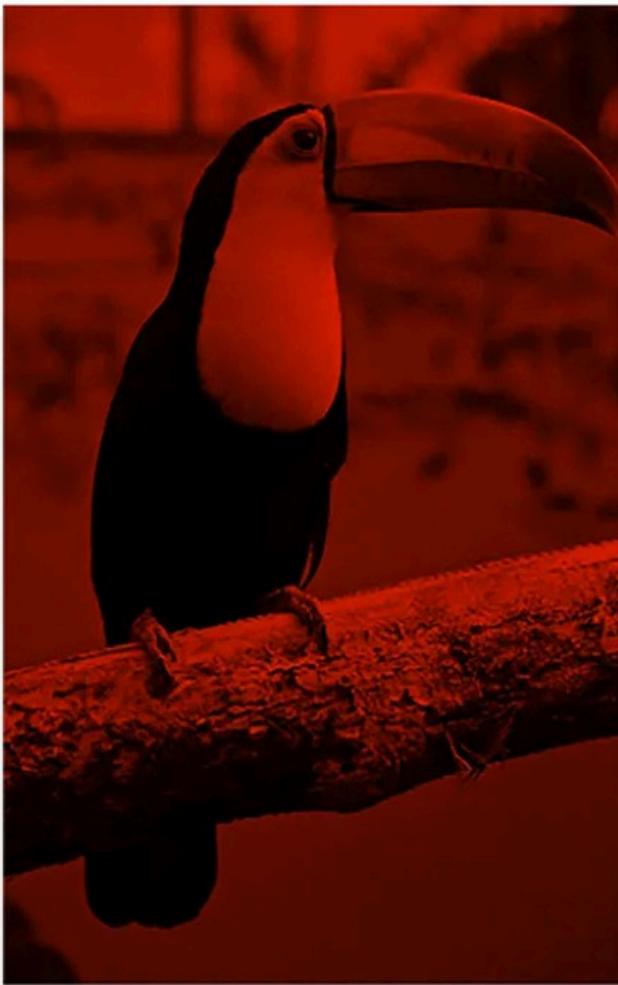
# Image Representation

---

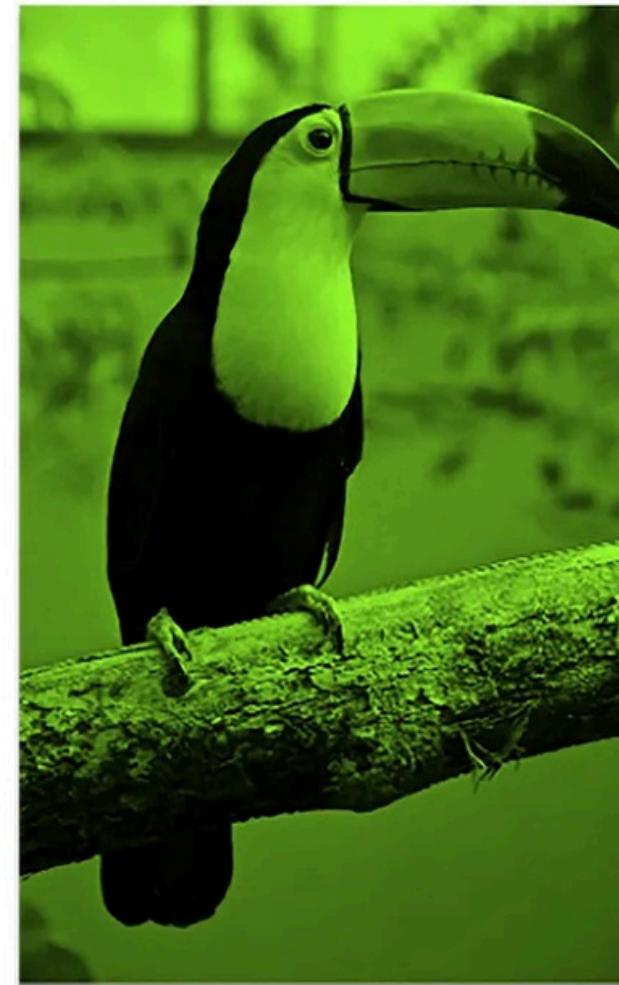
## ➤ Colored Images



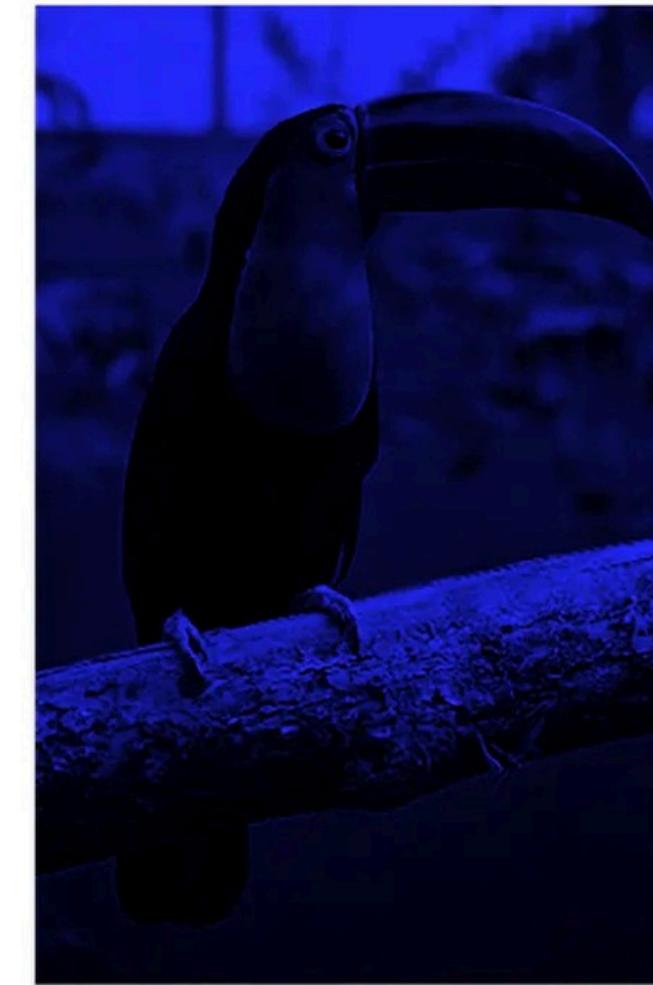
Original Image



Red Channel



Green Channel

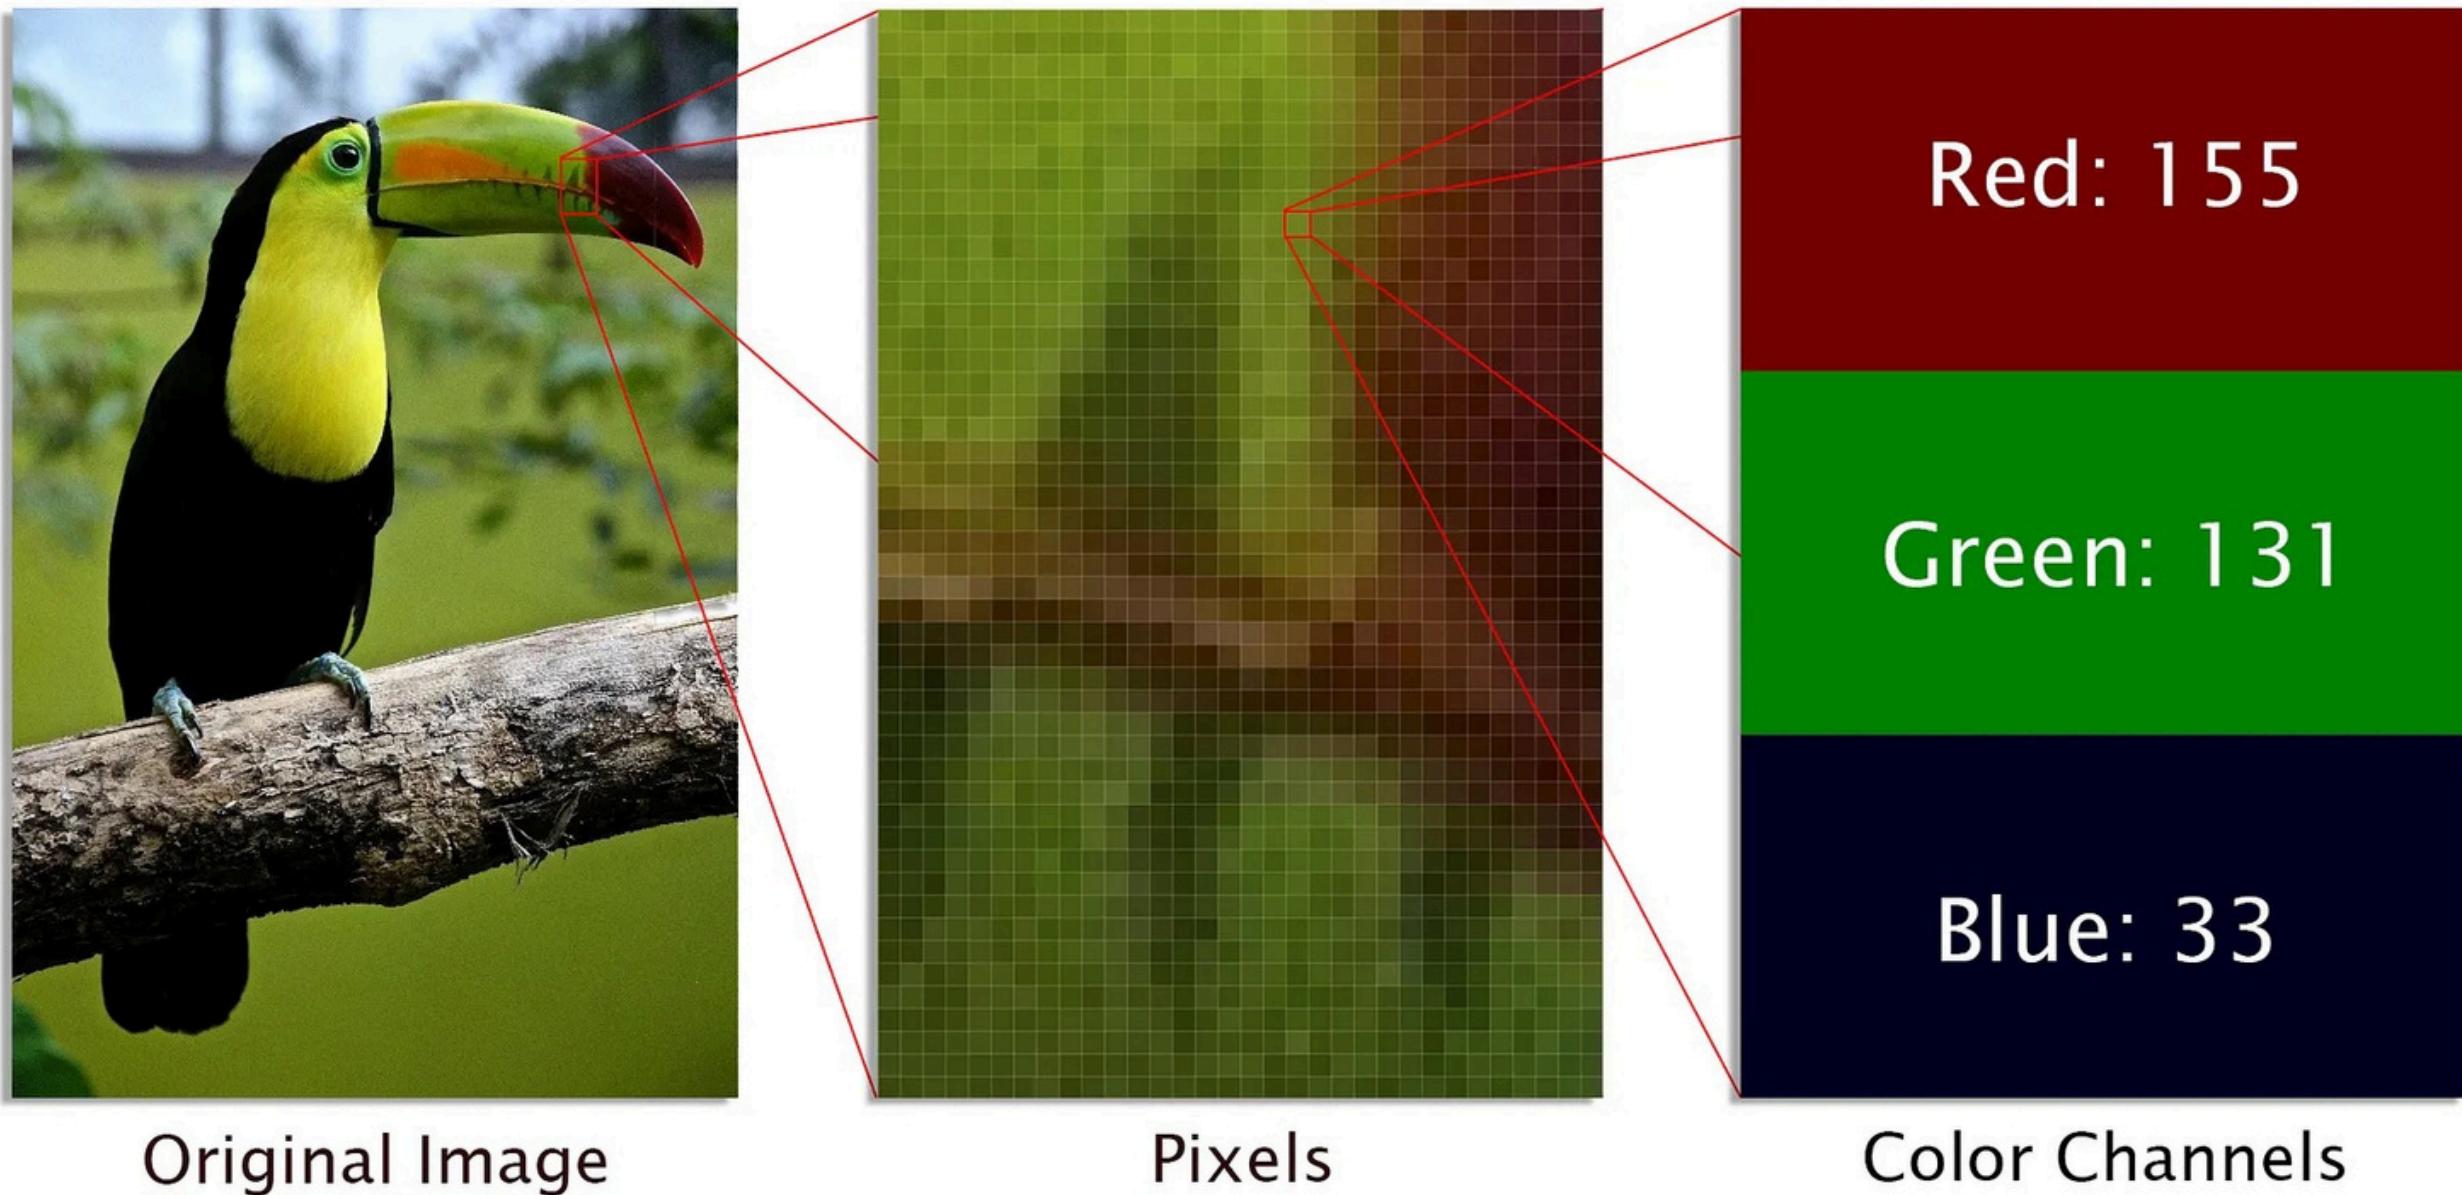


Blue Channel



# Image Representation

## ➤ Colored Images



\*

# Image Representation

**Tabular** = structured rows & columns

Size (m <sup>2</sup> )	Rooms	Type	Parking	Price (SAR)
120	3	Apartment	1	950,000
180	4	Villa	2	1,350,000
140	3	Apartment	1	980,000
220	5	Villa	3	1,850,000
160	4	Townhouse	2	1,200,000

**Image** = unstructured pixel grid

**JPG 260 X 194**



**260 X 194 X 3**

8,11,0, 55,13,25,19  
15,241,2,155,13,35,65  
14,211,0,255,23,45,11  
05,255,1,255,10,17,23  
77,167,9,112,56,16,90  
45,245,0,145,22,55,48

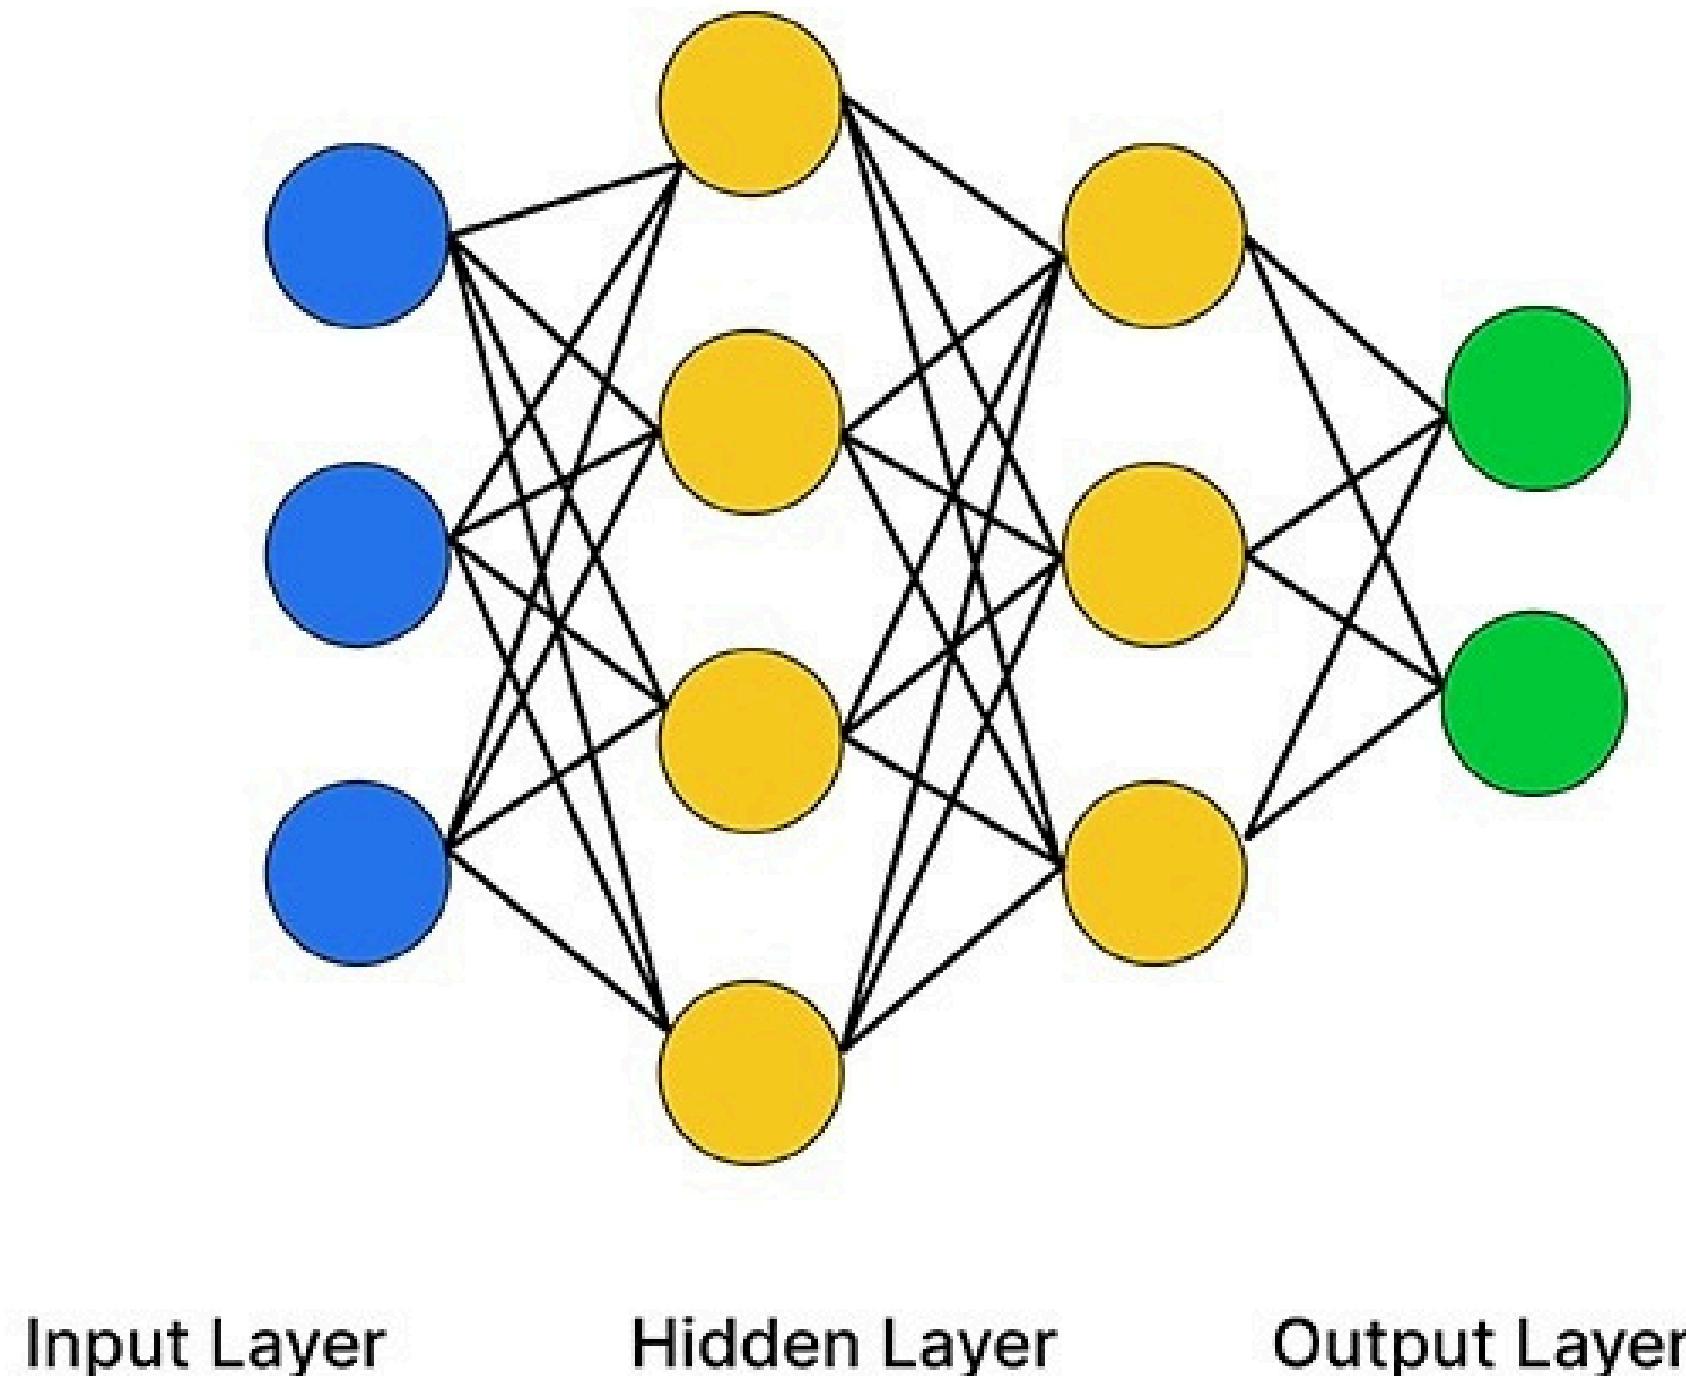


# Convolutional Neural Networks

# Convolutional Neural Networks

---

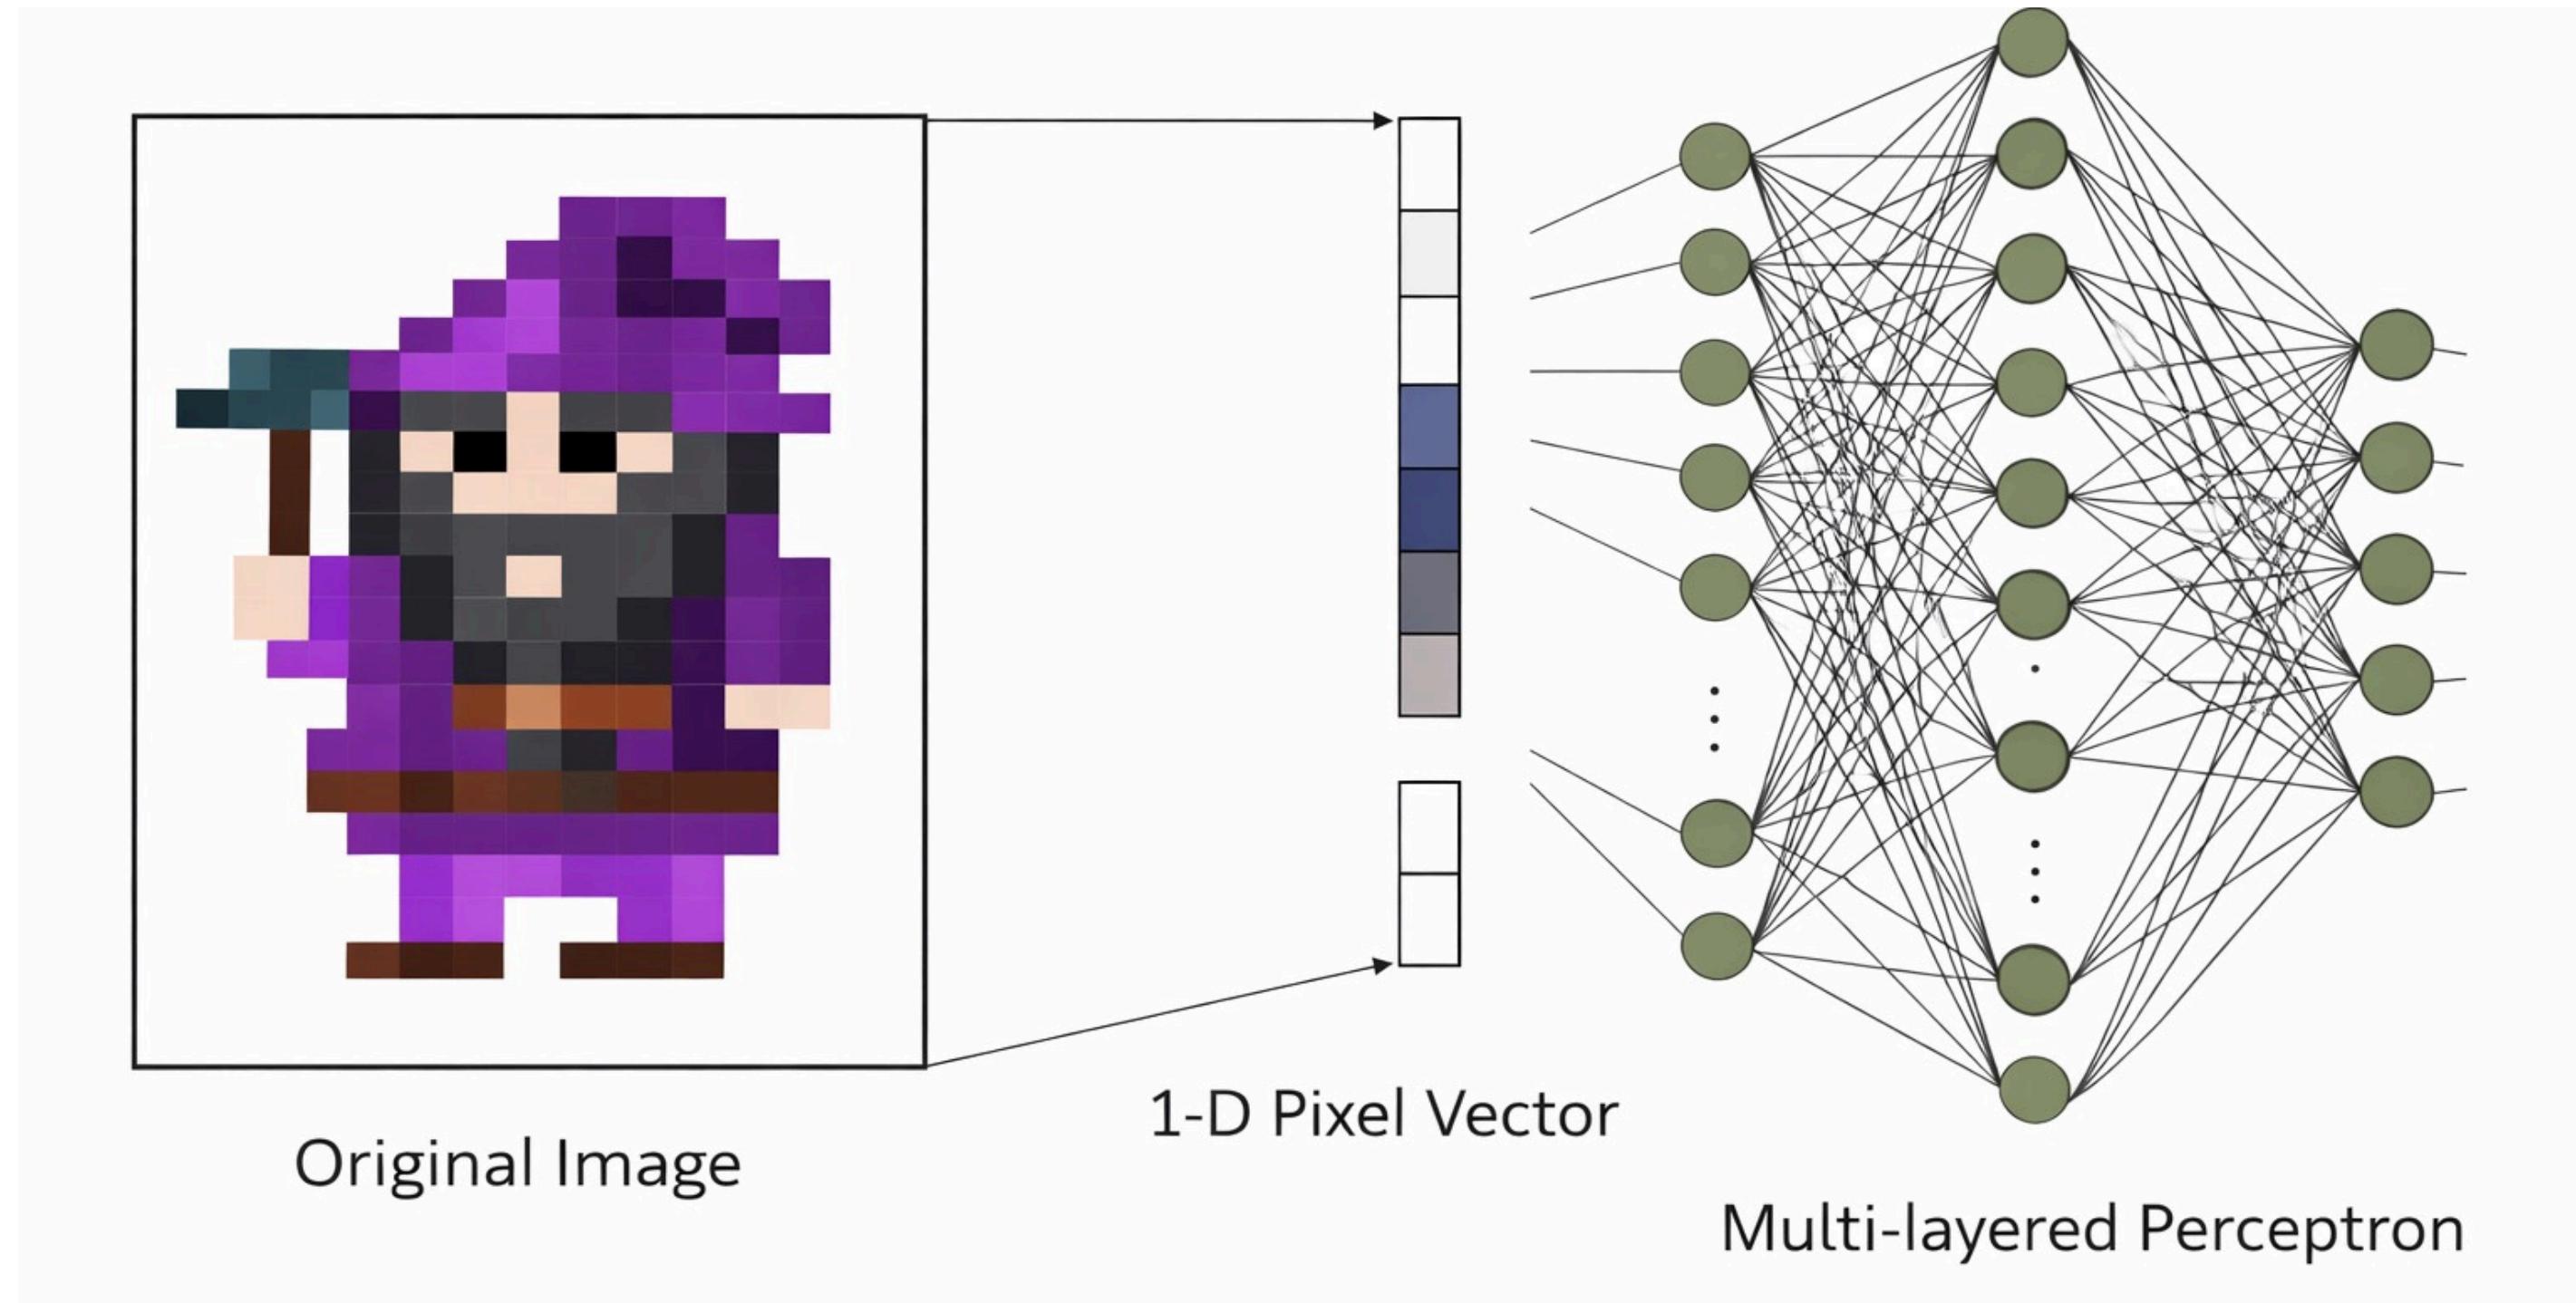
- How to build vision models?



- But can we apply the same approach to images?

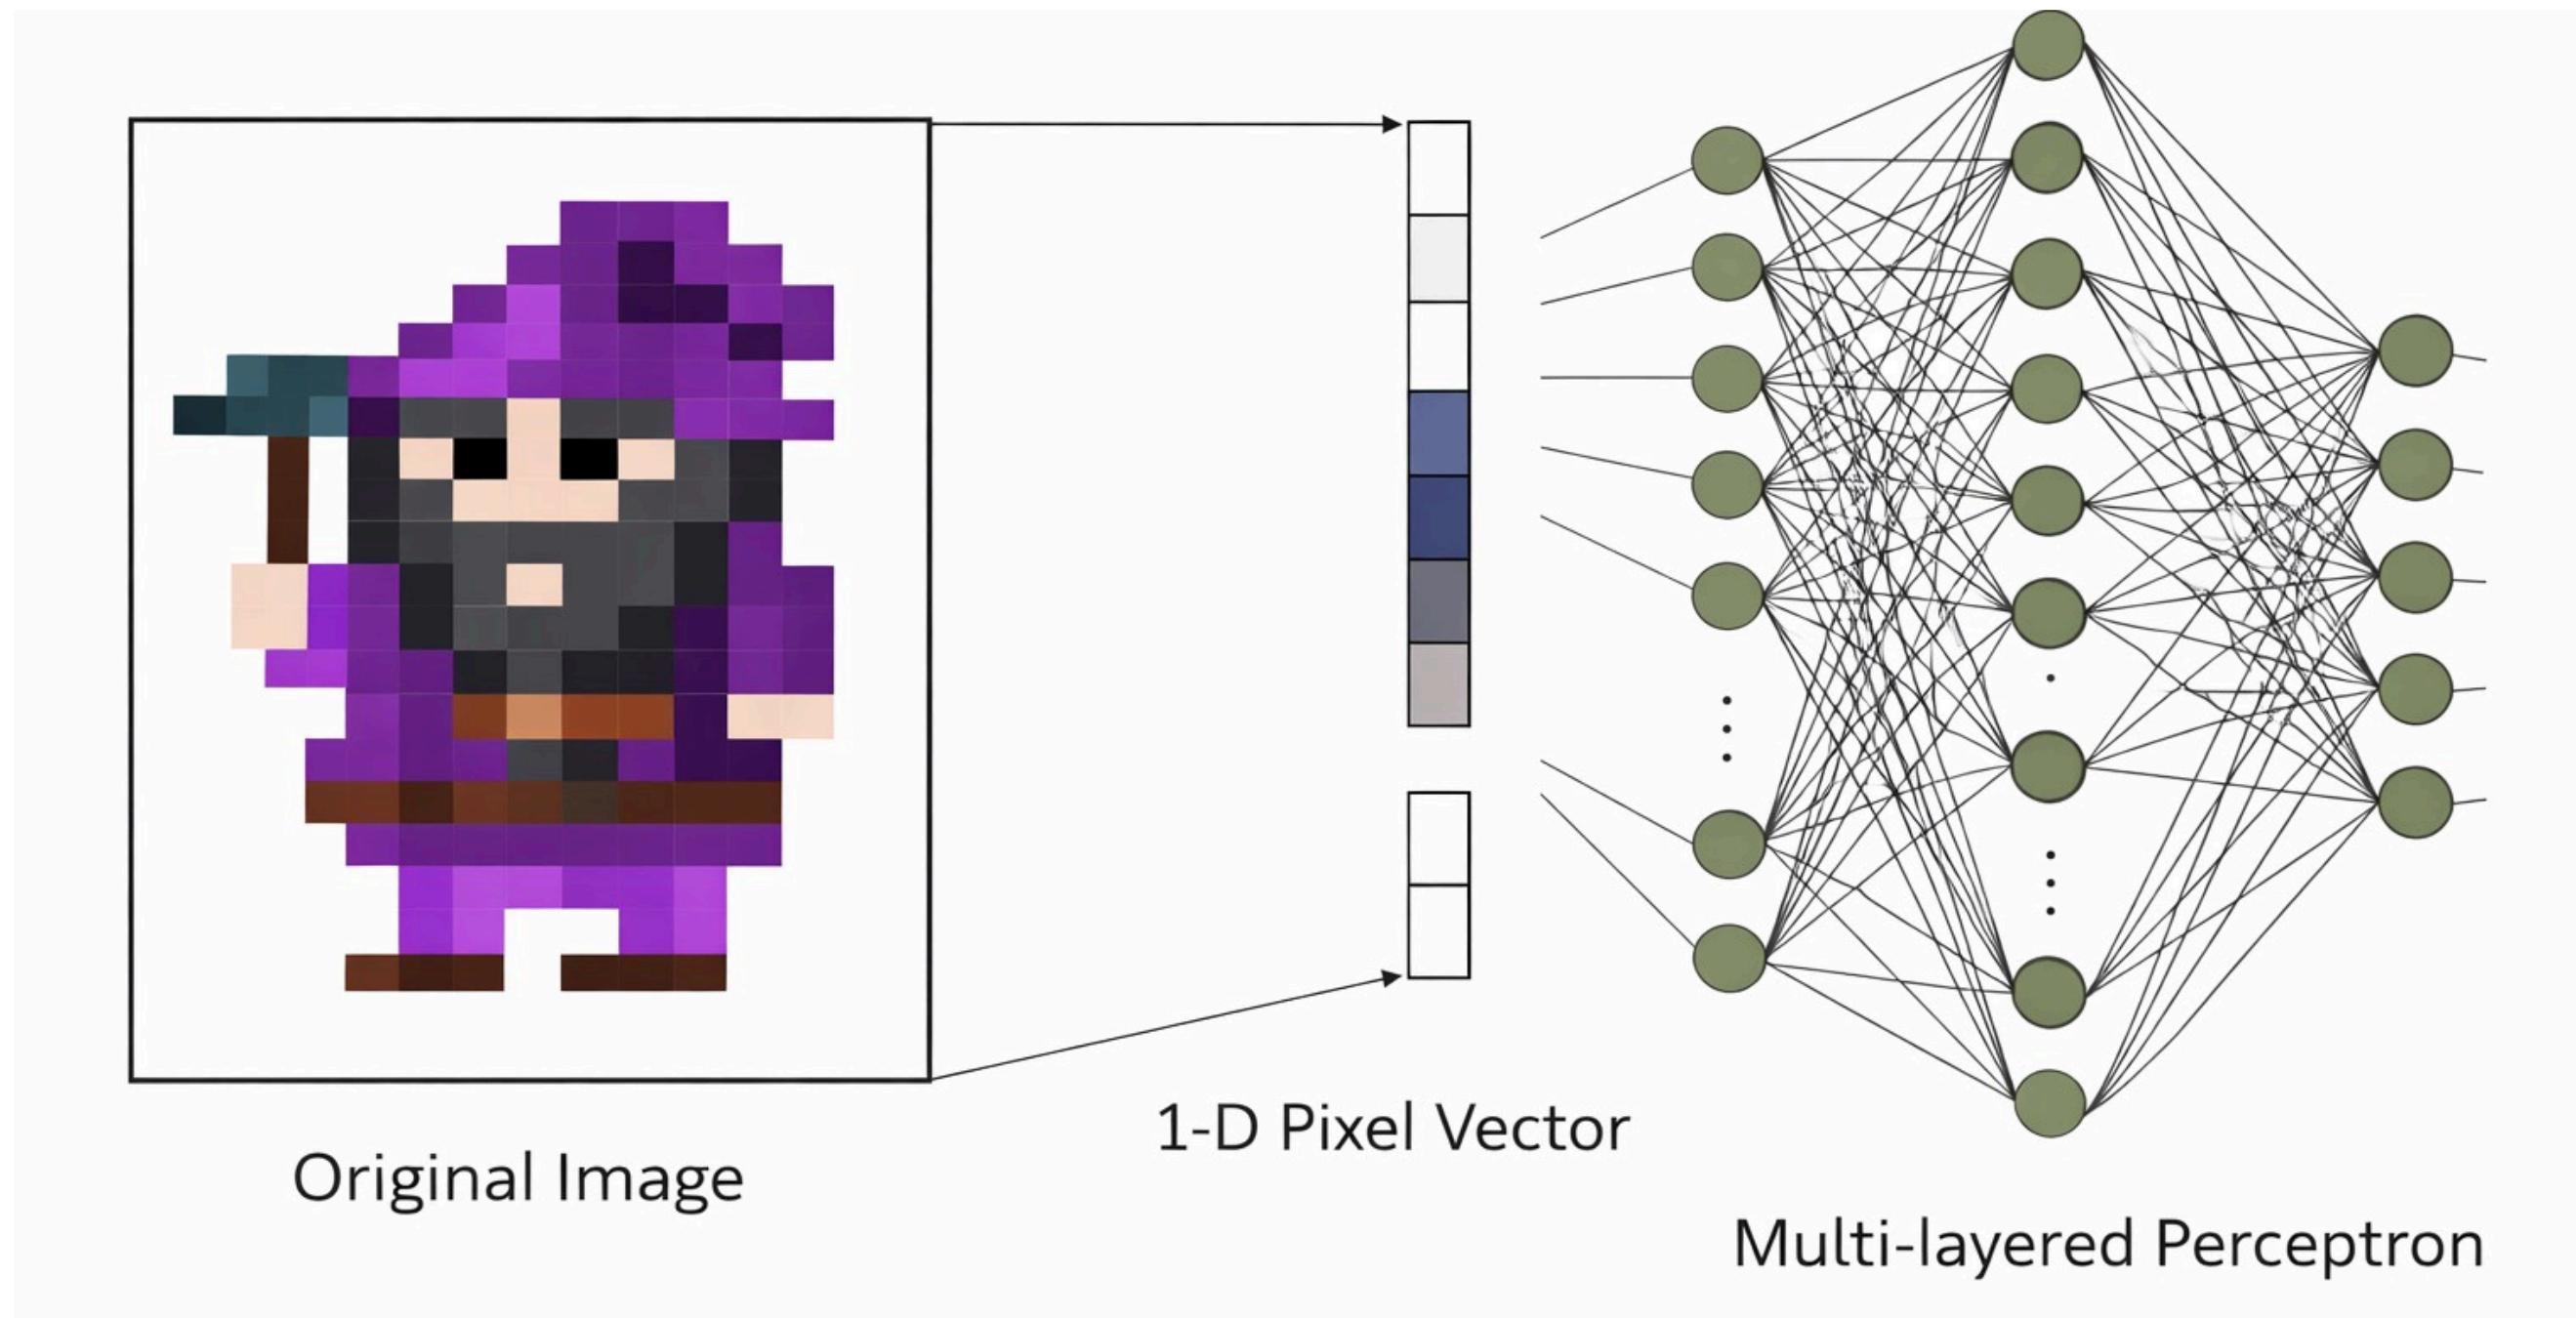
# Convolutional Neural Networks

› How to build vision models?



# Convolutional Neural Networks

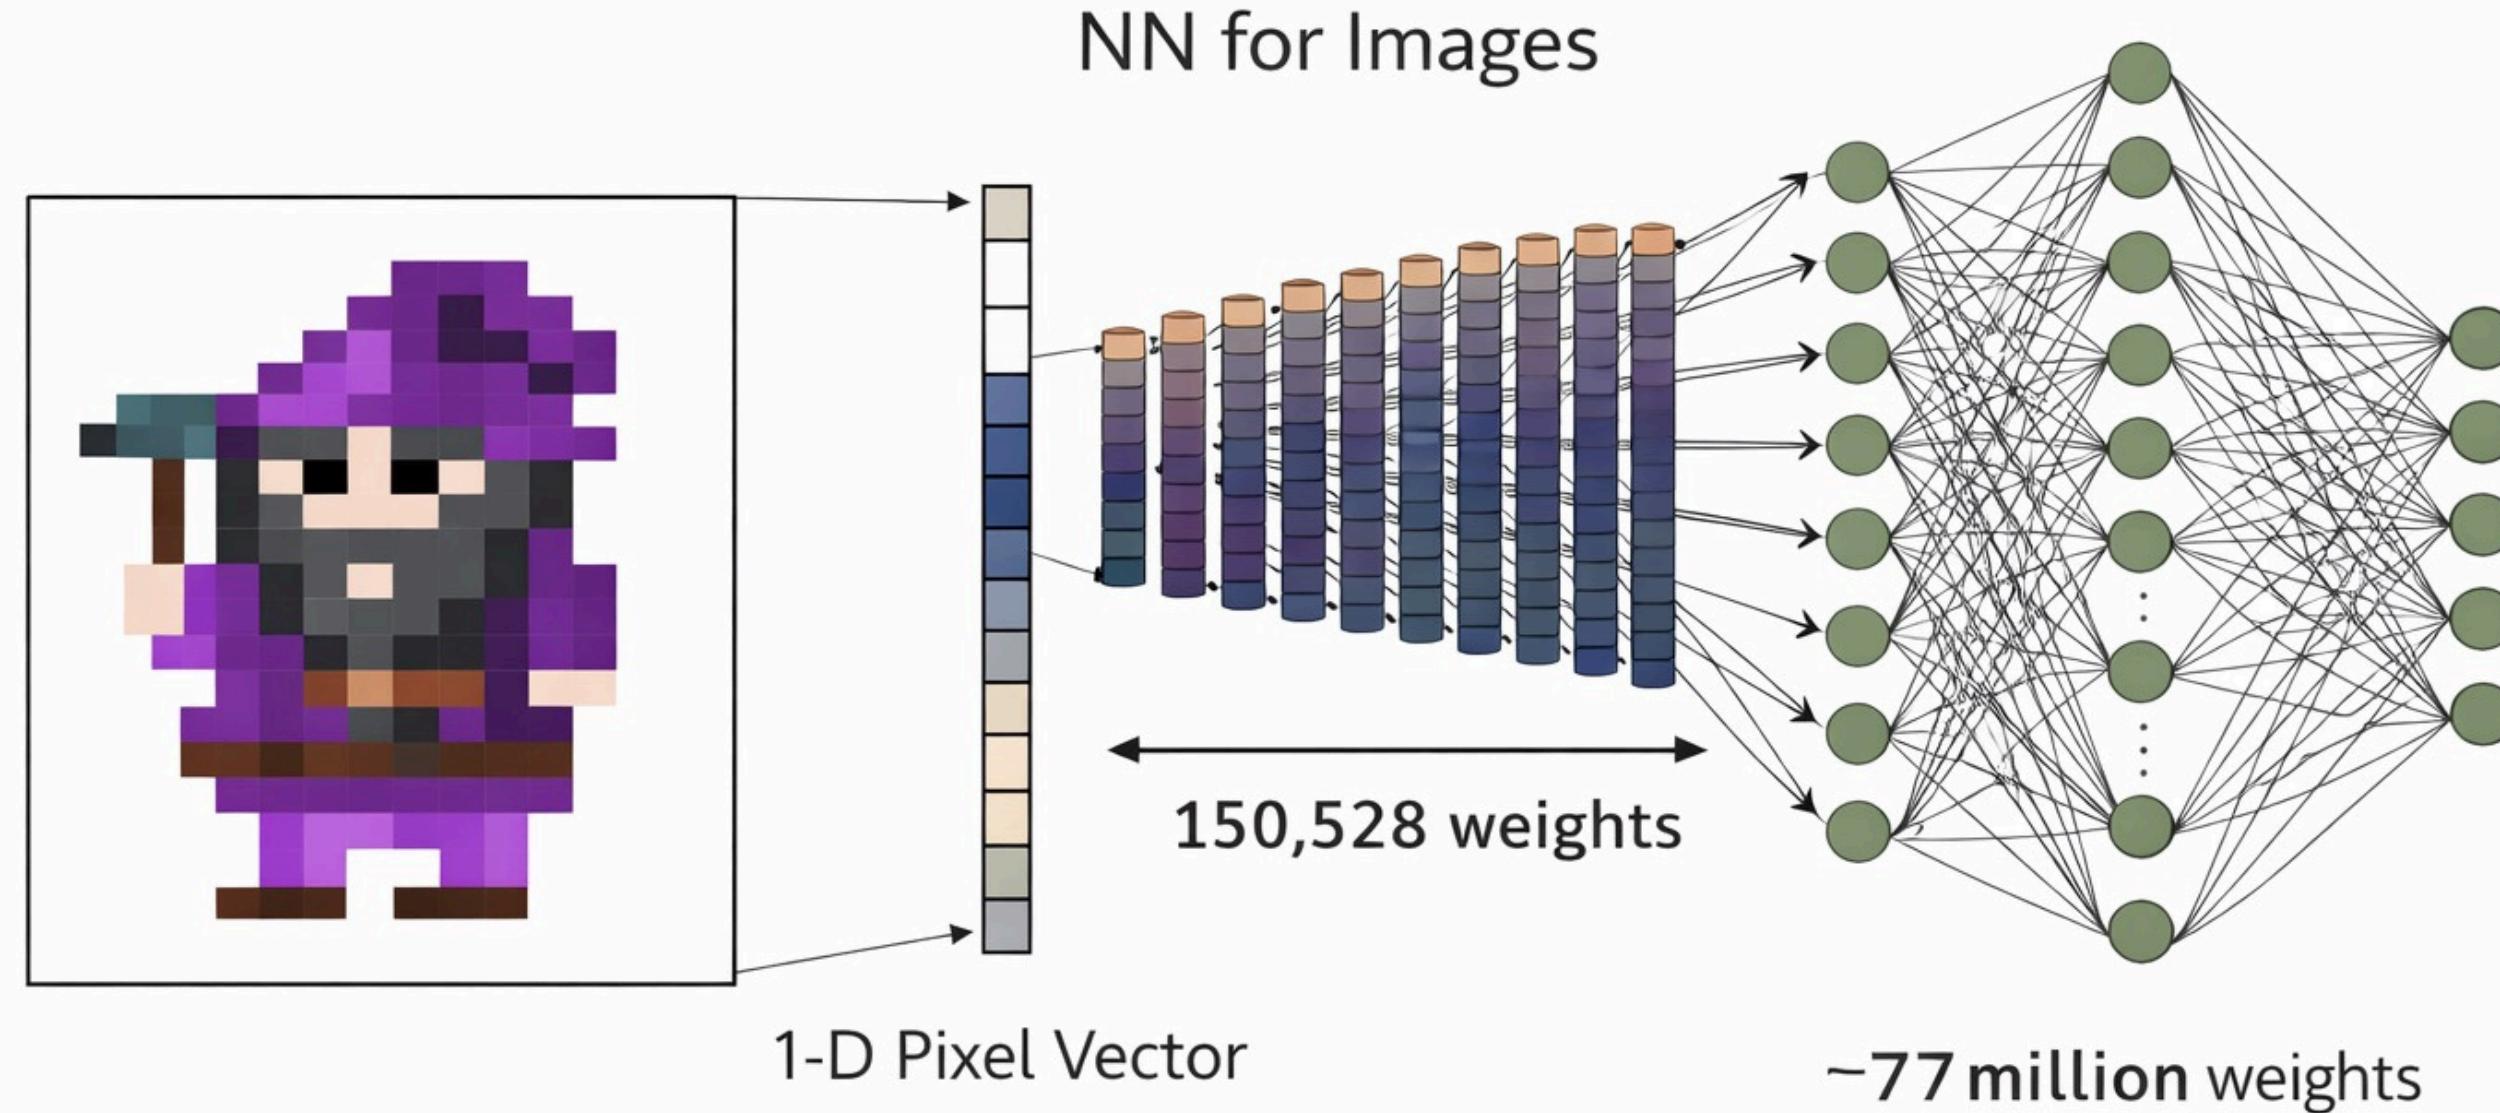
› How to build vision models?



› But!

# Convolutional Neural Networks

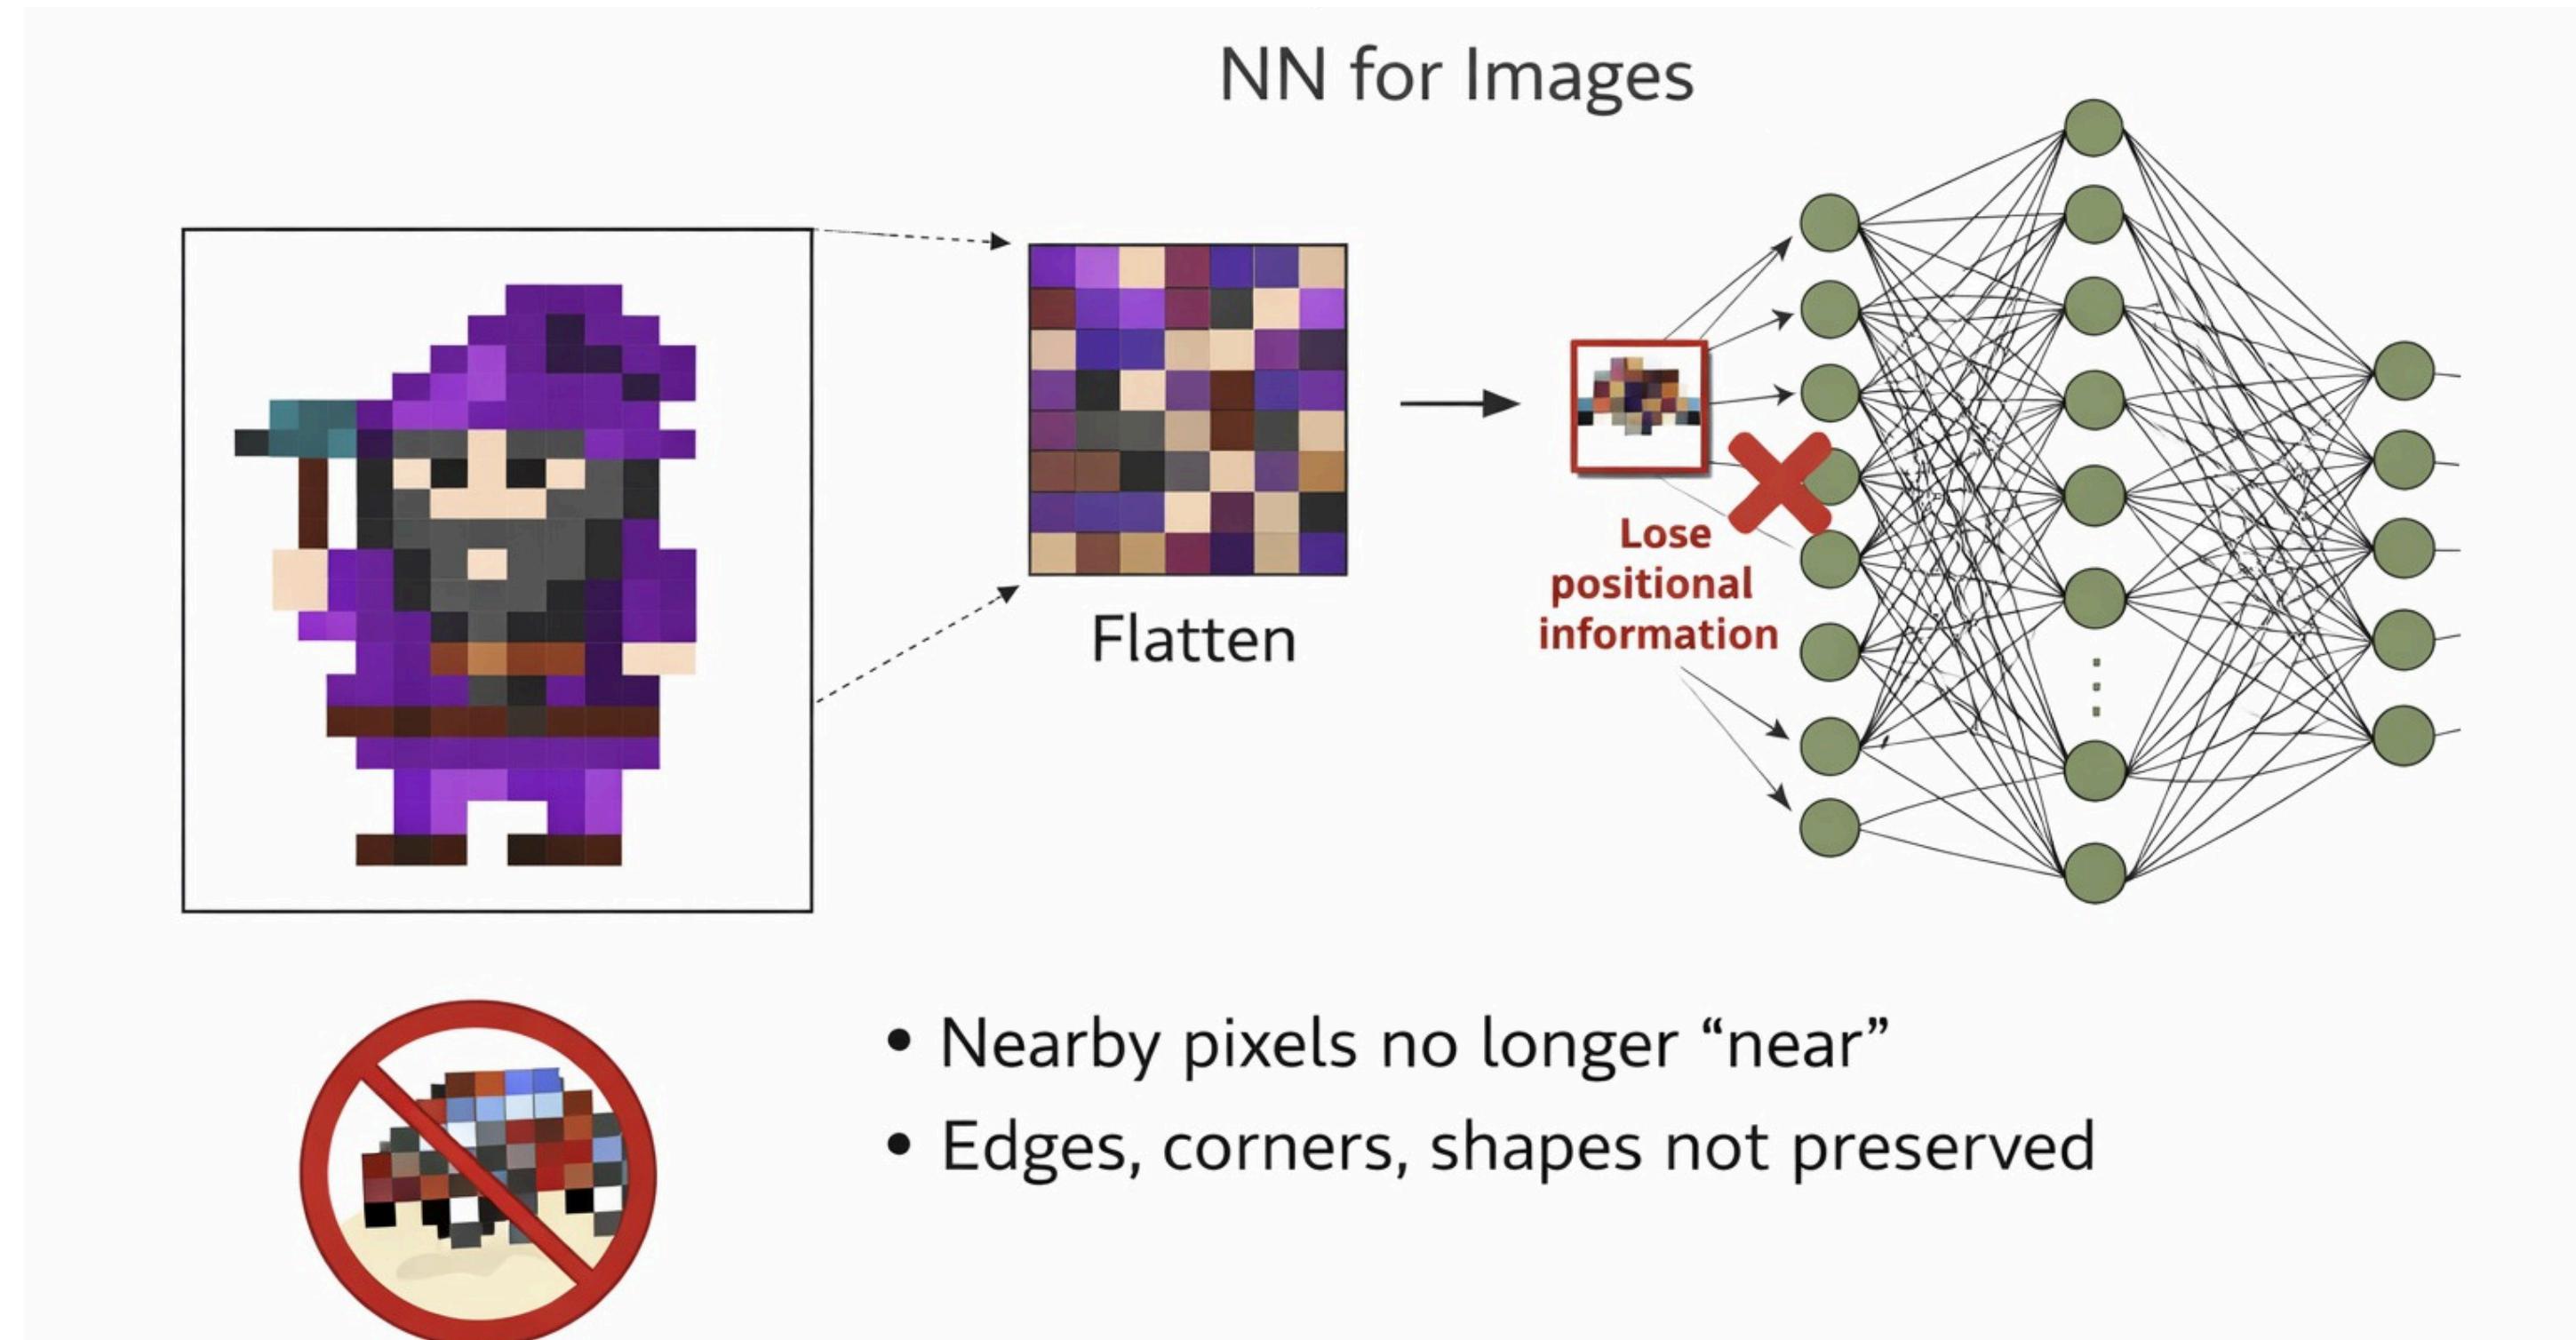
- Huge number of parameters



- High memory usage
- Slow training
- Overfitting
- Even small hidden layer (512 neurons) → ~77 million weights

# Convolutional Neural Networks

## > Loss of spatial structure



# Convolutional Neural Networks

- Little or no invariance to shifting, scaling, and other forms of distortion

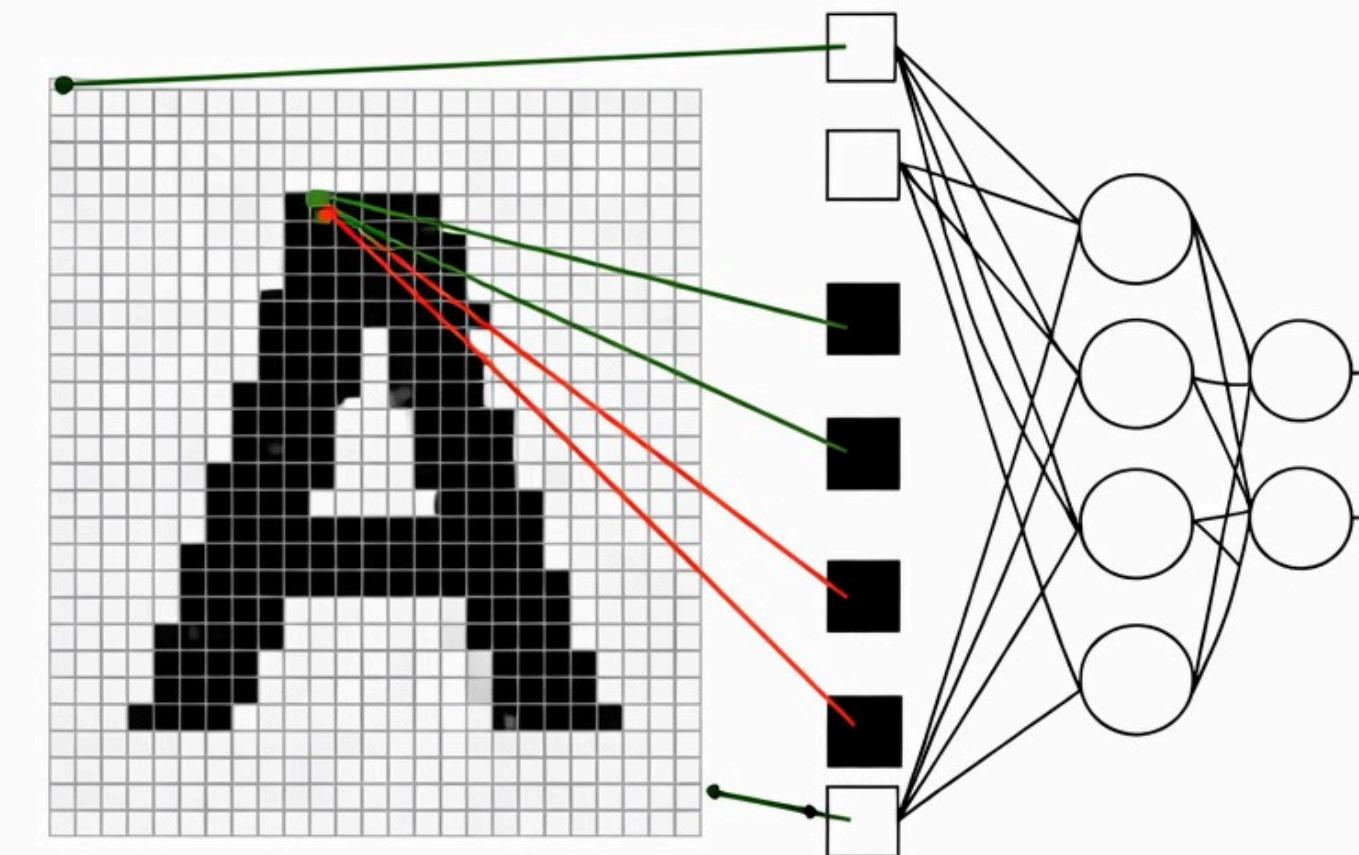


Figure 2: Original "A" character.

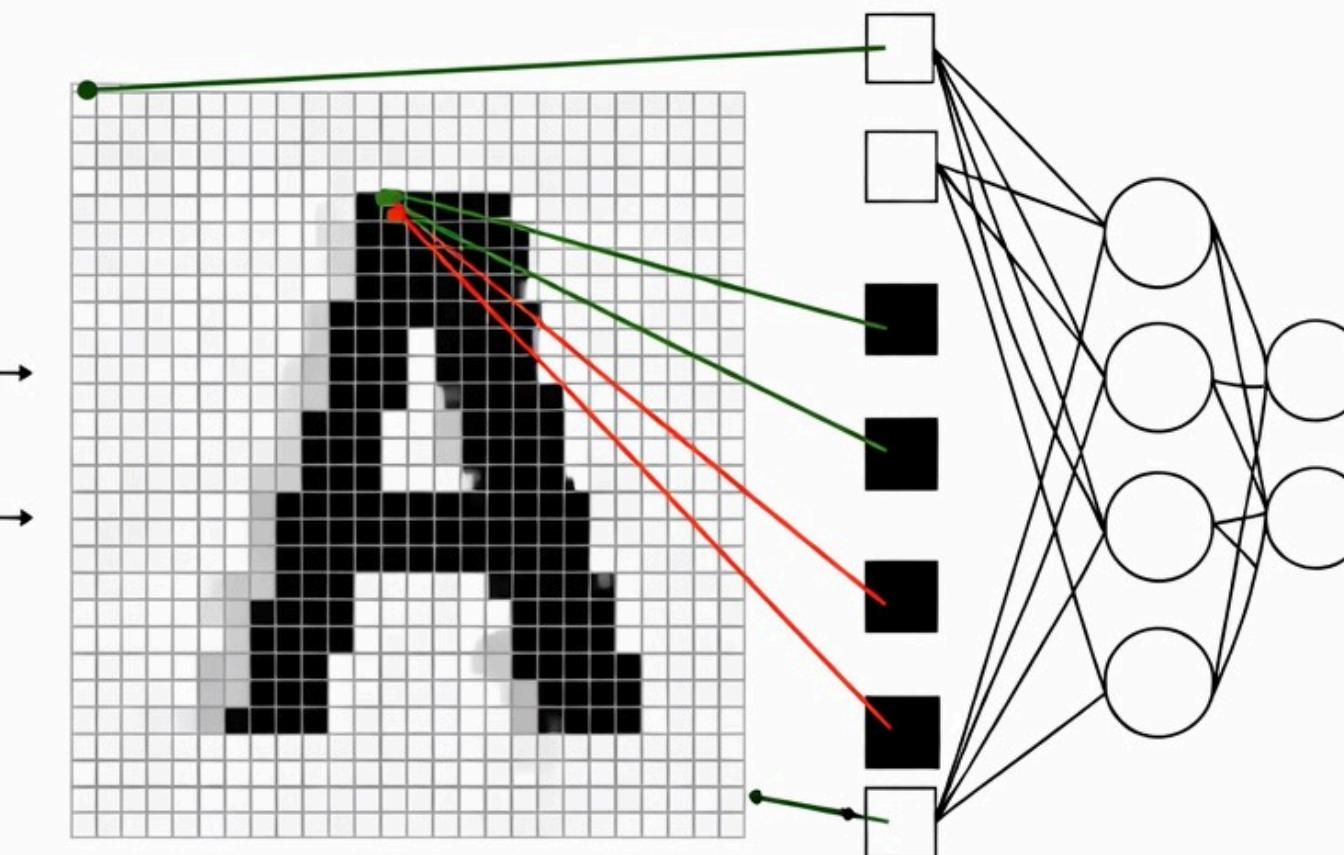
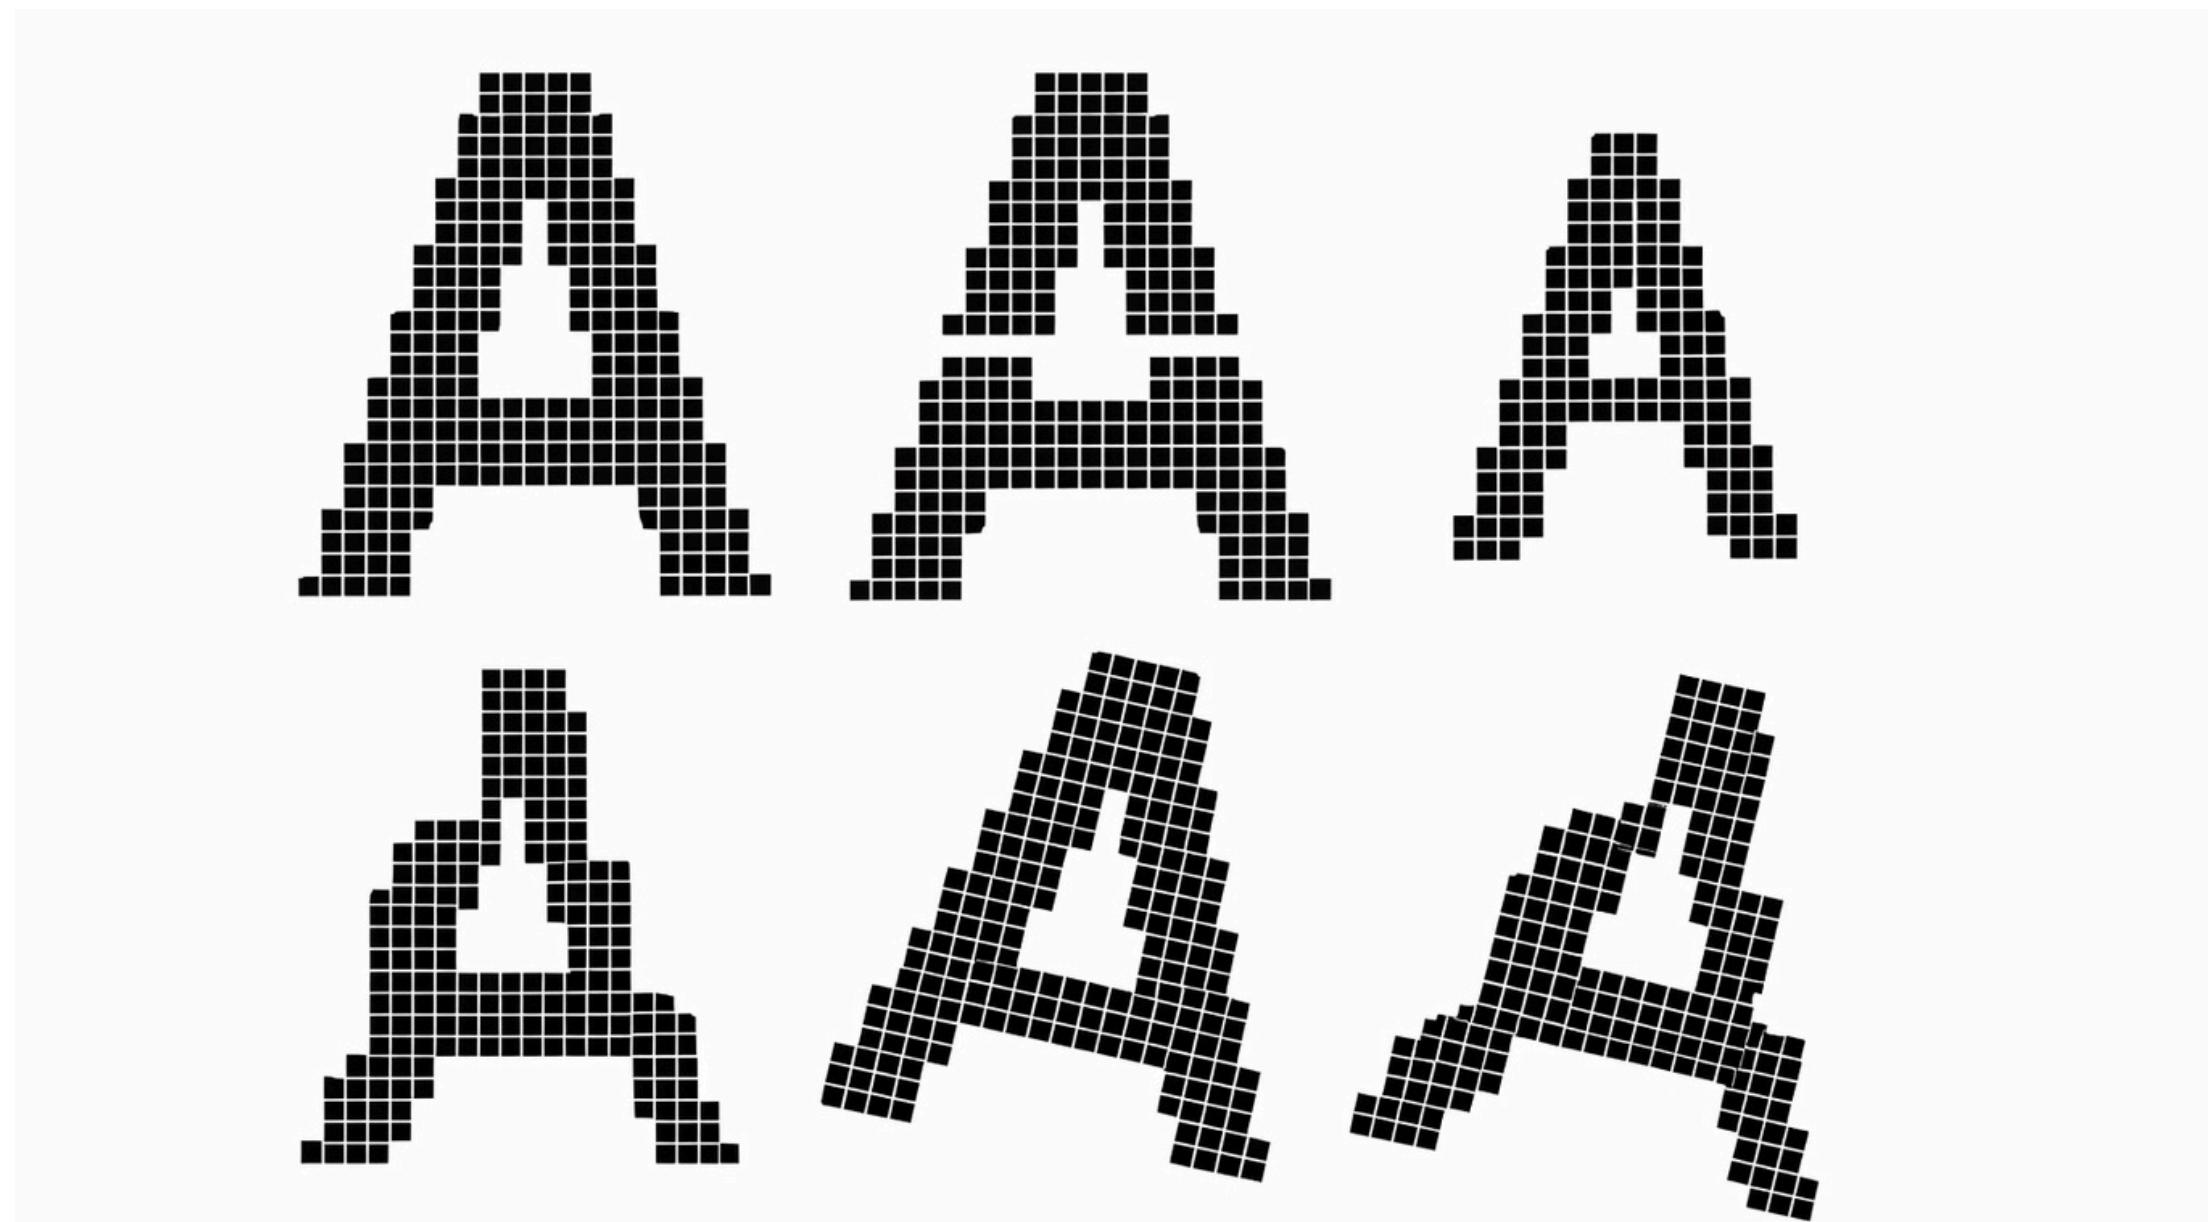


Figure 3: Shifted "A" character.

# Convolutional Neural Networks

---

- It treats pixels as independent features rather than structured patterns.
  - Black and white patterns  $2^{32 \times 32}$
  - Grayscale patterns  $2^{32 \times 32}$



# Convolutional Neural Networks

---

- We need a model that can:
  - **Learn visual patterns:** Automatically detect local features such as edges, corners, and shapes across the image.
  - **Operate efficiently:** Reduce computational cost by processing groups of pixels together instead of treating each pixel independently.
  - **Be robust to variations:** Correctly recognize objects despite changes in position, orientation, scale, or appearance.

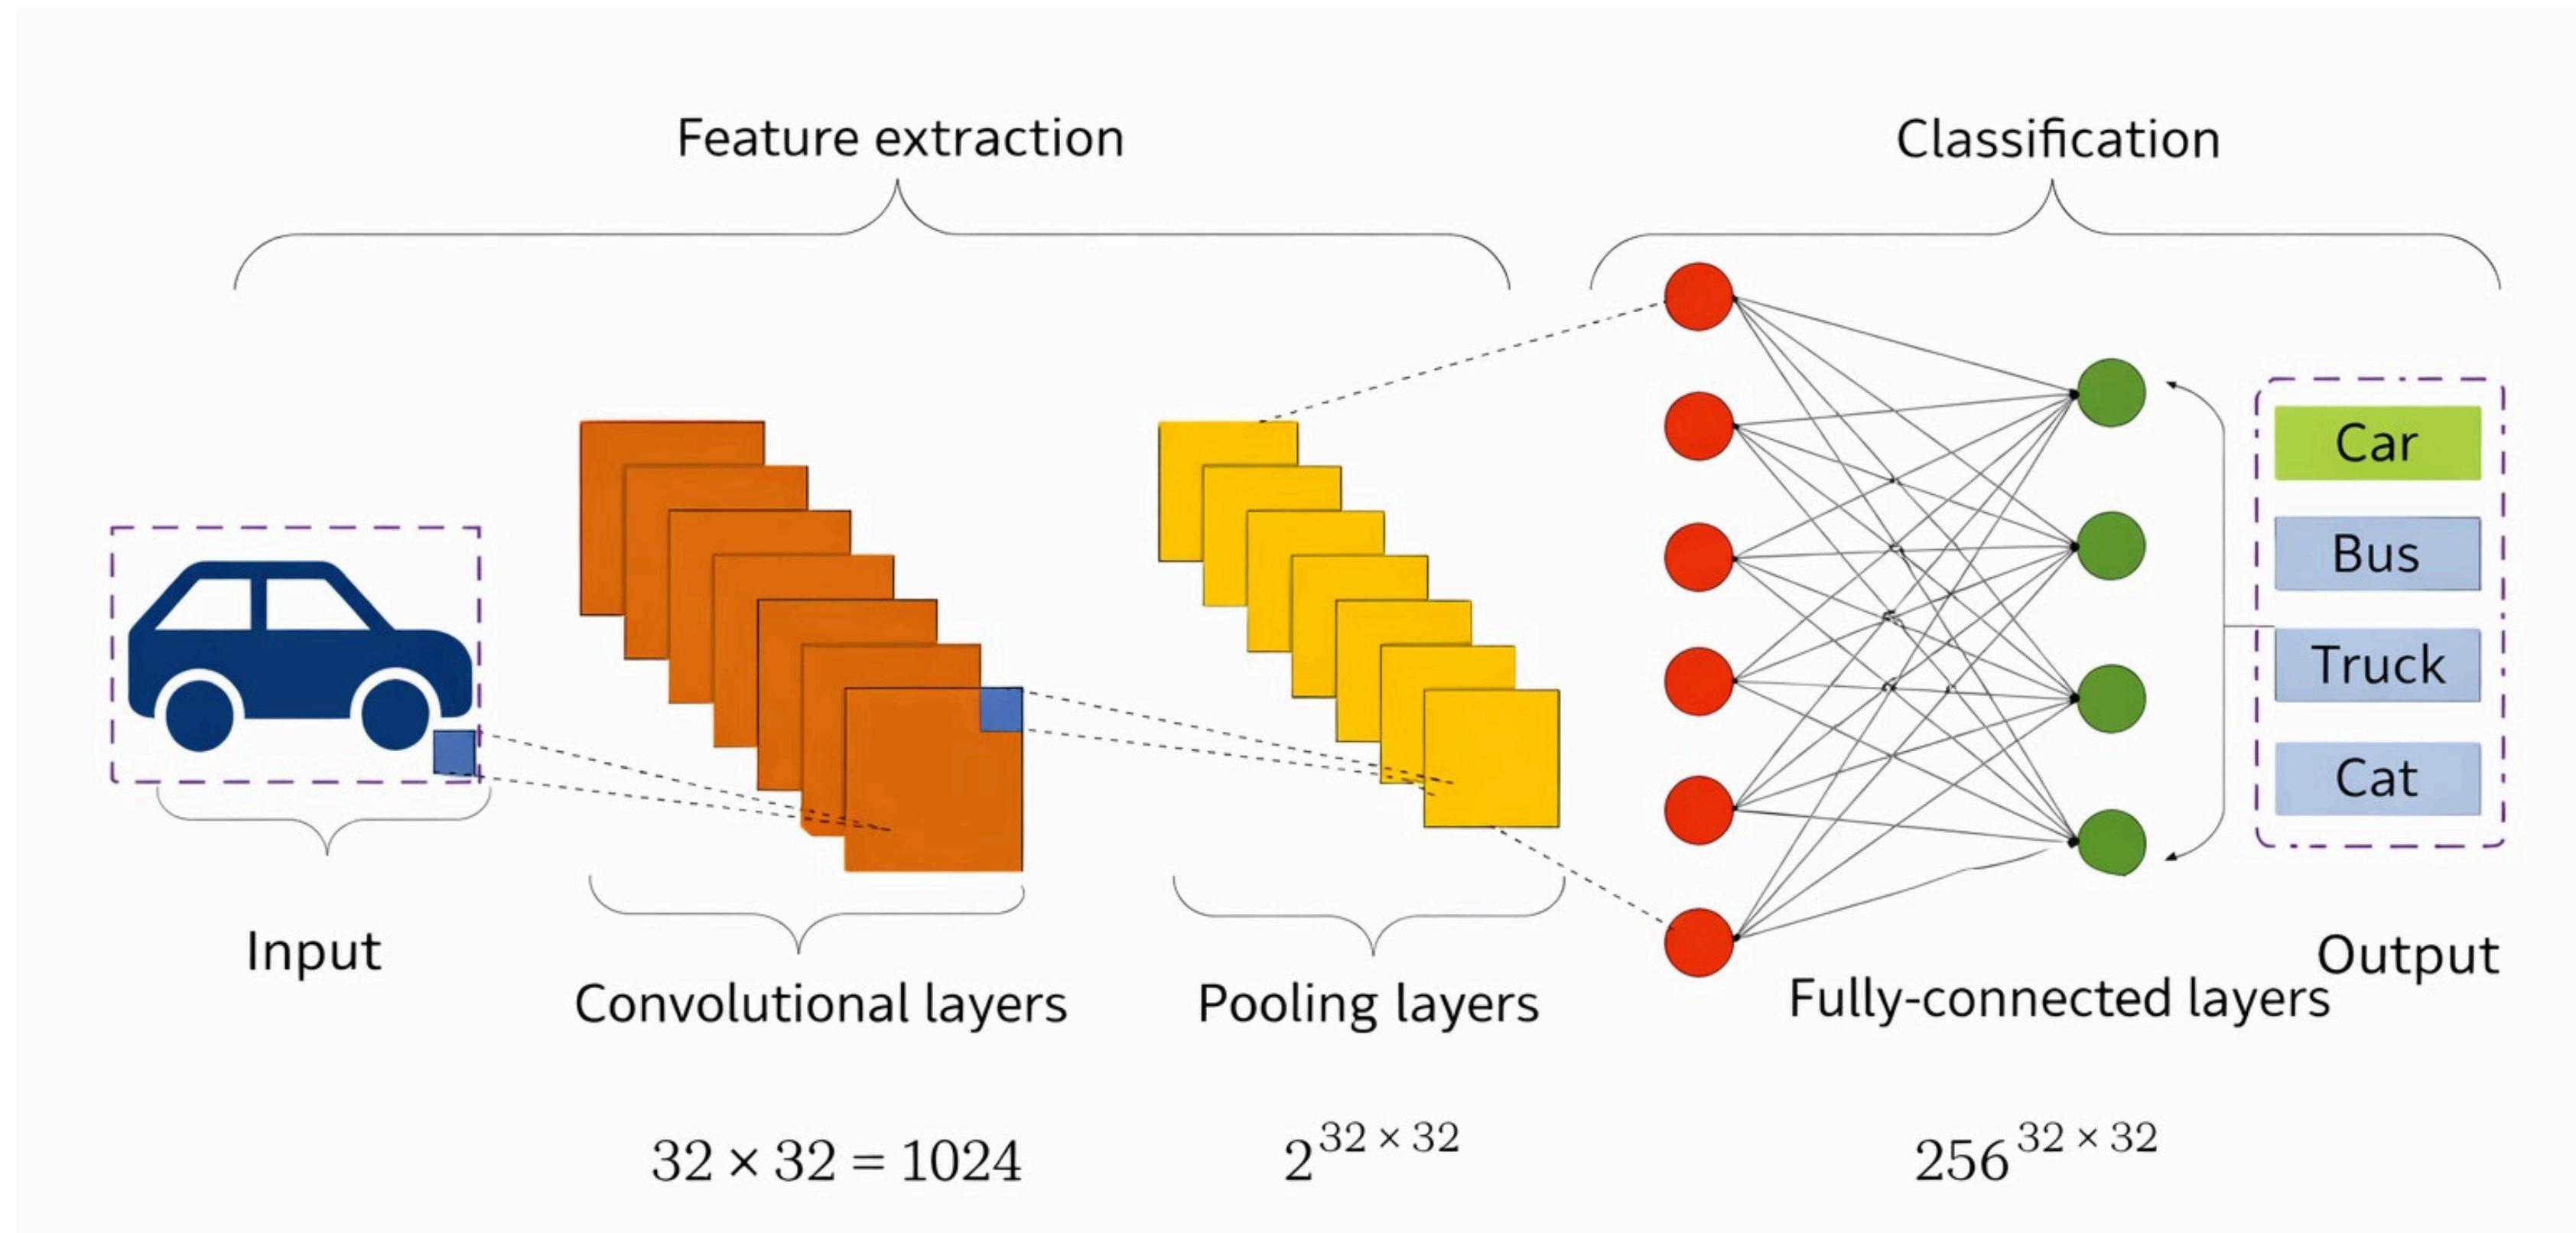
# Convolutional Neural Networks

---

- We need a model that can:
  - **Learn visual patterns:** Automatically detect local features such as edges, corners, and shapes across the image.
  - **Operate efficiently:** Reduce computational cost by processing groups of pixels together instead of treating each pixel independently.
  - **Be robust to variations:** Correctly recognize objects despite changes in position, orientation, scale, or appearance.

*We need Convolutional Neural Networks (CNNs)!*

# Convolutional Neural Networks



# Convolutional Neural Networks

---

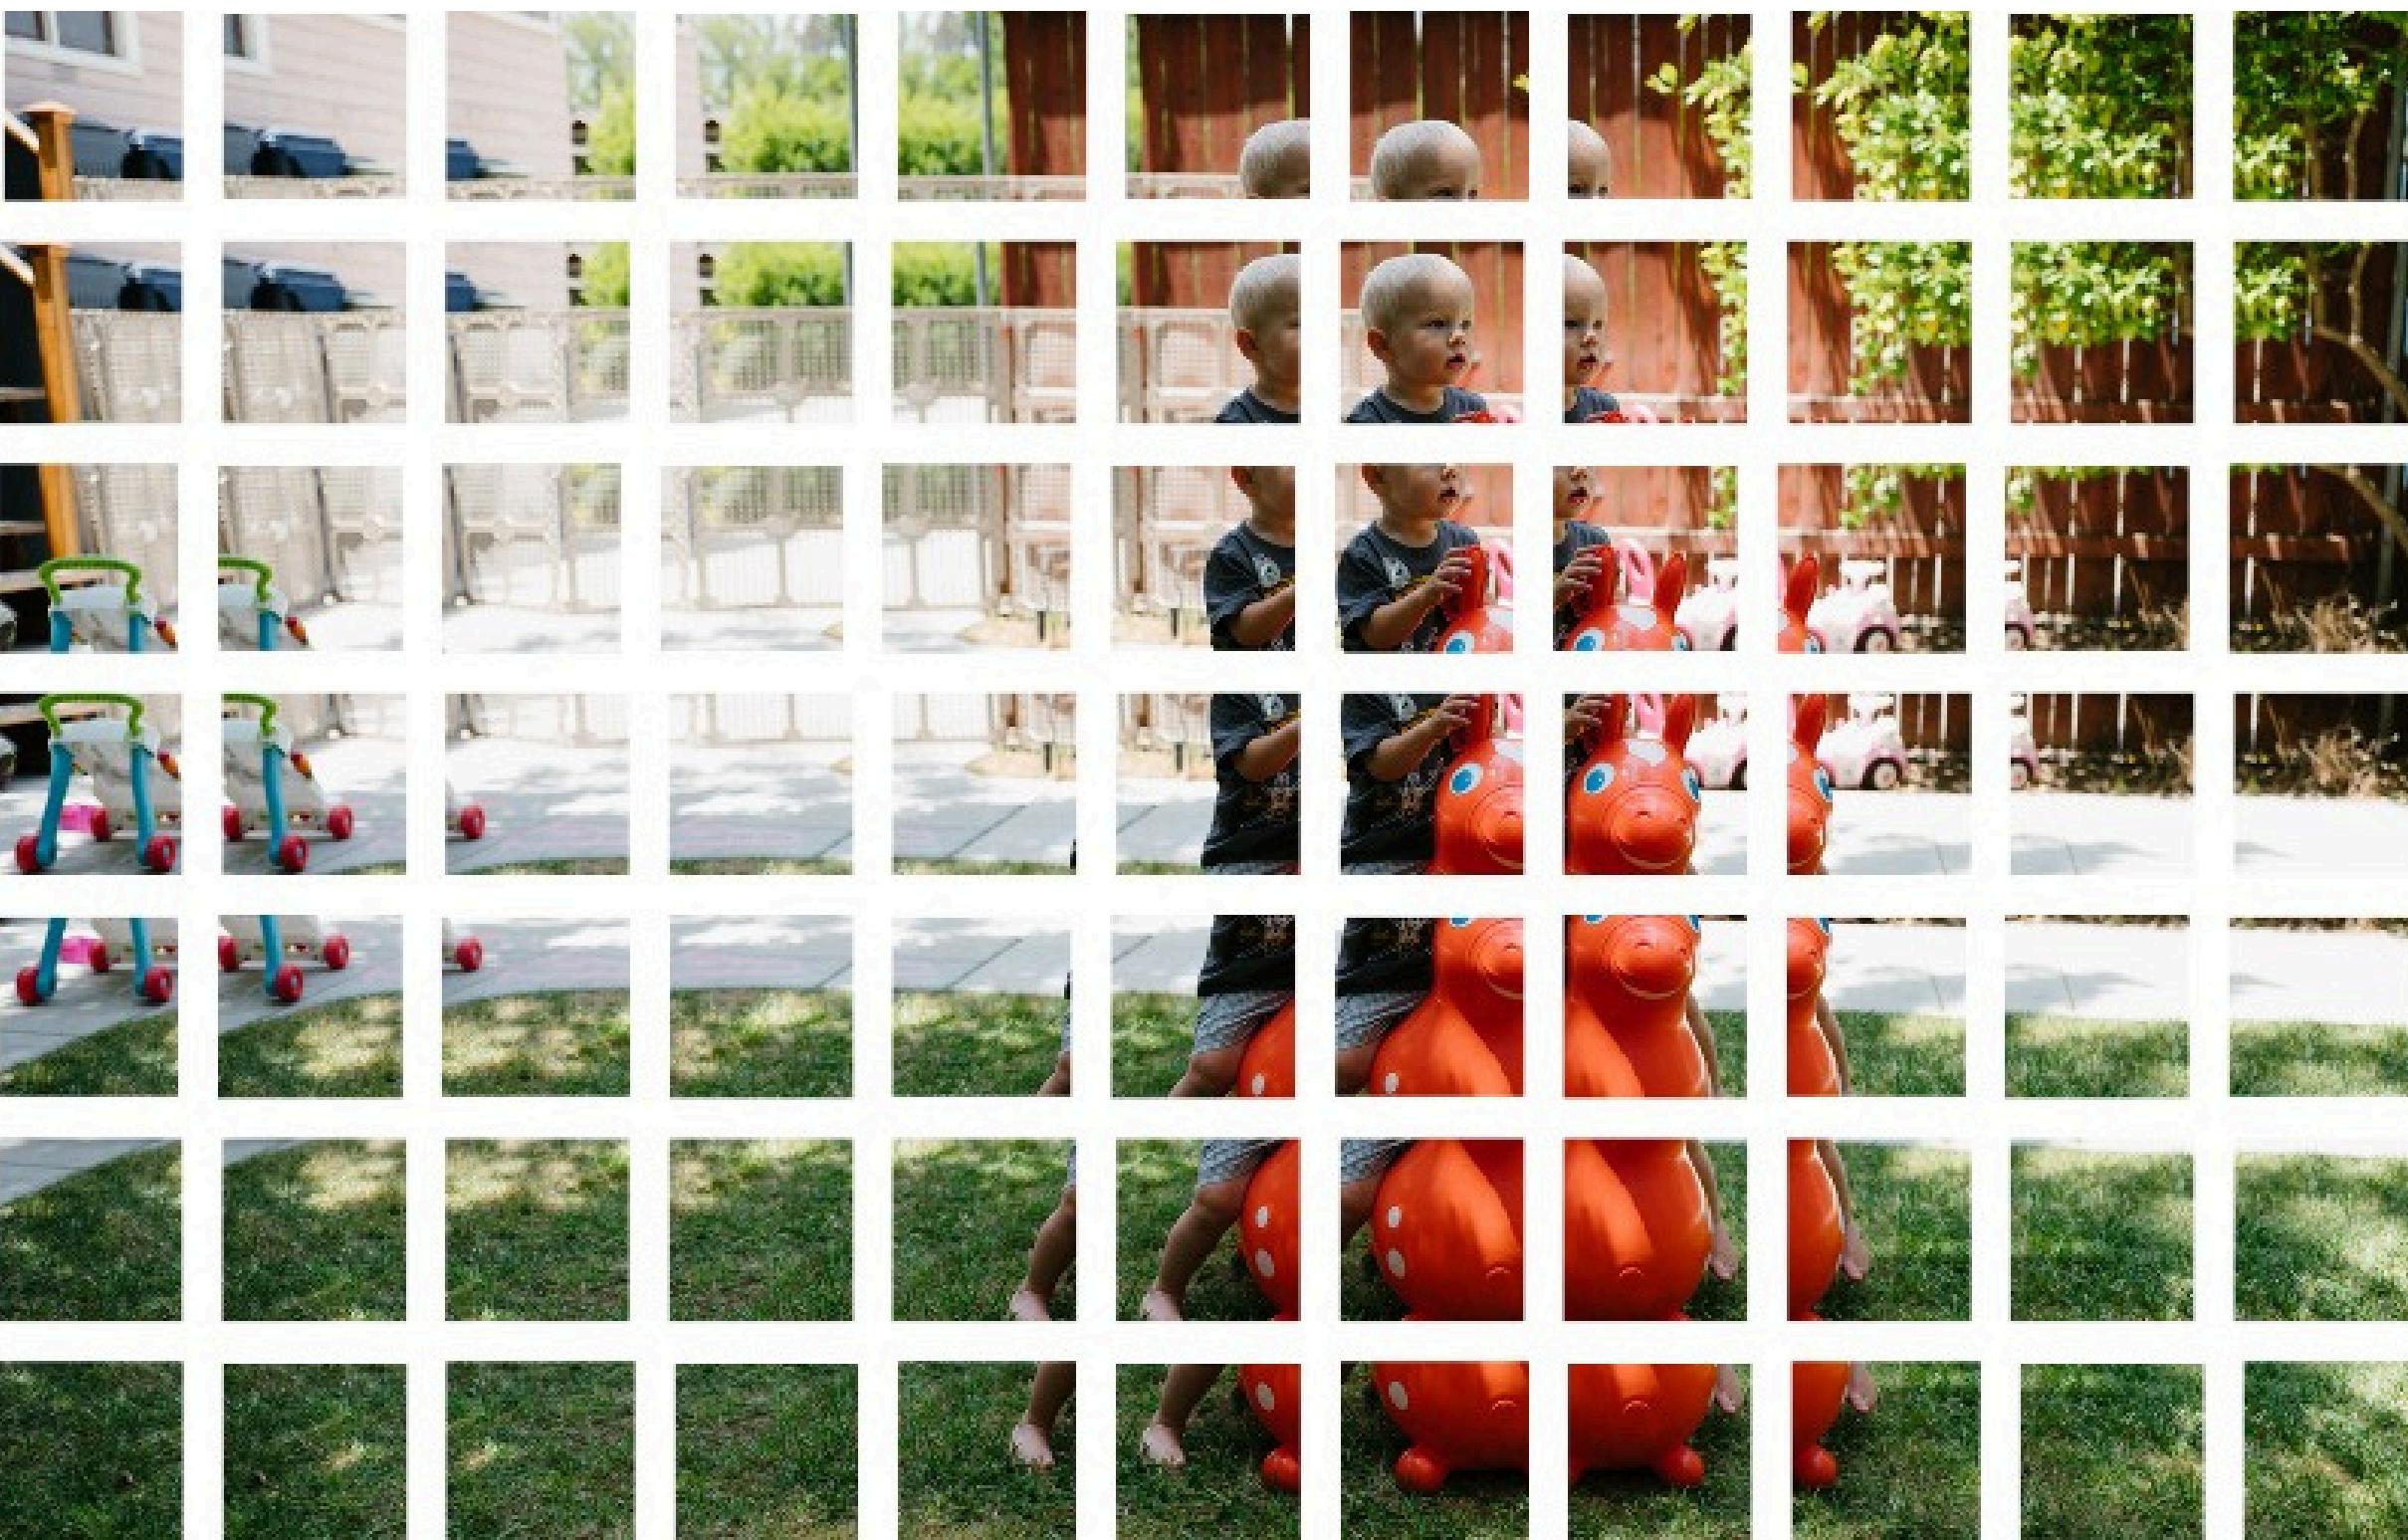
## > Convolution Layer



# Convolutional Neural Networks

---

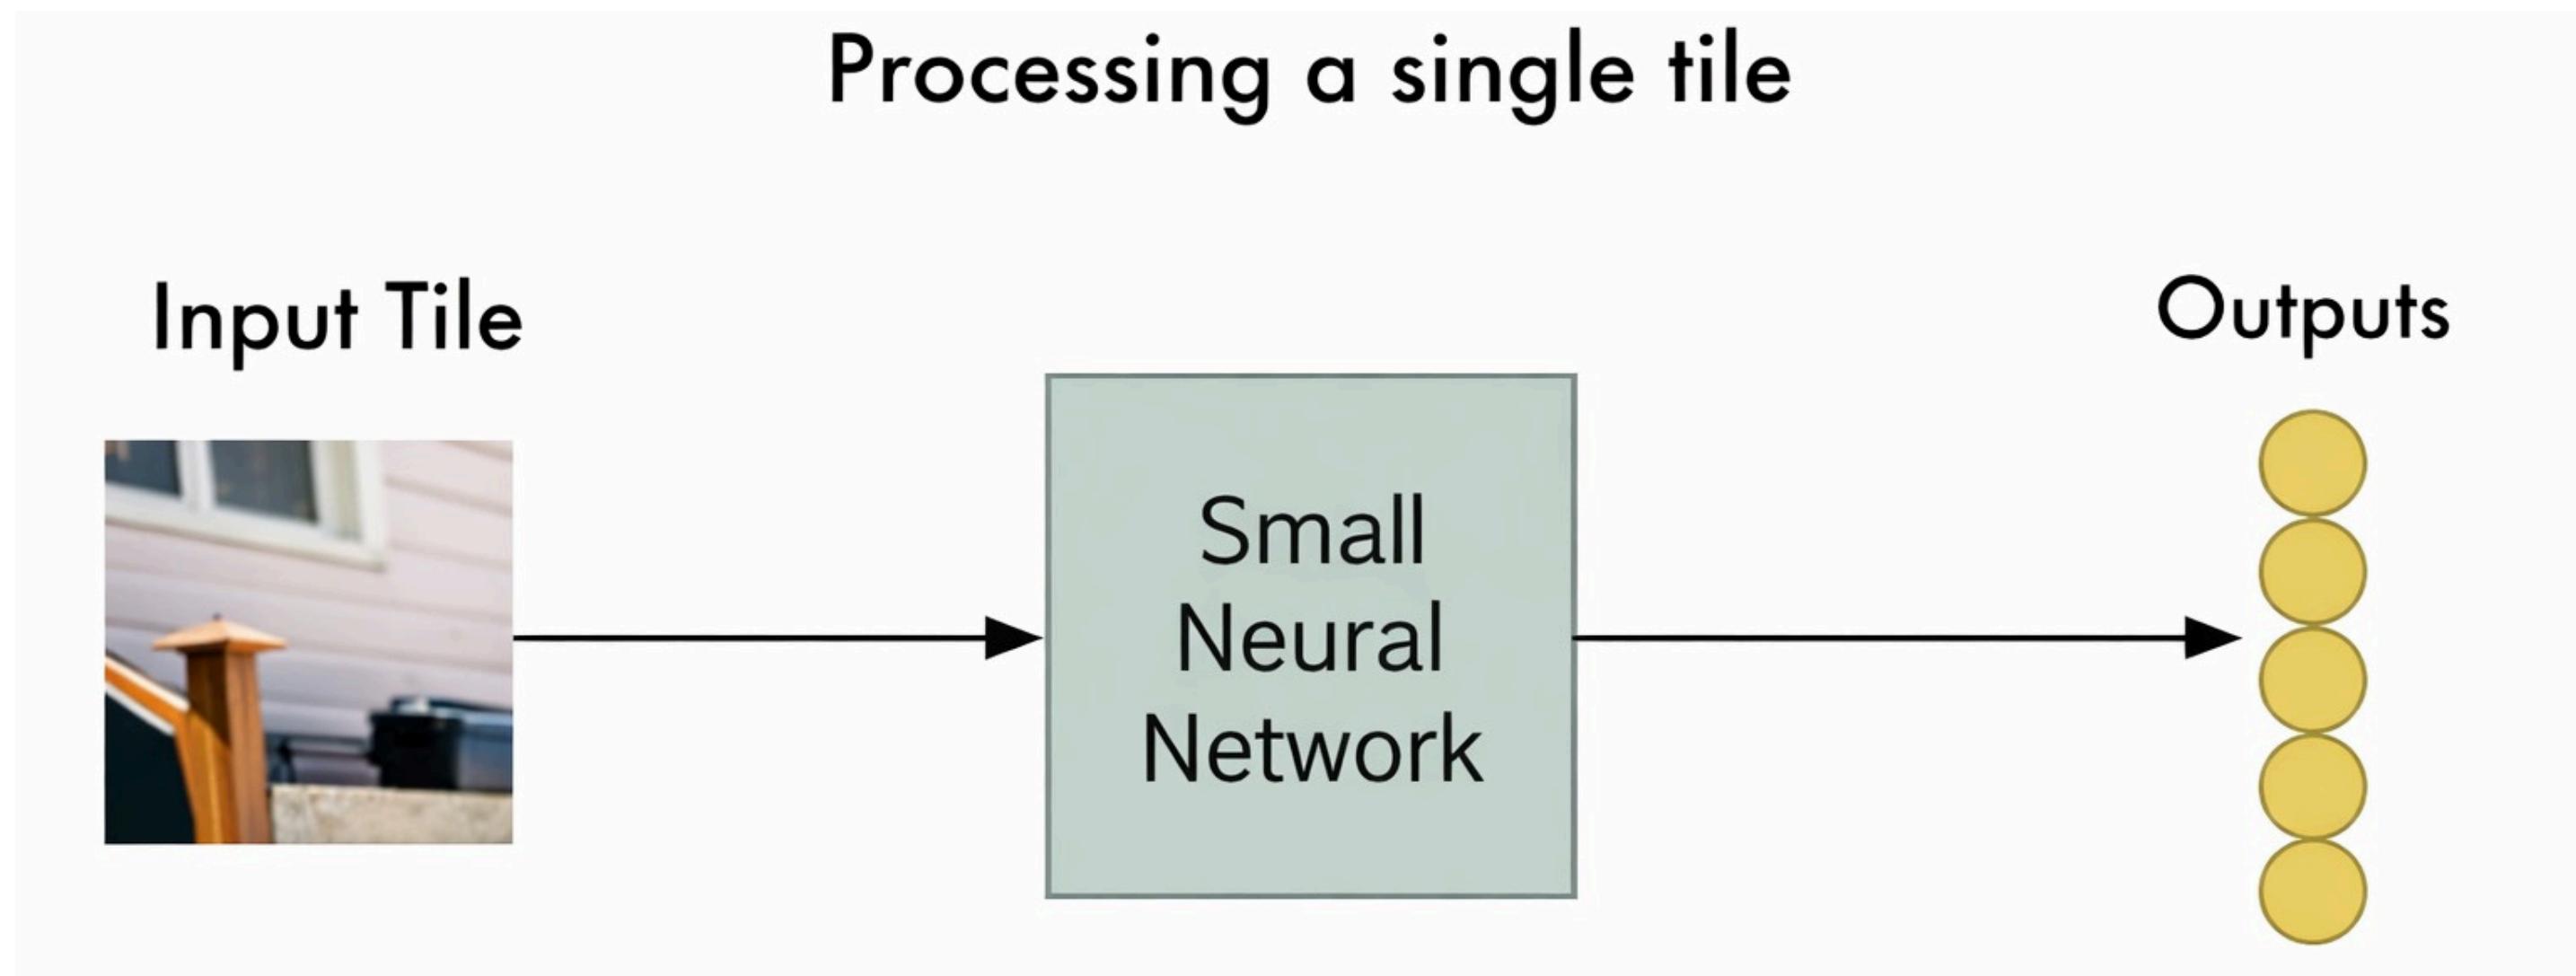
## > Convolution Layer



# Convolutional Neural Networks

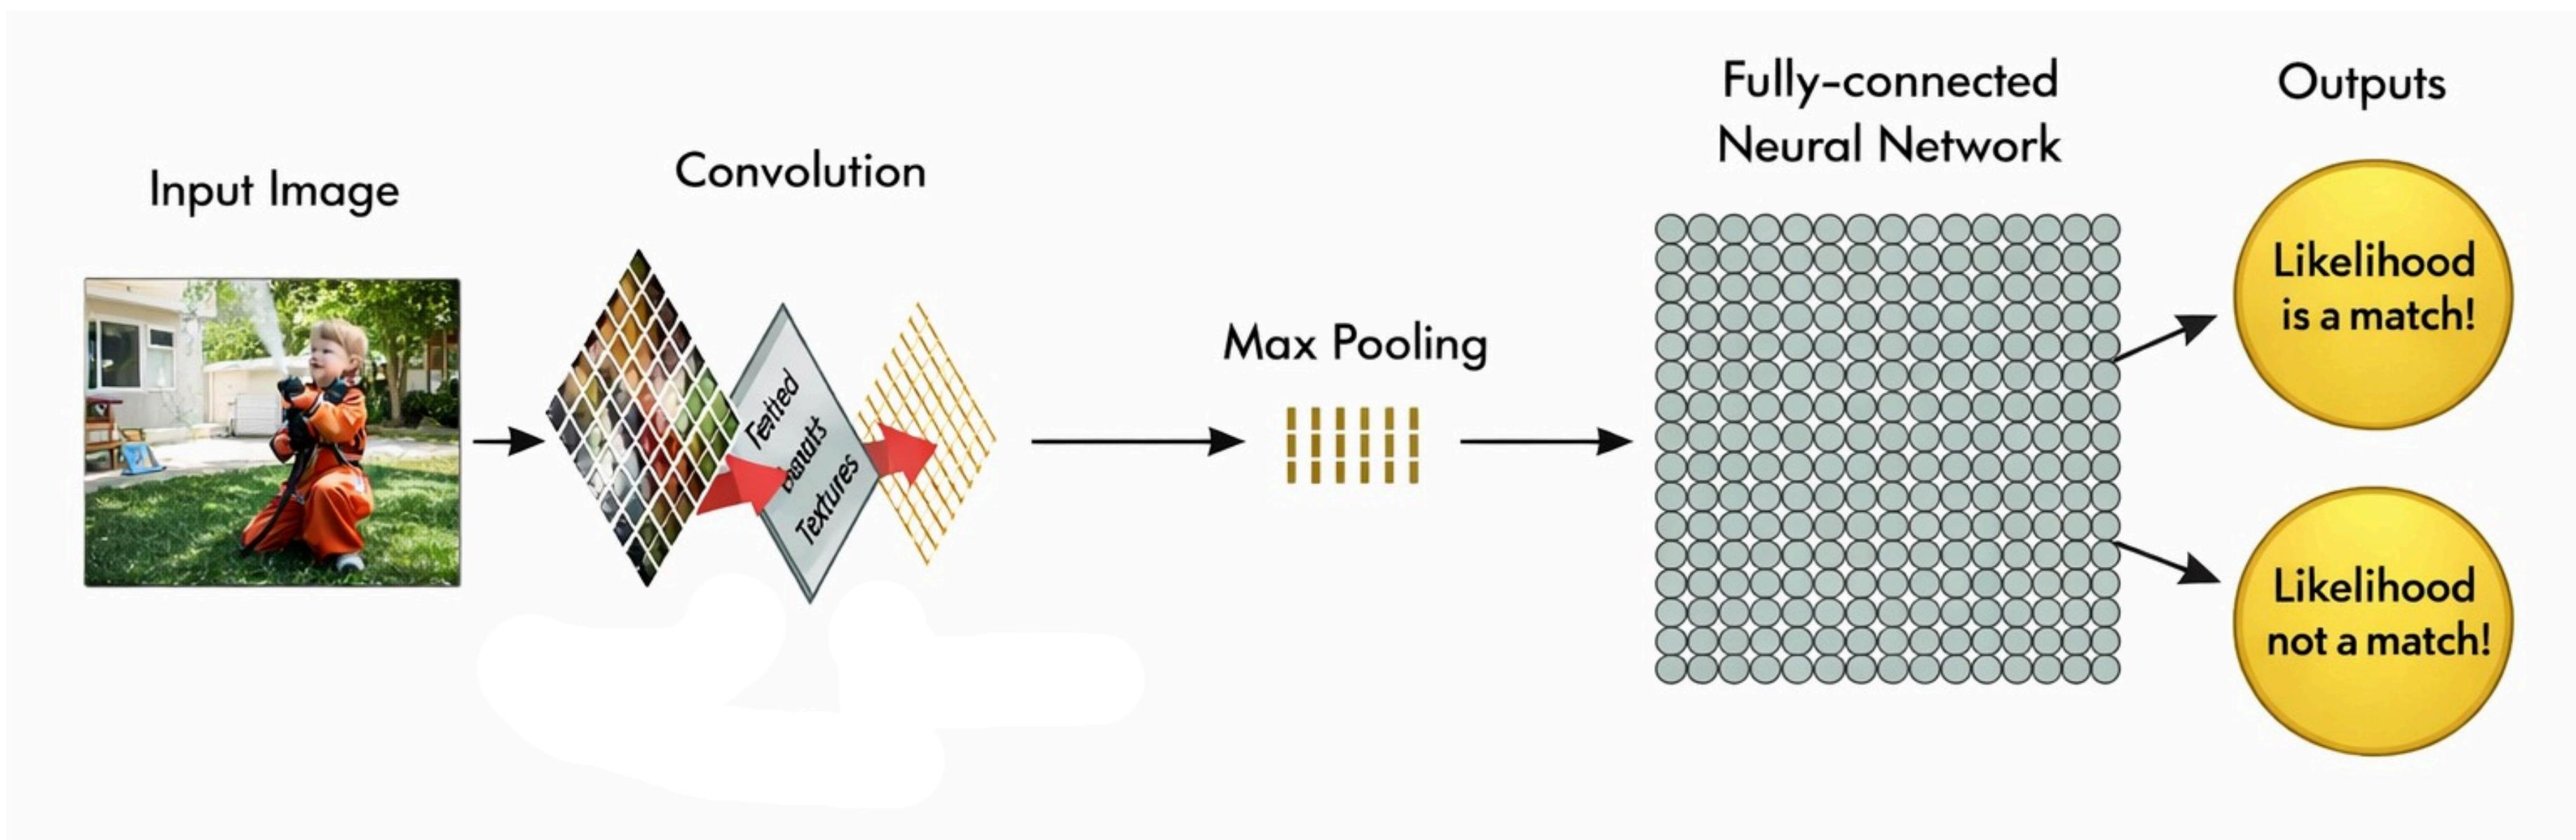
---

- We process each tile using the weights matrix of a small neural network.



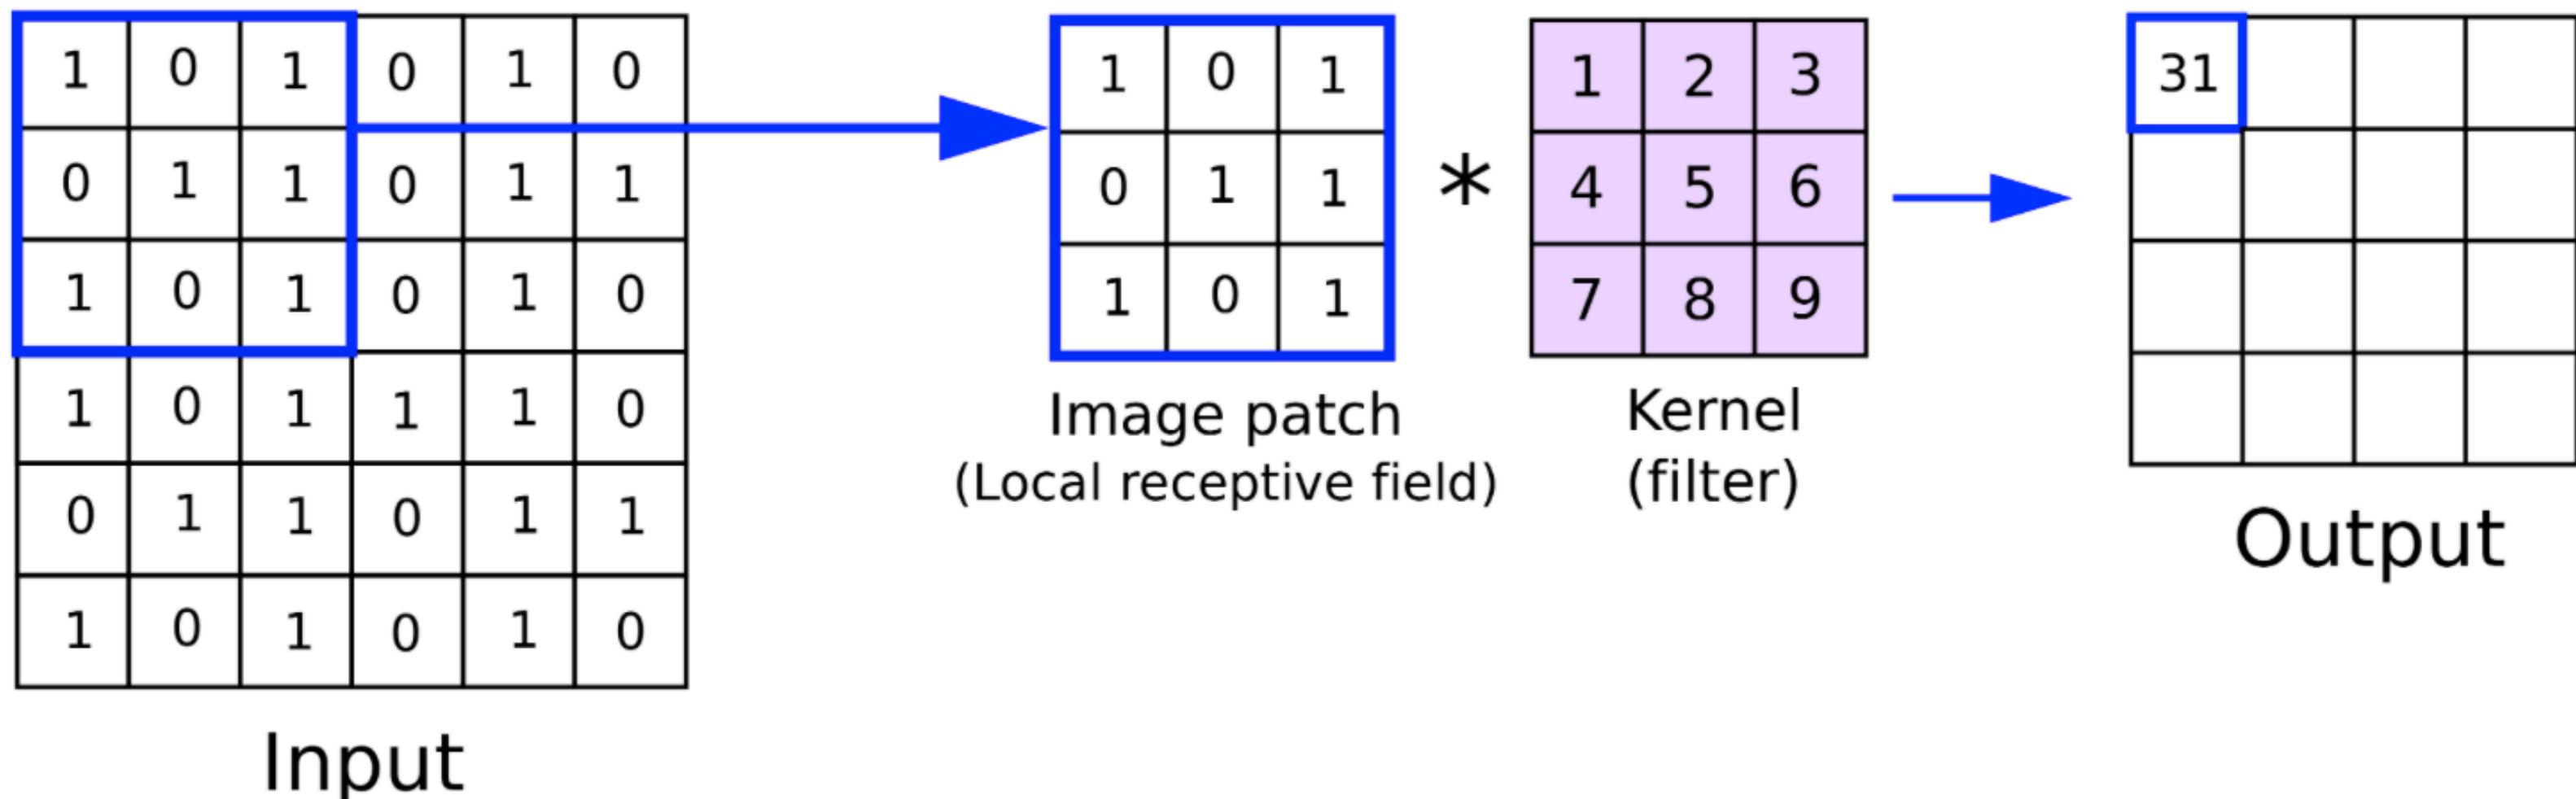
# Convolutional Neural Networks

- › We process each tile using the weights matrix of a small neural network.



# Convolutional Neural Networks

- Convolution involves applying a kernel (filter) over an image to extract features.



# Convolutional Neural Networks

- Convolution involves applying a kernel (filter) over an image to extract features.

$$\begin{array}{|c|c|c|c|c|c|c|c|c|} \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ \hline \end{array} \quad * \quad \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline \end{array} \quad = \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline 0 & 0 & 30 & 30 & 0 & 0 \\ \hline \end{array}$$

Vertical

The diagram illustrates the convolution process for vertical edge detection. The input image is an 8x8 grid of values, mostly 10 with some 0s. A 3x3 kernel with values 1, 0, -1 (repeated) is applied. The result is an 8x8 output image where the bottom row has values 30, 30, 0, 0, 0, 0, 0, 0, and the bottom-left cell is highlighted in green. The input image has three colored boxes: a blue one at the top-left, a pink one in the middle-left, and a green one at the bottom-left. The output image has three colored boxes: a blue one at the top-left, a pink one in the middle-left, and a green one at the bottom-left.

*An example of vertical edge detection*

# Convolutional Neural Networks

- Convolution involves applying a kernel (filter) over an image to extract features.

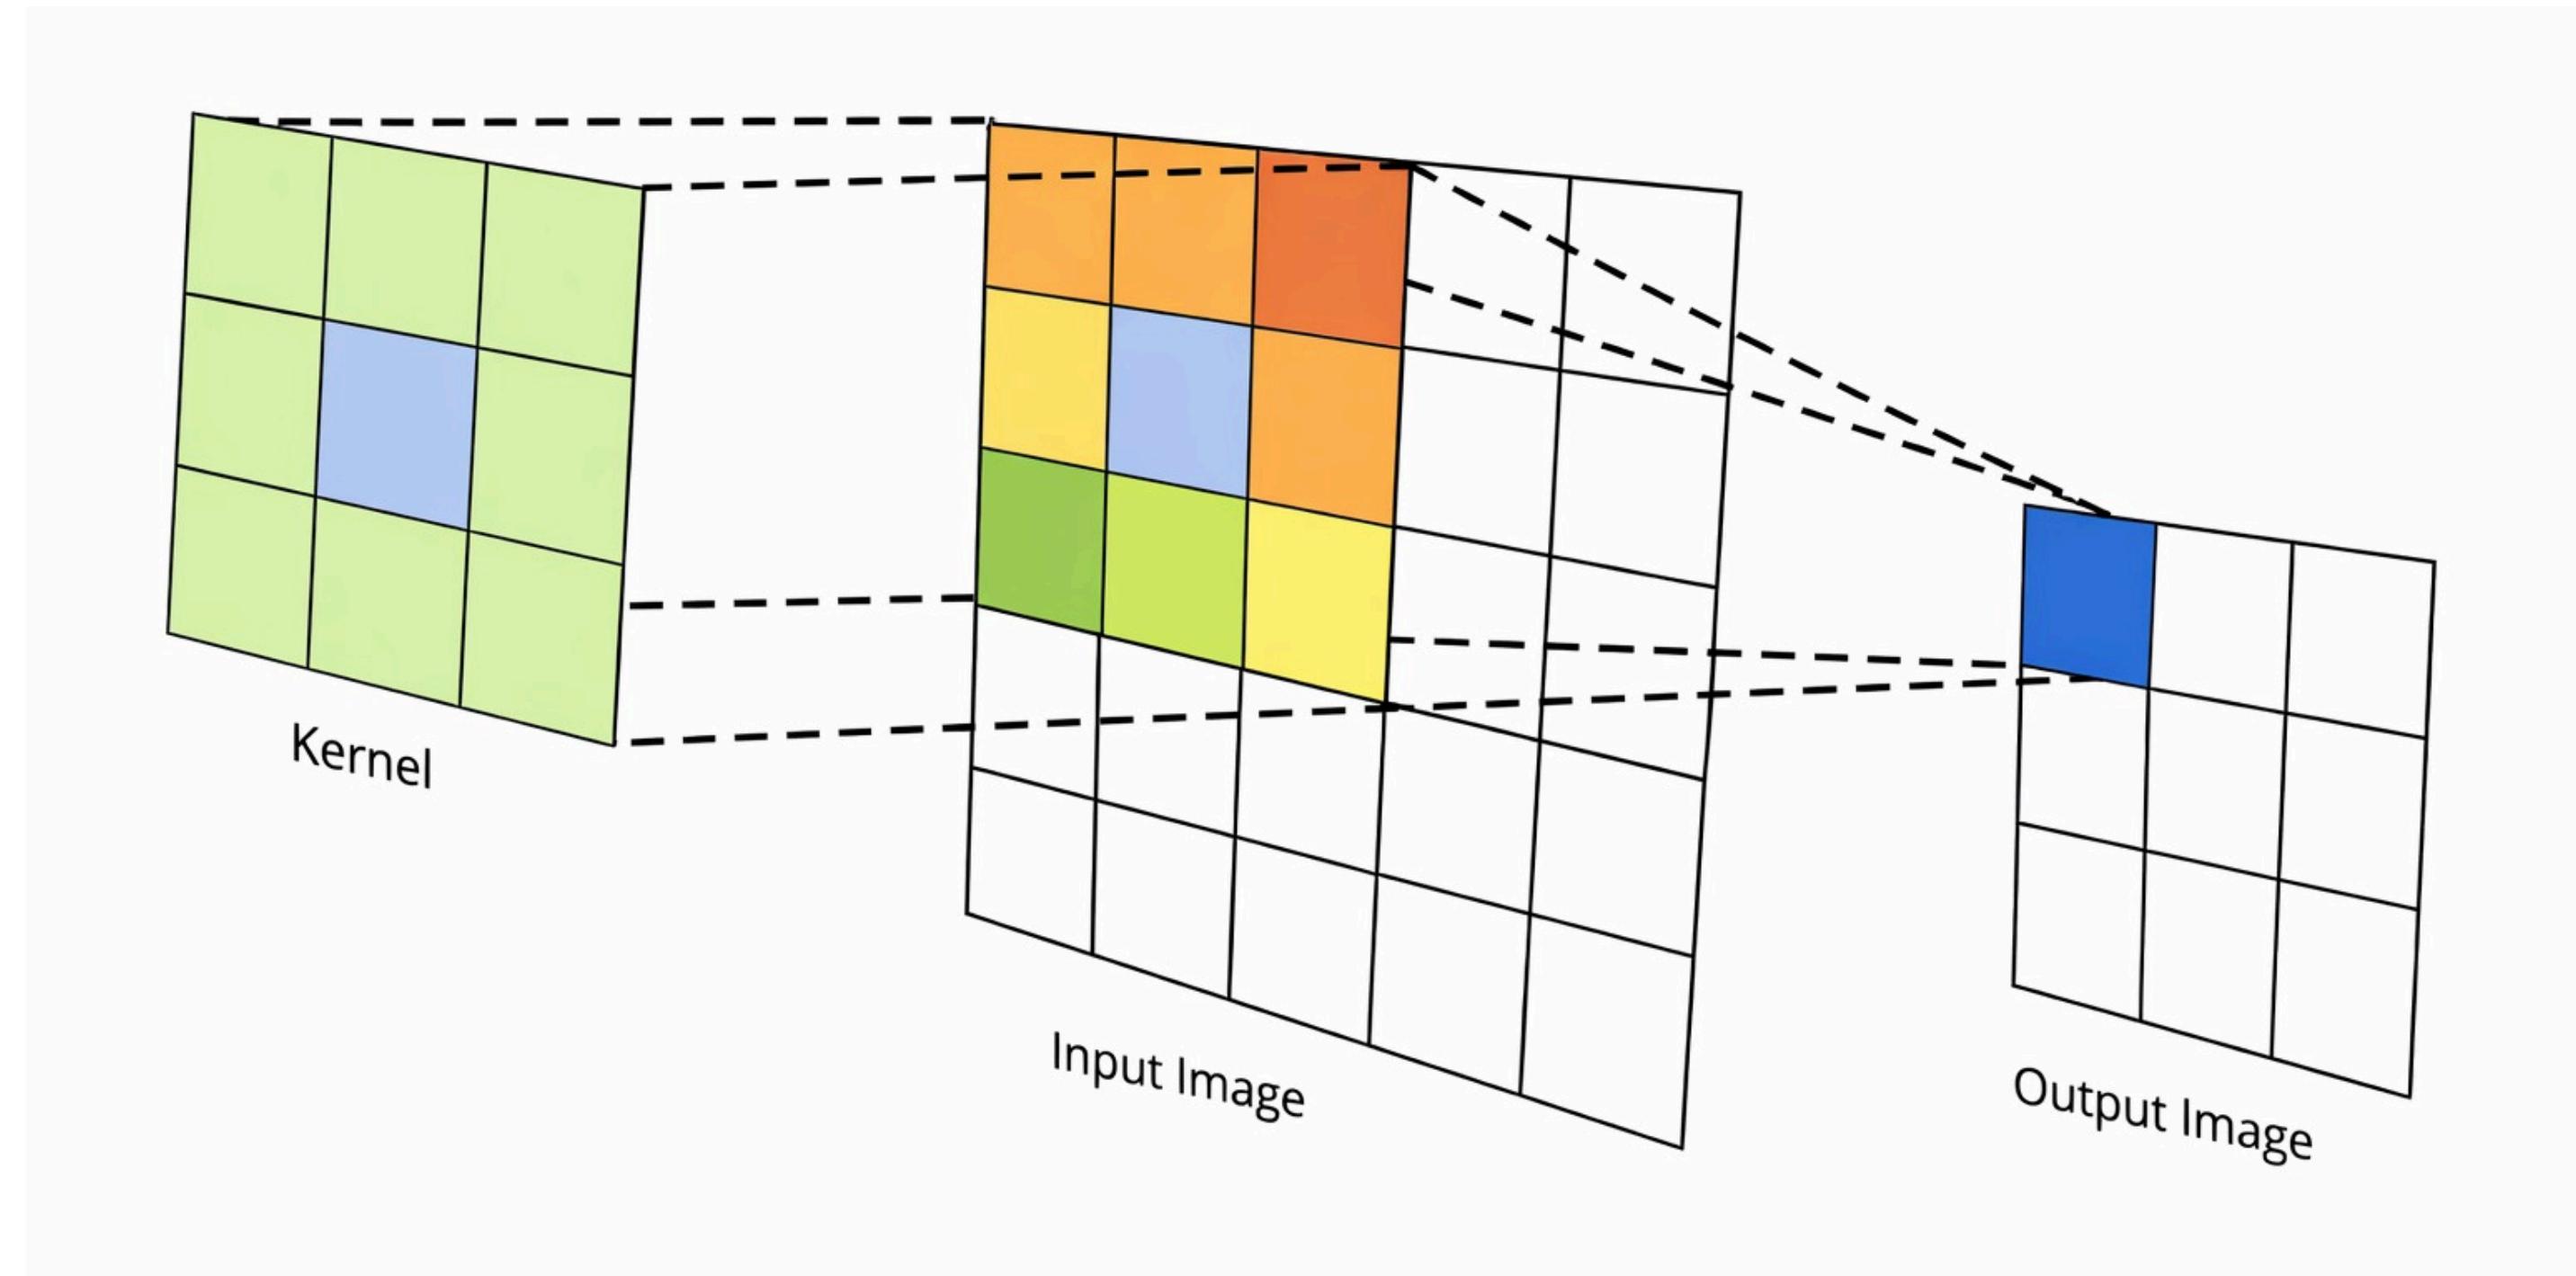
$$\begin{matrix} 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 10 & 10 & 10 & 10 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 10 & 10 & 10 & 10 \\ 0 & 0 & 0 & 0 & 10 & 10 & 10 & 10 \\ 0 & 0 & 0 & 0 & 10 & 10 & 10 & 10 \\ 0 & 0 & 0 & 0 & 10 & 10 & 10 & 10 \end{matrix} * \begin{matrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{matrix}_{\text{Horizontal}} = \begin{matrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 30 & 30 & 10 & -10 & -30 & -30 \\ 30 & 30 & 10 & -10 & -30 & -30 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{matrix}$$

*An example of horizontal edge detection*

# Convolutional Neural Networks

---

- The kernel is slid across the image, computing a dot product with local regions at each position.



# Convolutional Neural Networks

---

- This can be represented mathematically by:

$$z_{i,j} = (W * X)_{i,j}$$

$$z_{i,j} = \sum_{a=0}^{m-1} \sum_{b=0}^{n-1} W_{a,b} X_{i+a, j+b}$$

$z_{i,j}$  : output of the convolution at position  $(i, j)$

$W$  : filter (kernel) matrix of size  $m \times n$

$X_{i,j}$  : input value at position  $(i, j)$

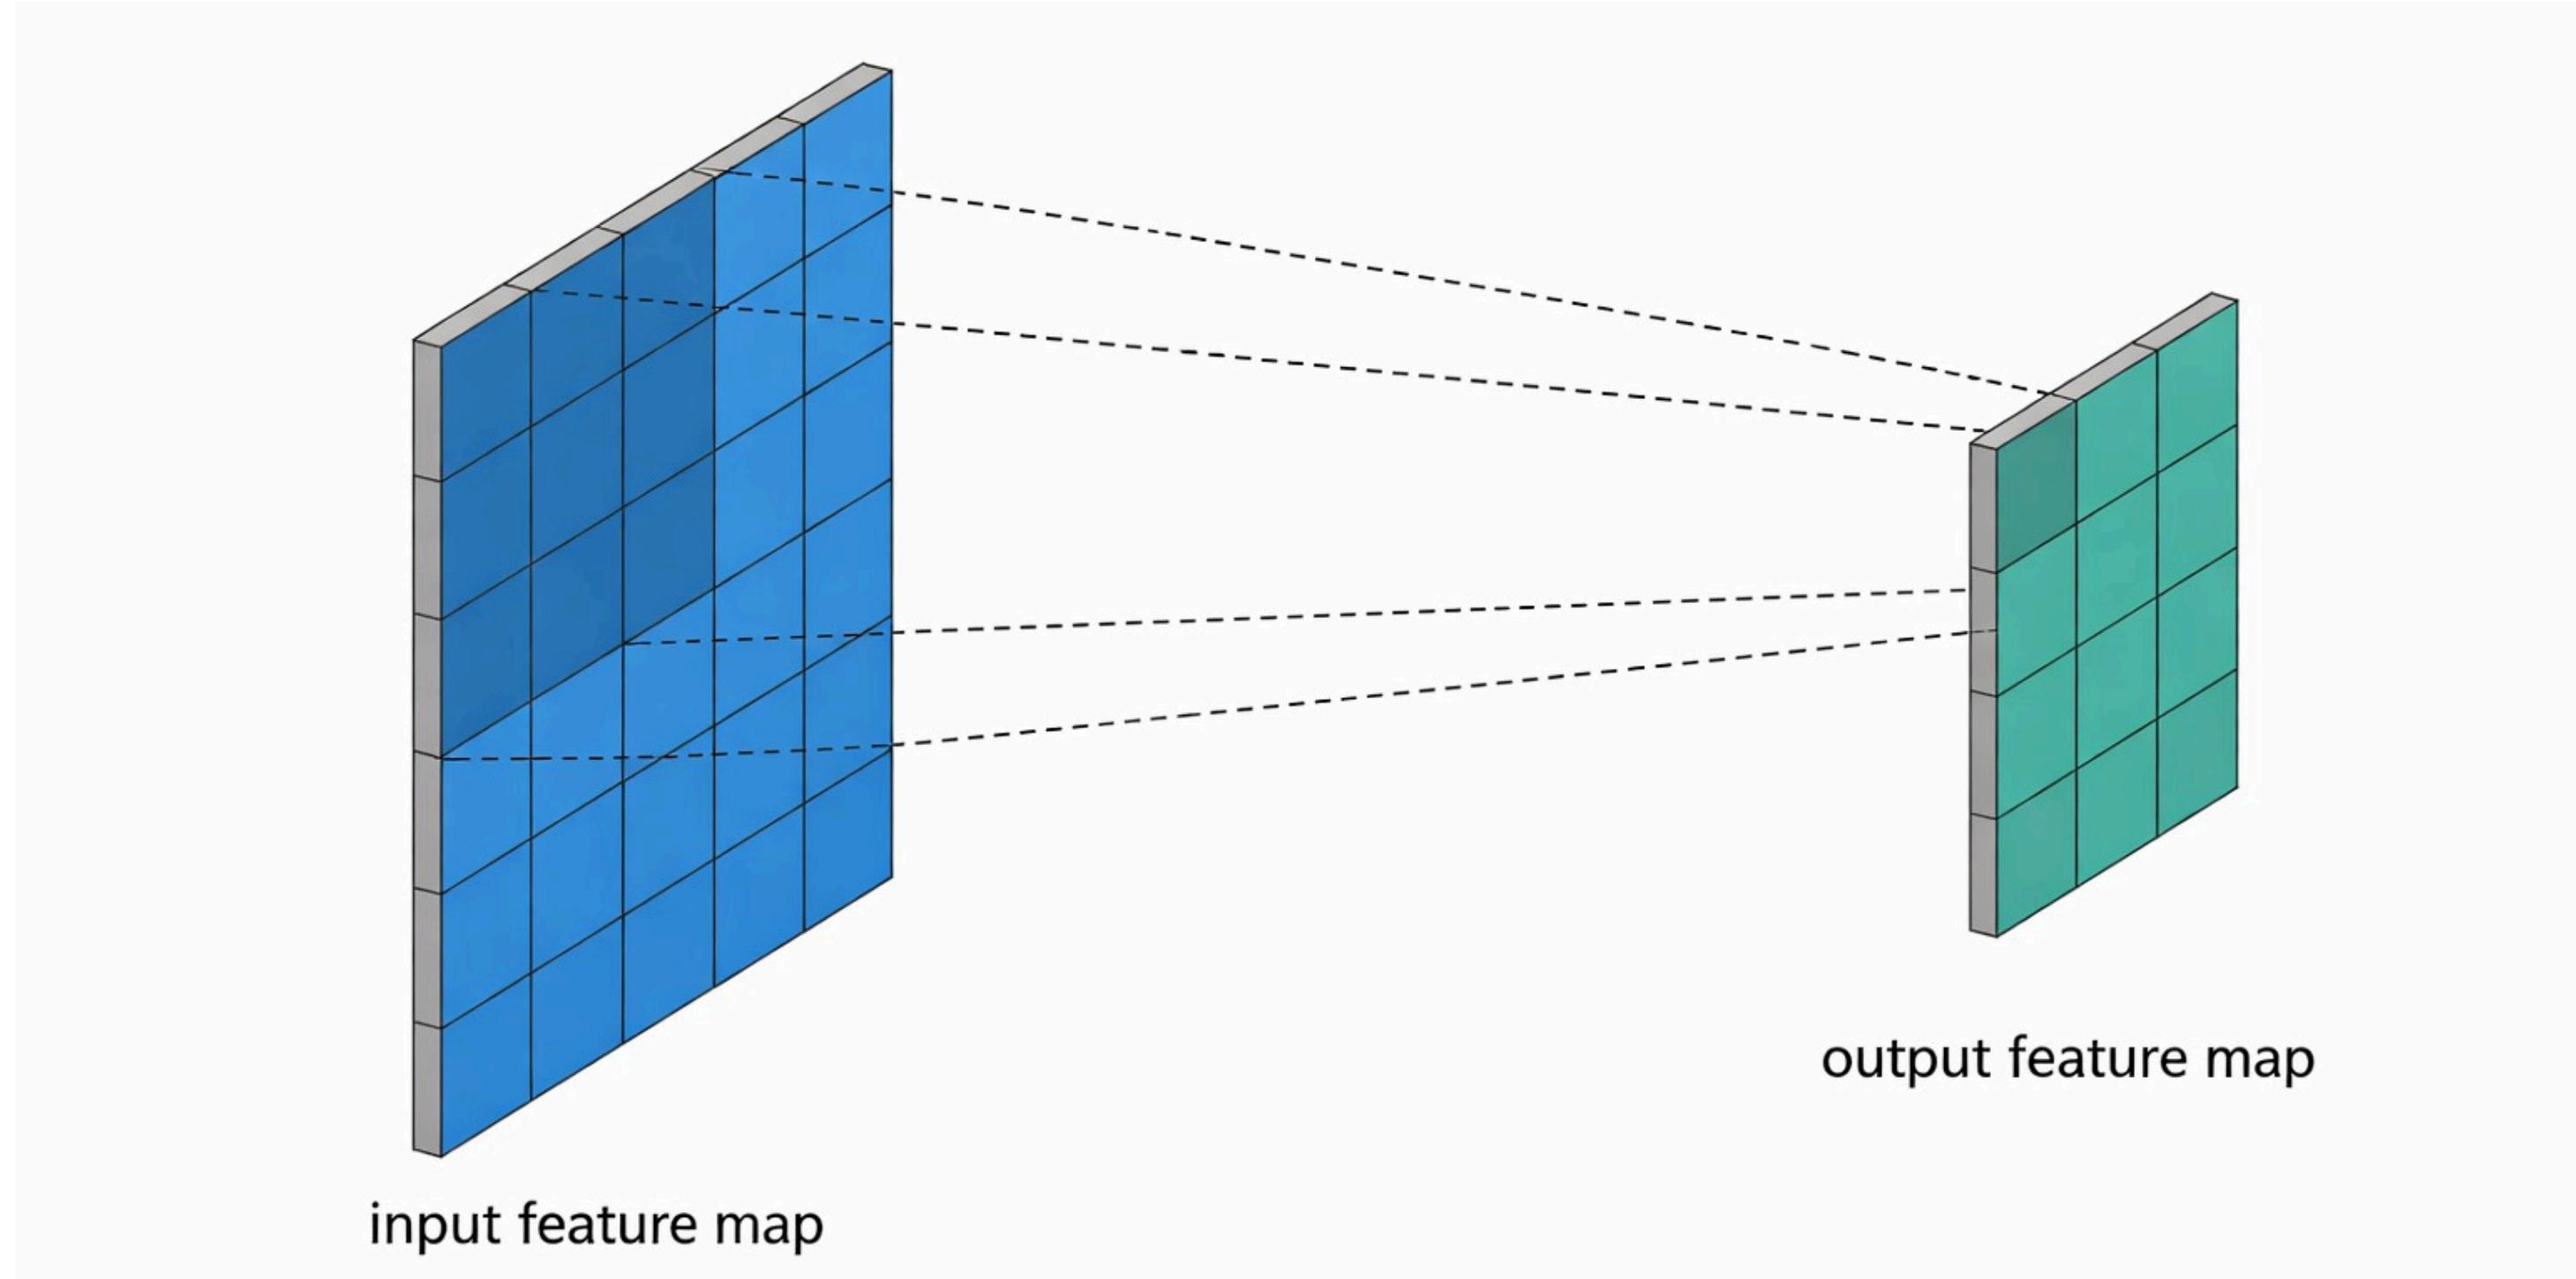
$W_{a,b}$  : weight of the filter at position  $(a, b)$

$X_{i+a, j+b}$  : input value in the local region at  $(i + a, j + b)$

# Convolutional Neural Networks

---

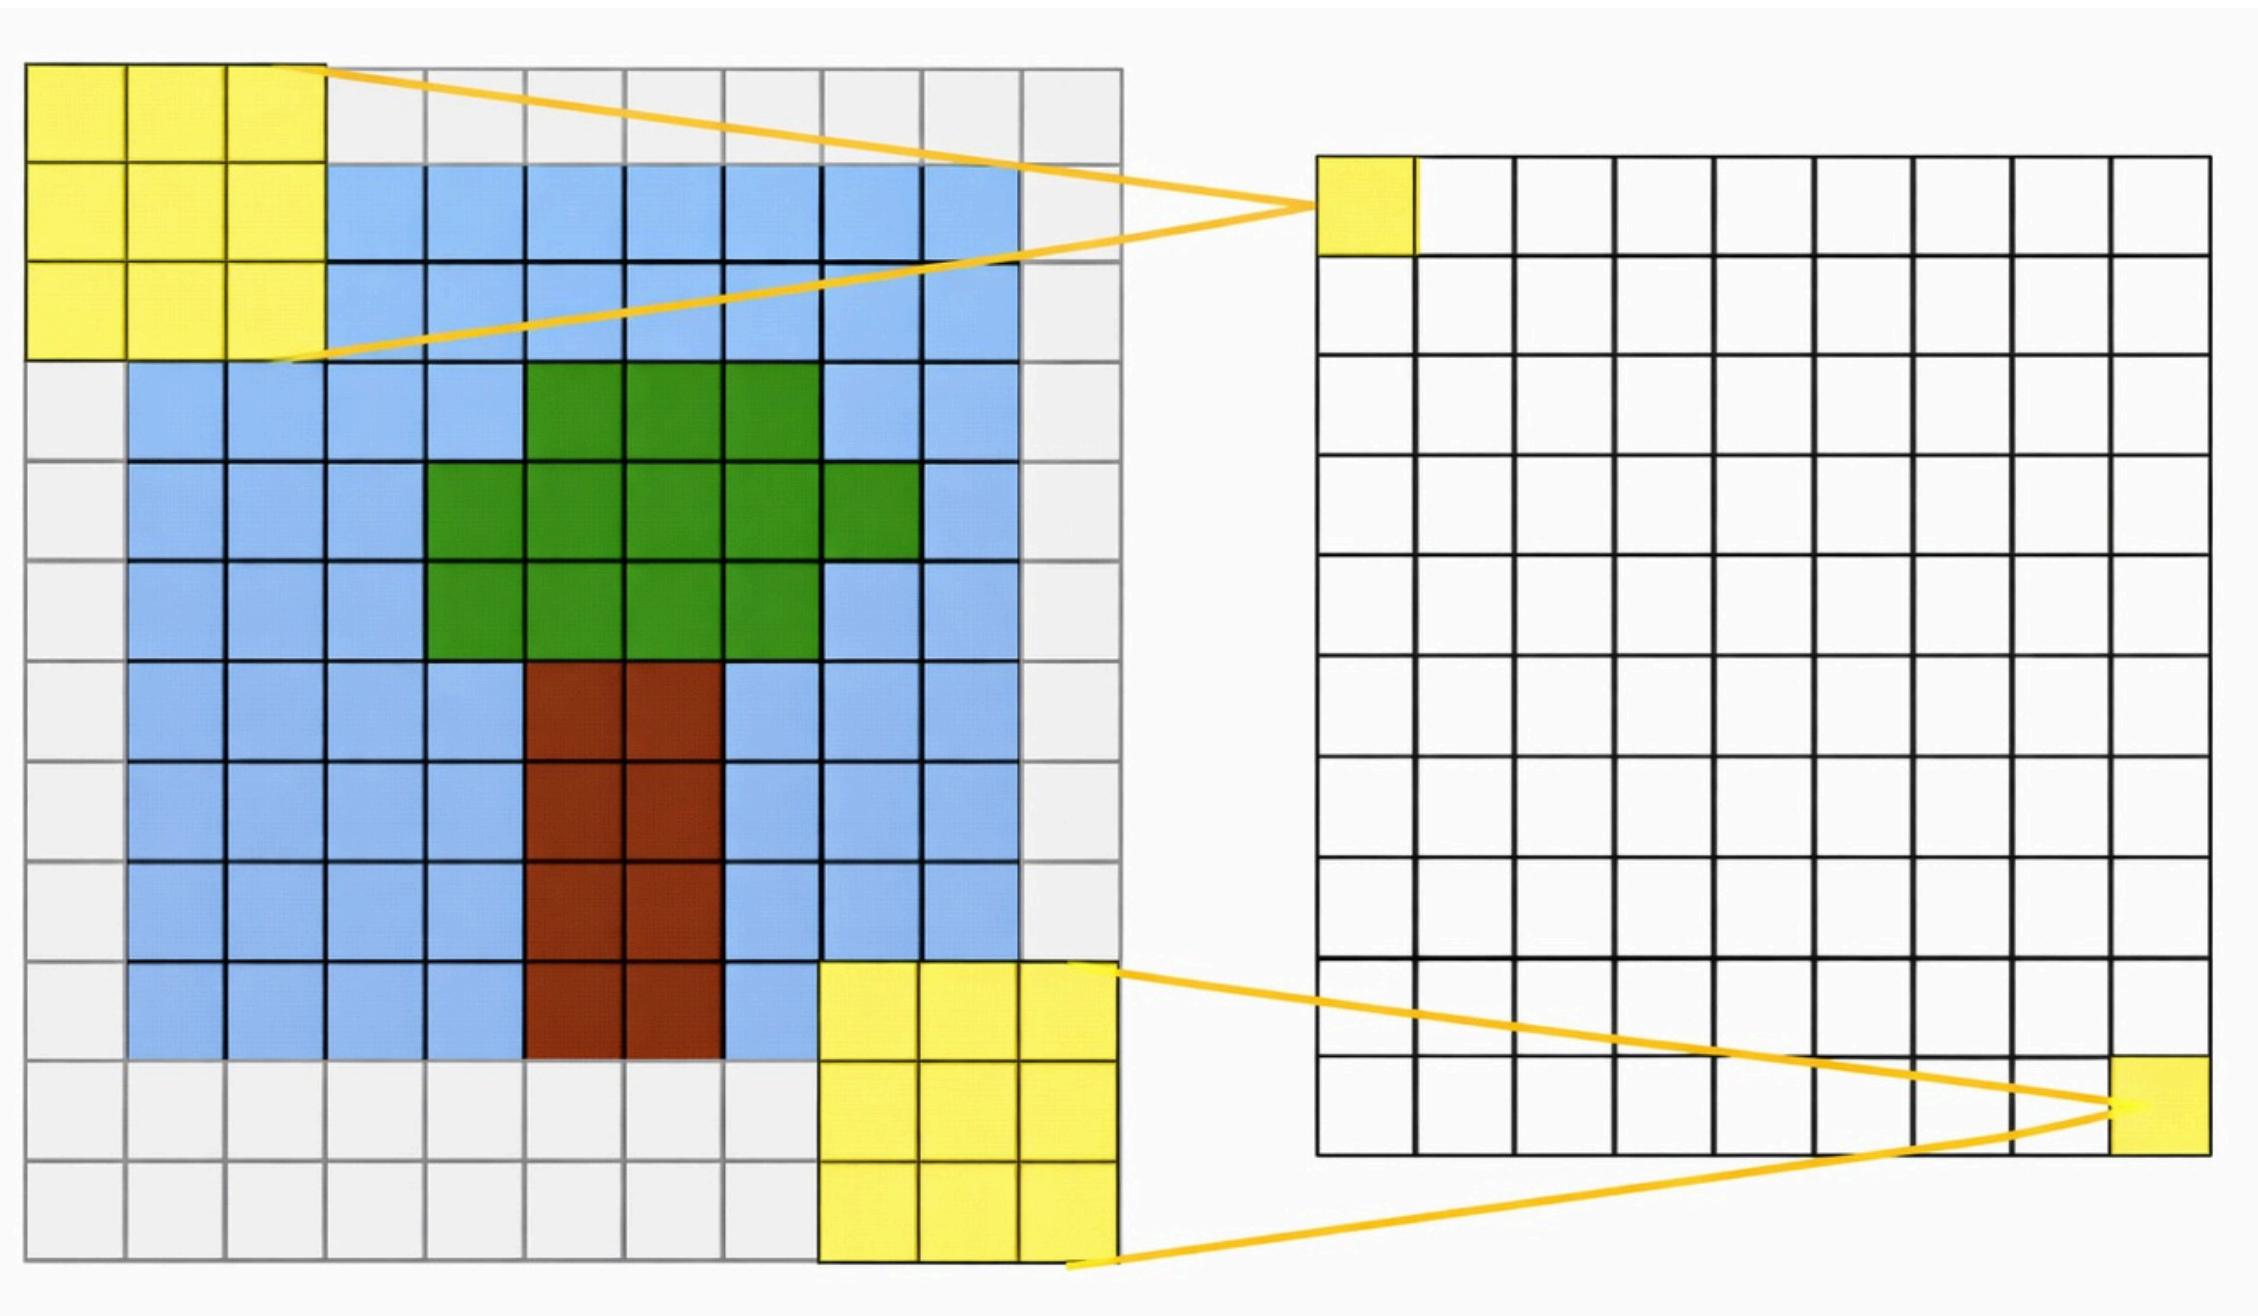
- Applying convolution without padding reduces the feature map size because the kernel cannot fully overlap the input at the borders.



# Convolutional Neural Networks

---

- **Padding** is employed by extending the input image with zeros along its boundaries to control the spatial dimensions of the convolution output.
- *Same Convolution*: Padding is added so the output size equals the input size.

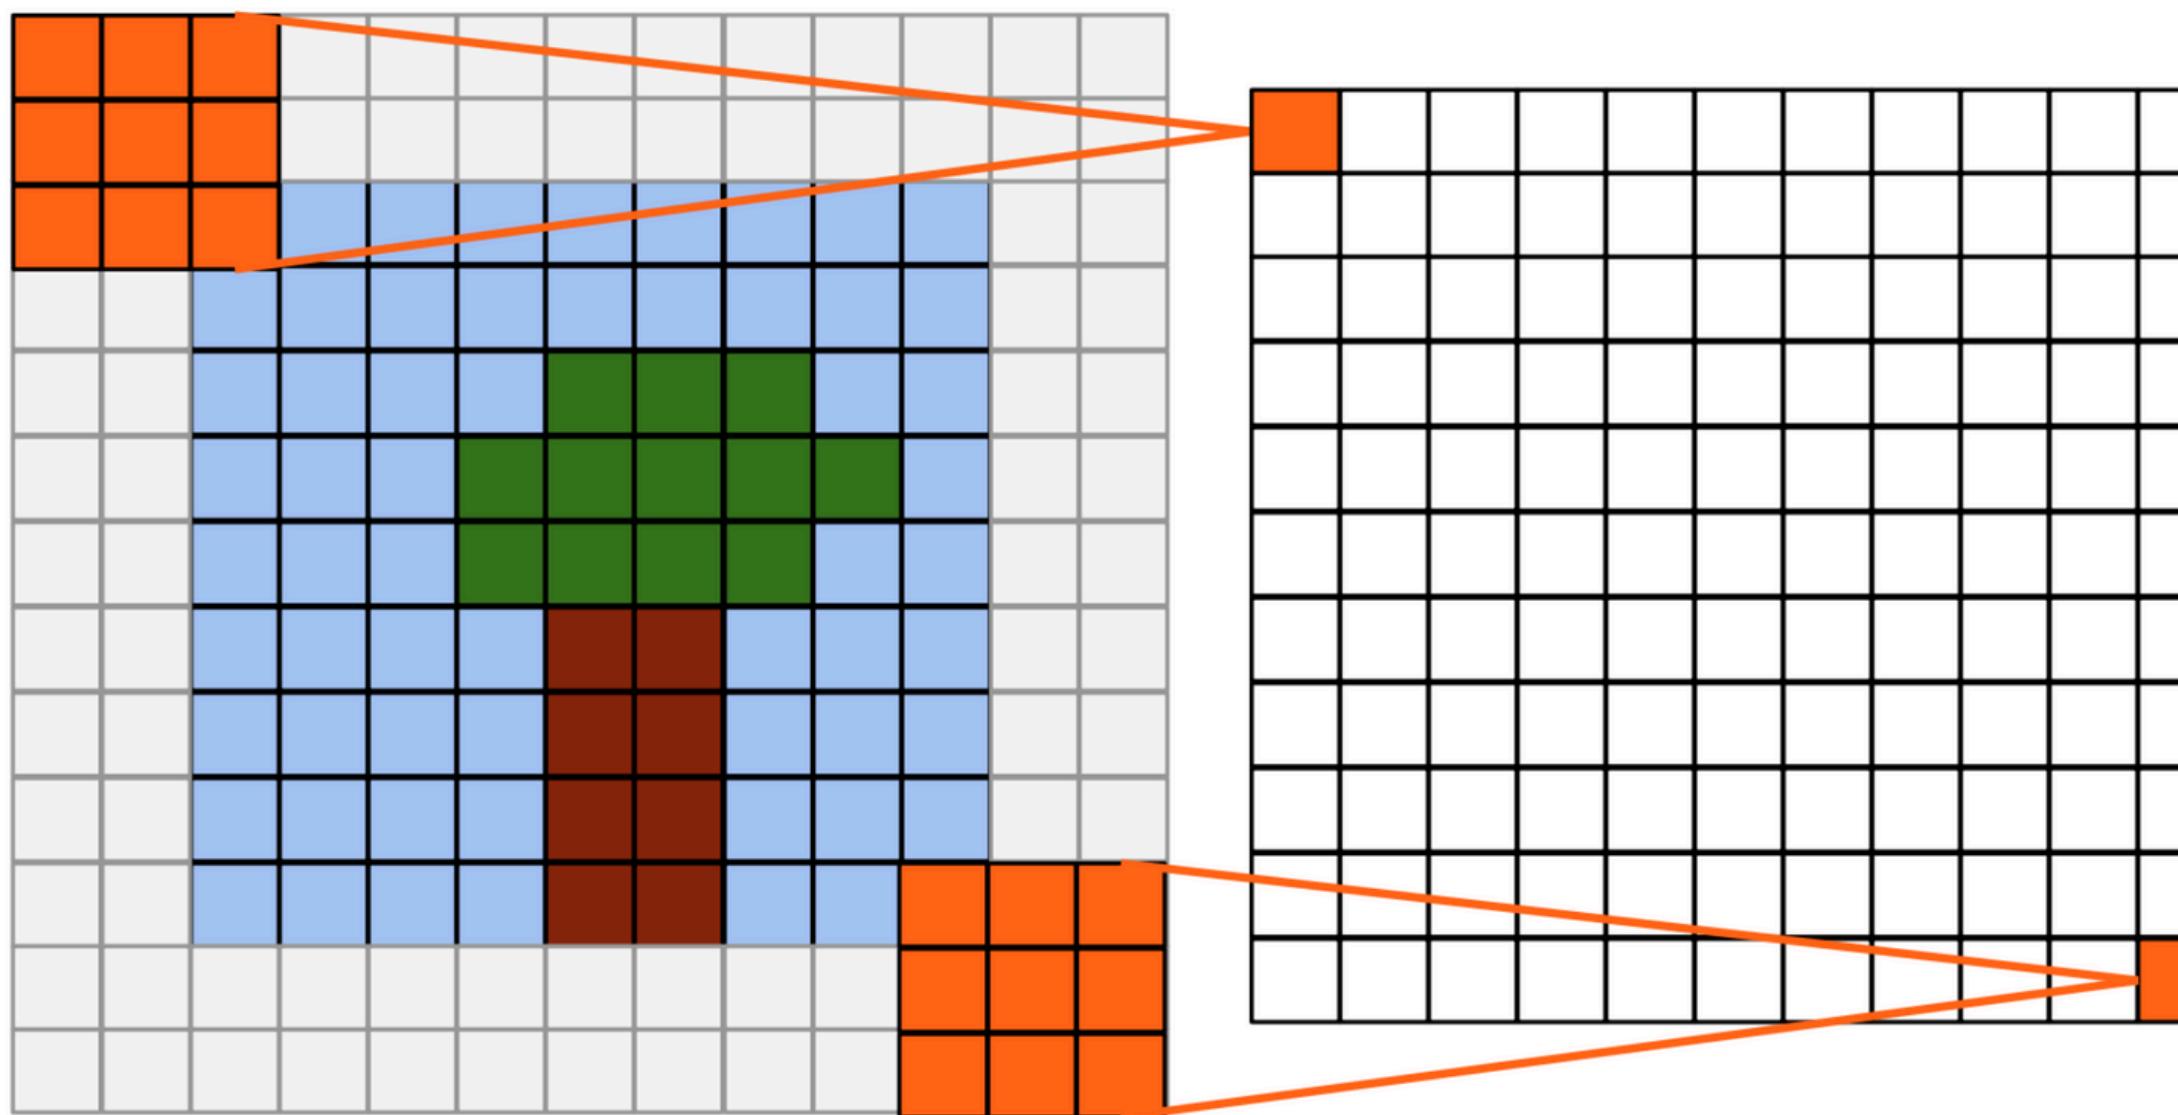


# Convolutional Neural Networks

---

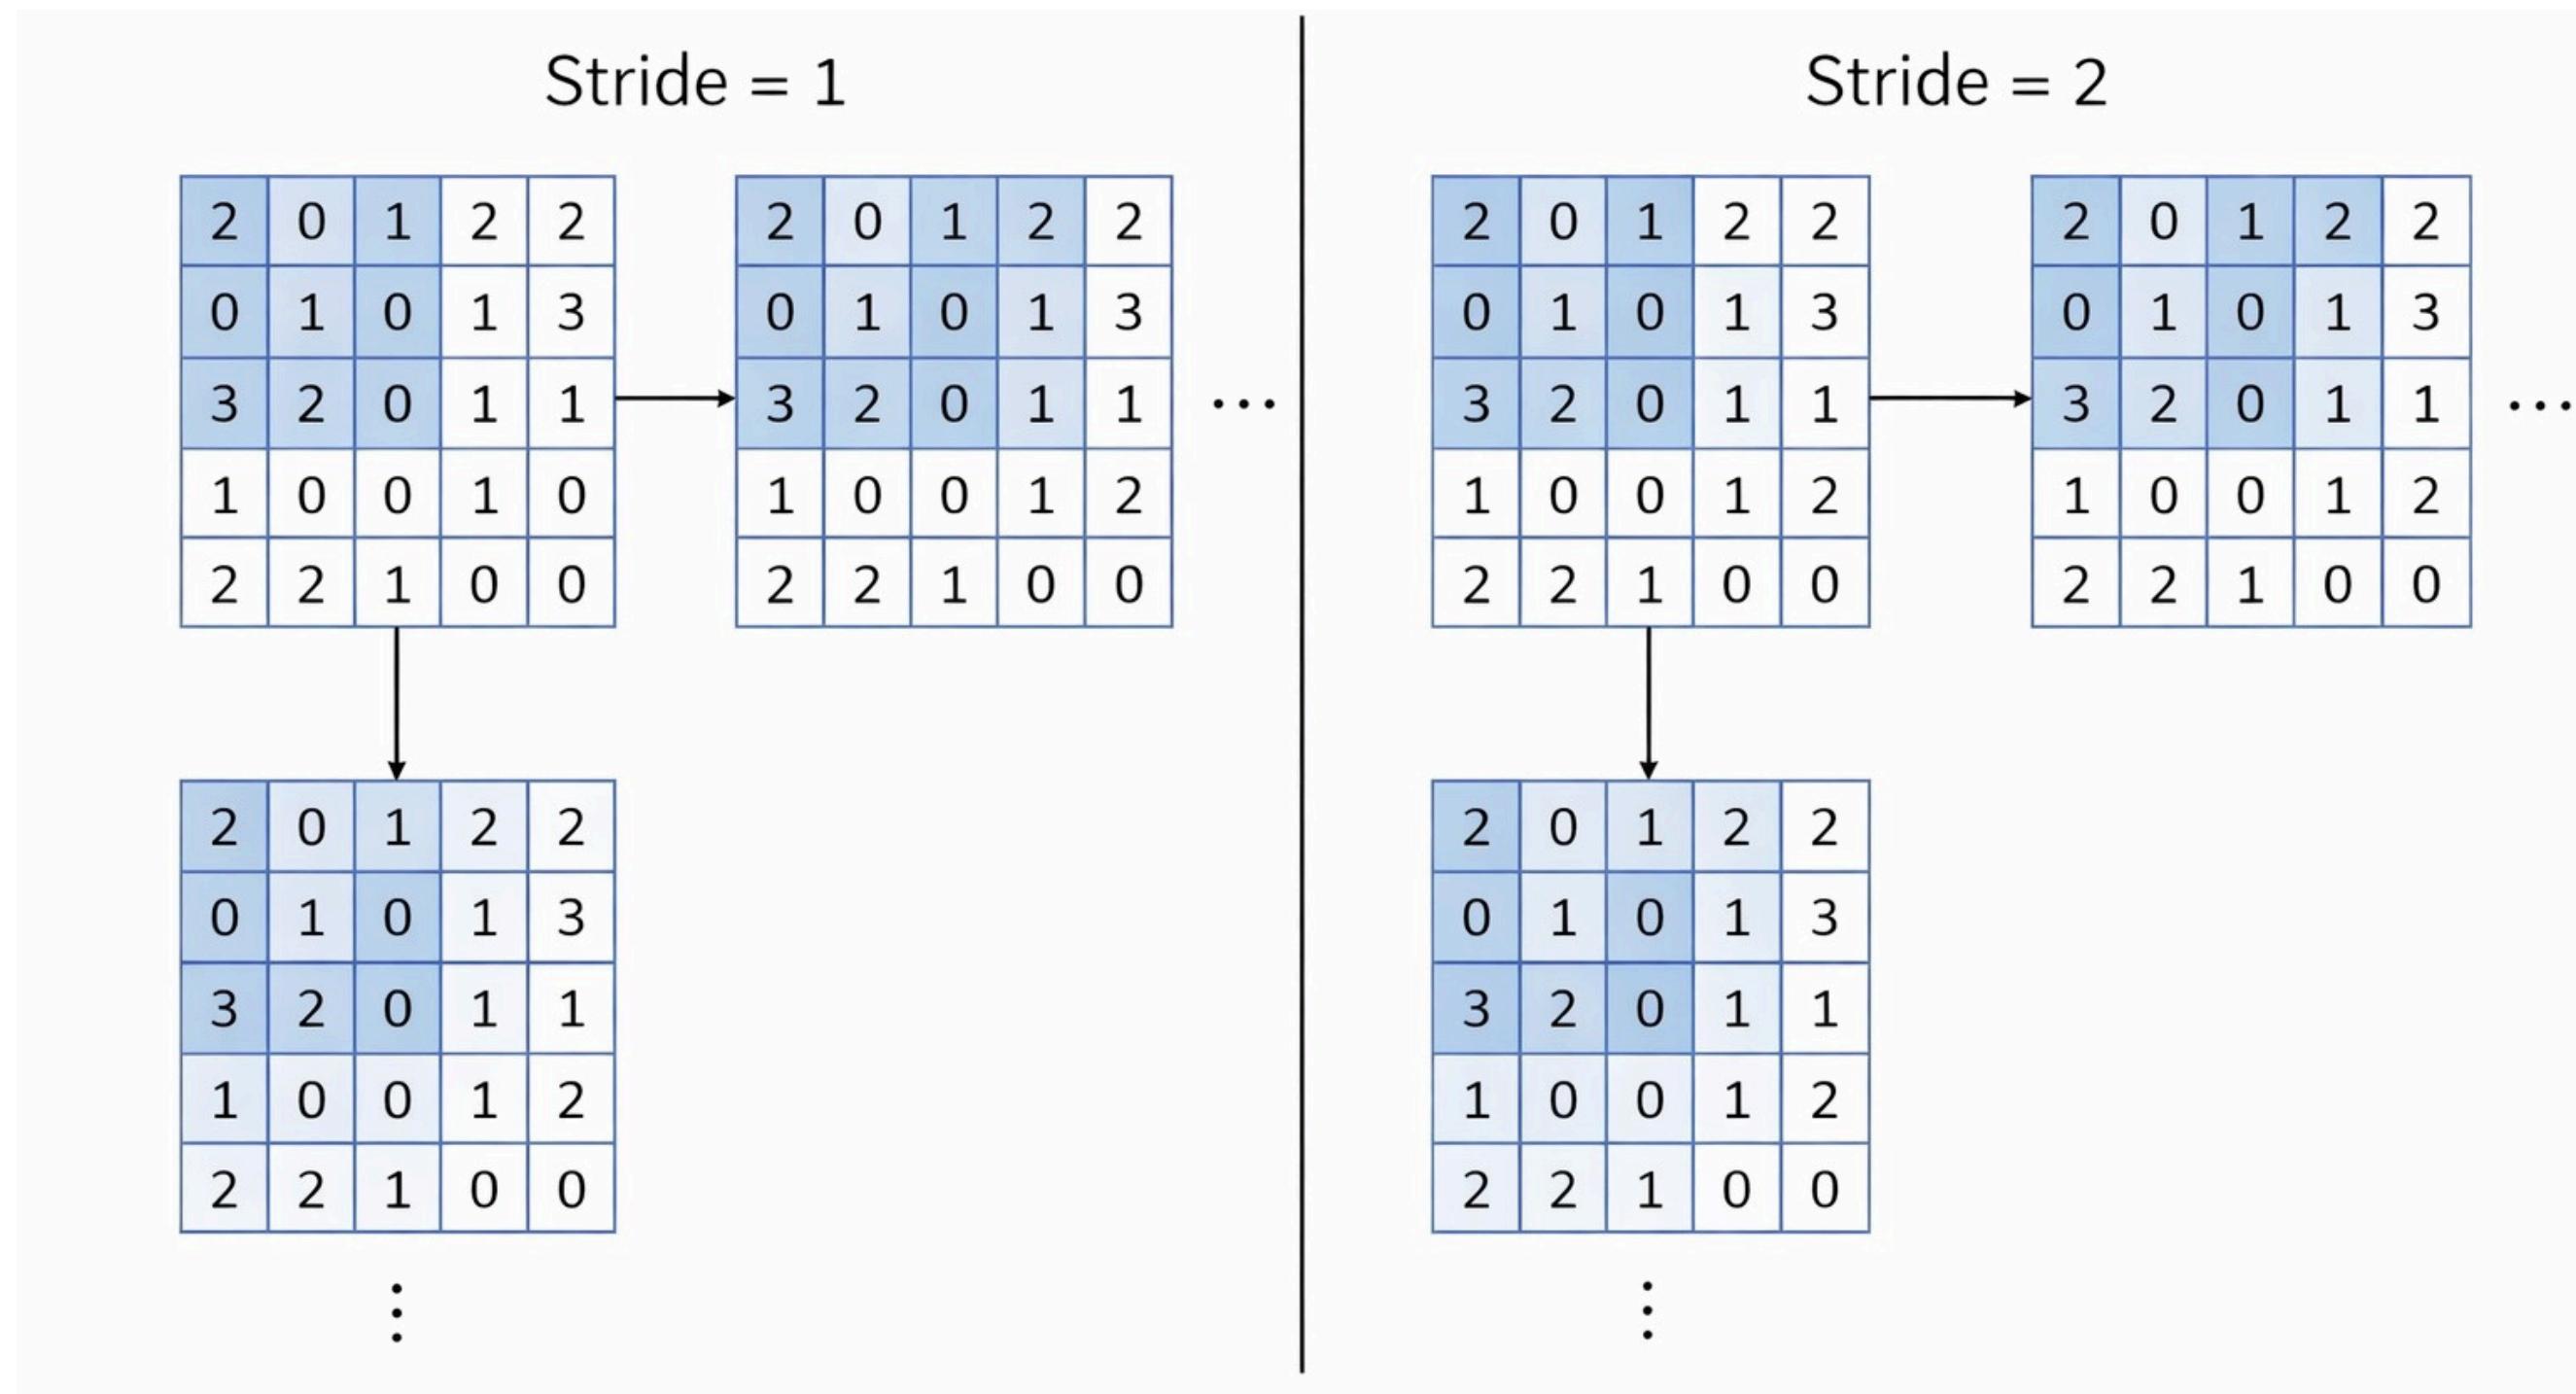
- *Full convolution* adds padding so the kernel can slide over the entire input, including the edges.

$$\text{output size} = \text{input size} + \text{kernel size} - 1$$



# Convolutional Neural Networks

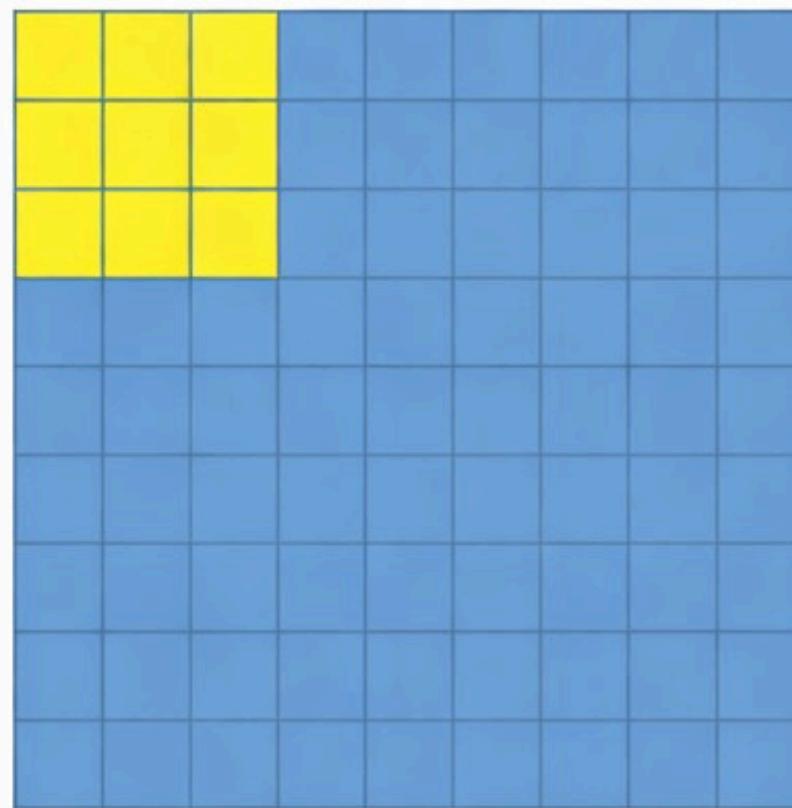
- **Strided Convolution:** Kernel slides along the image with a step  $>1$ 
  - Larger strides reduce output size.



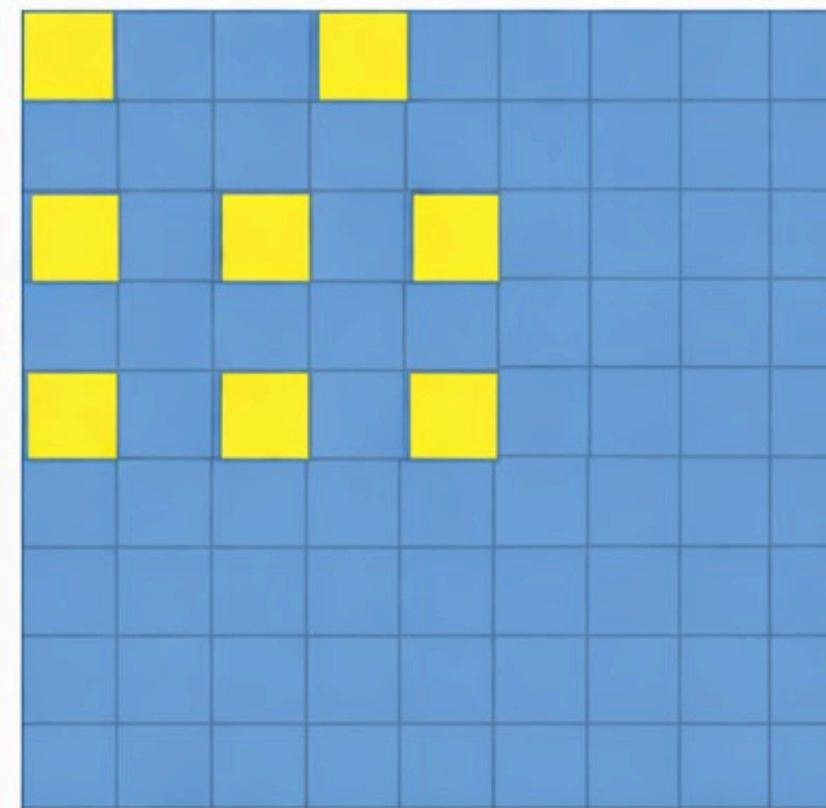
# Convolutional Neural Networks

---

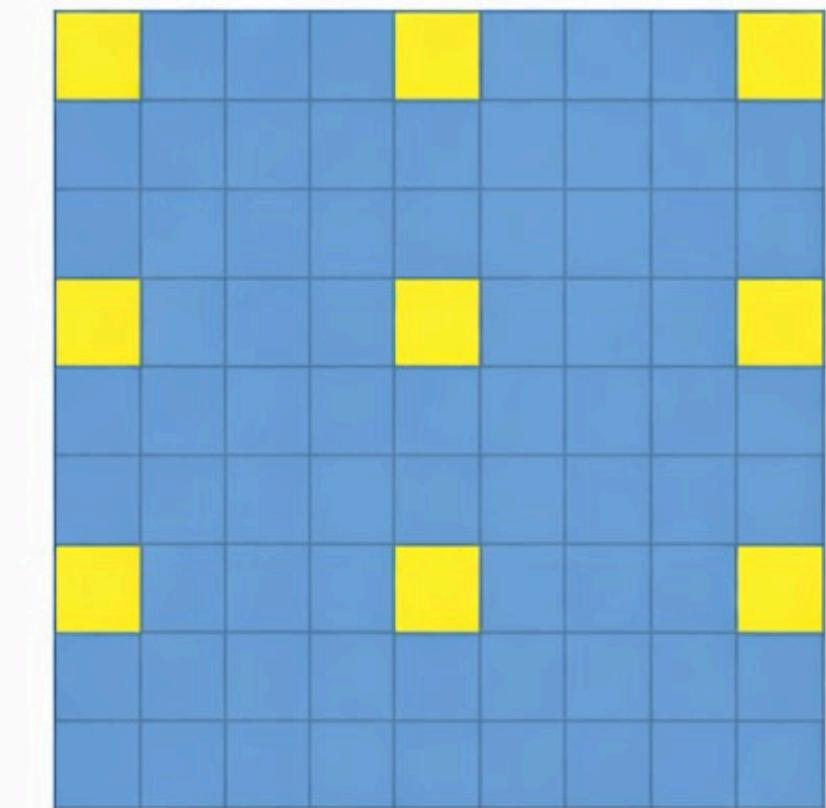
➤ Dilated Convolution: Kernel is spread out, step > 1 between kernel elements.



Kernel size:  $3 \times 3$   
Dilation rate: 1



Kernel size:  $3 \times 3$   
Dilation rate: 2



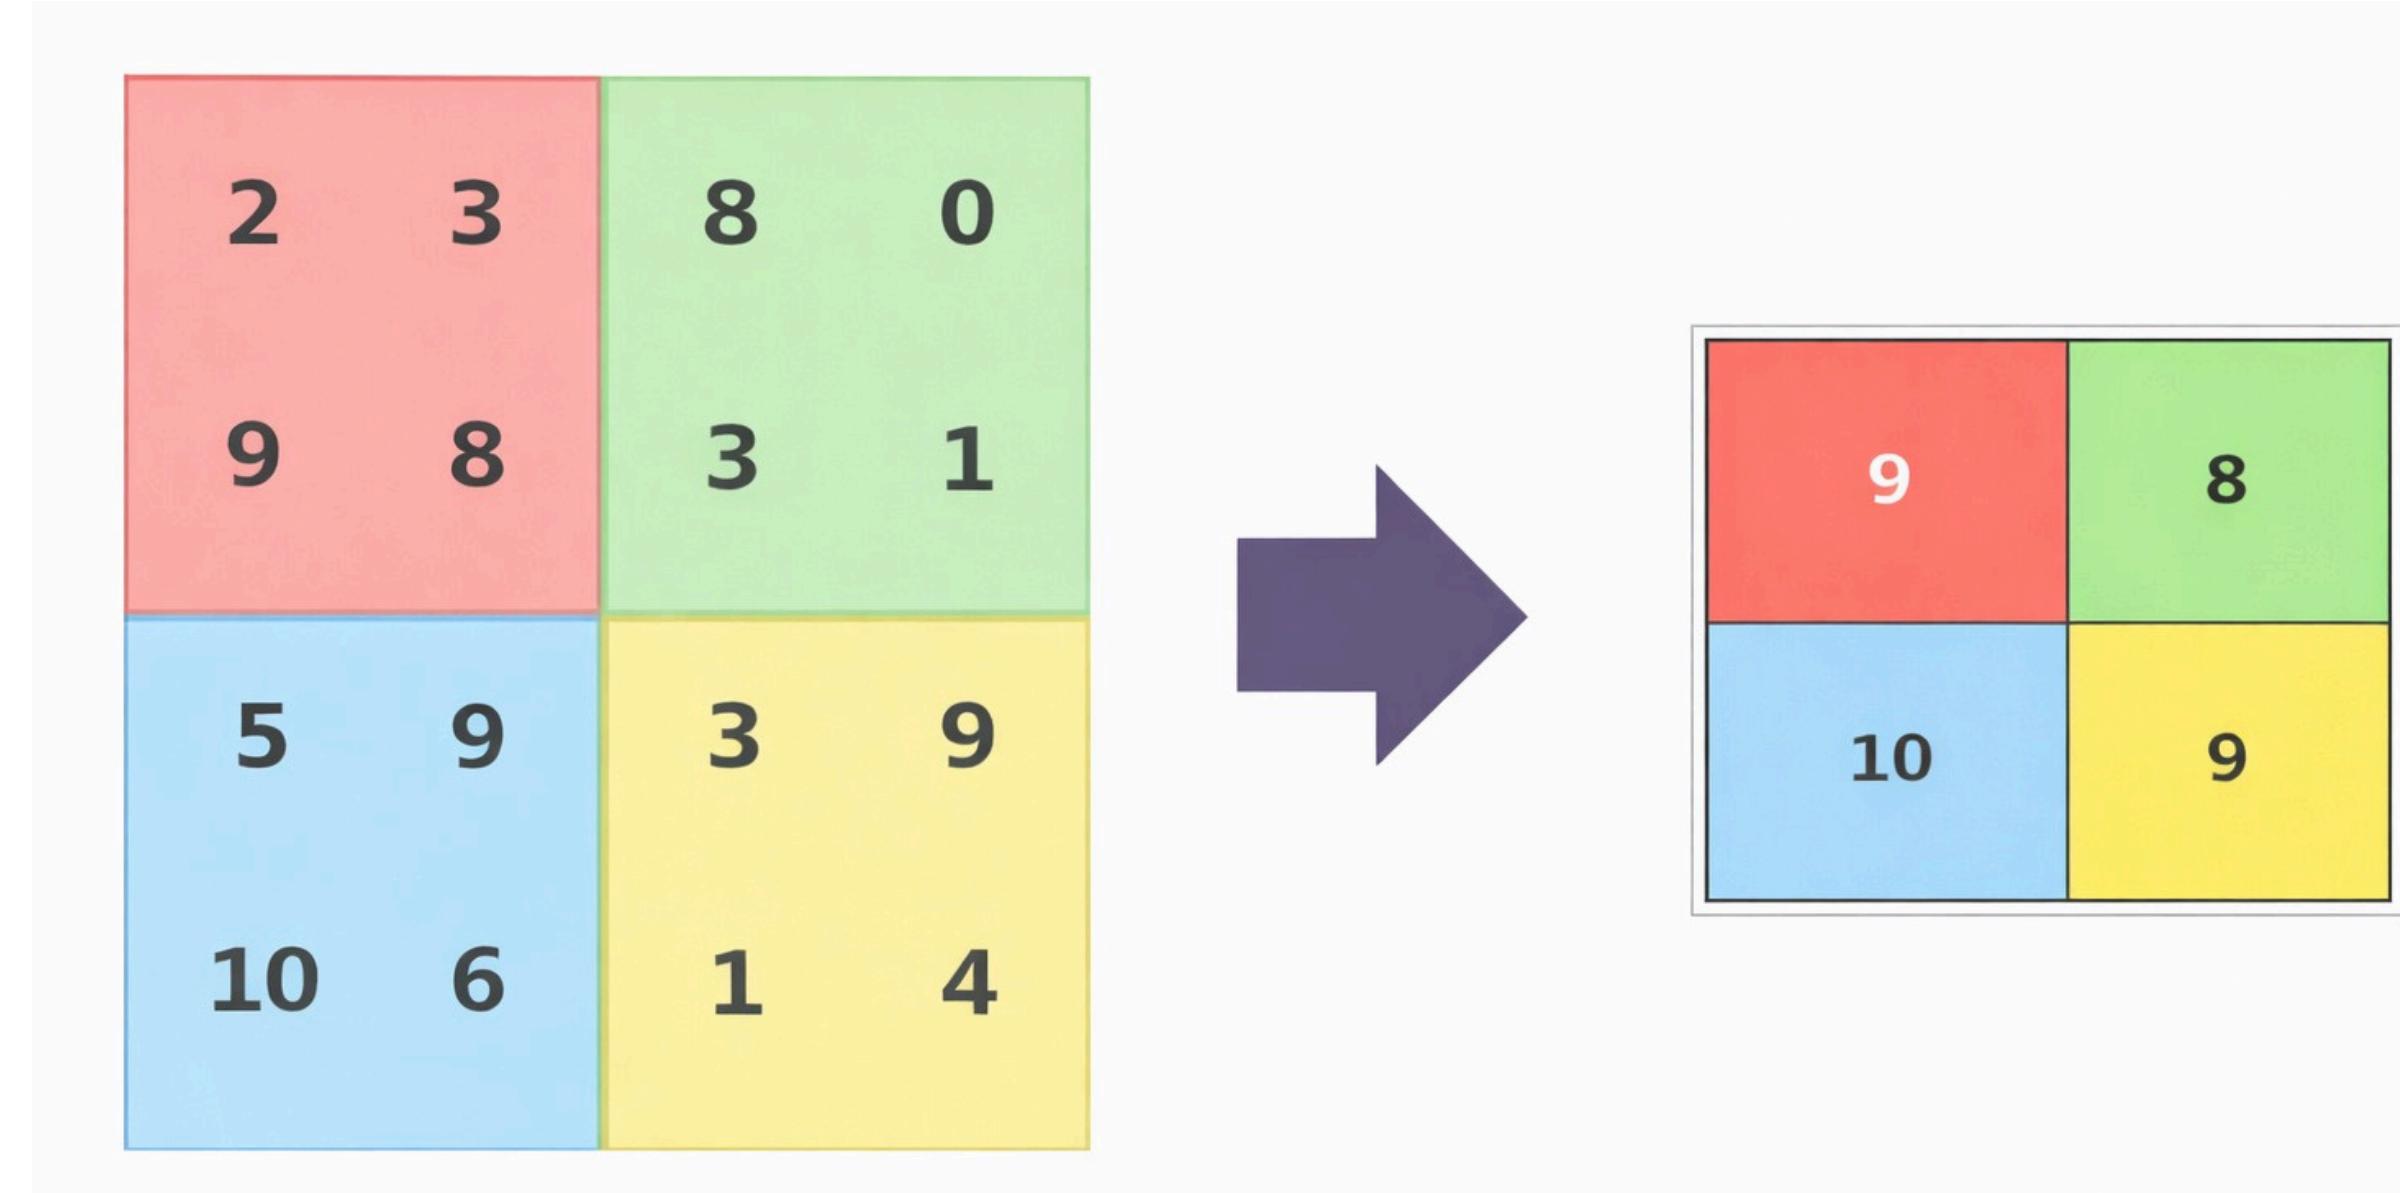
Kernel size:  $3 \times 3$   
Dilation rate: 3

# Pooling Layer

# Pooling Layer

- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

- Max Pooling

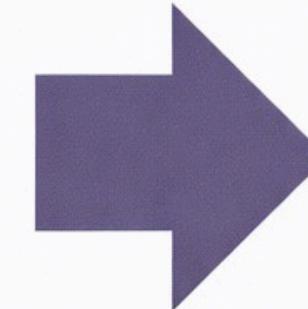


# Pooling Layer

- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

- Average Pooling

2	3	8	0
9	8	3	1
5	9	3	9
10	6	1	4



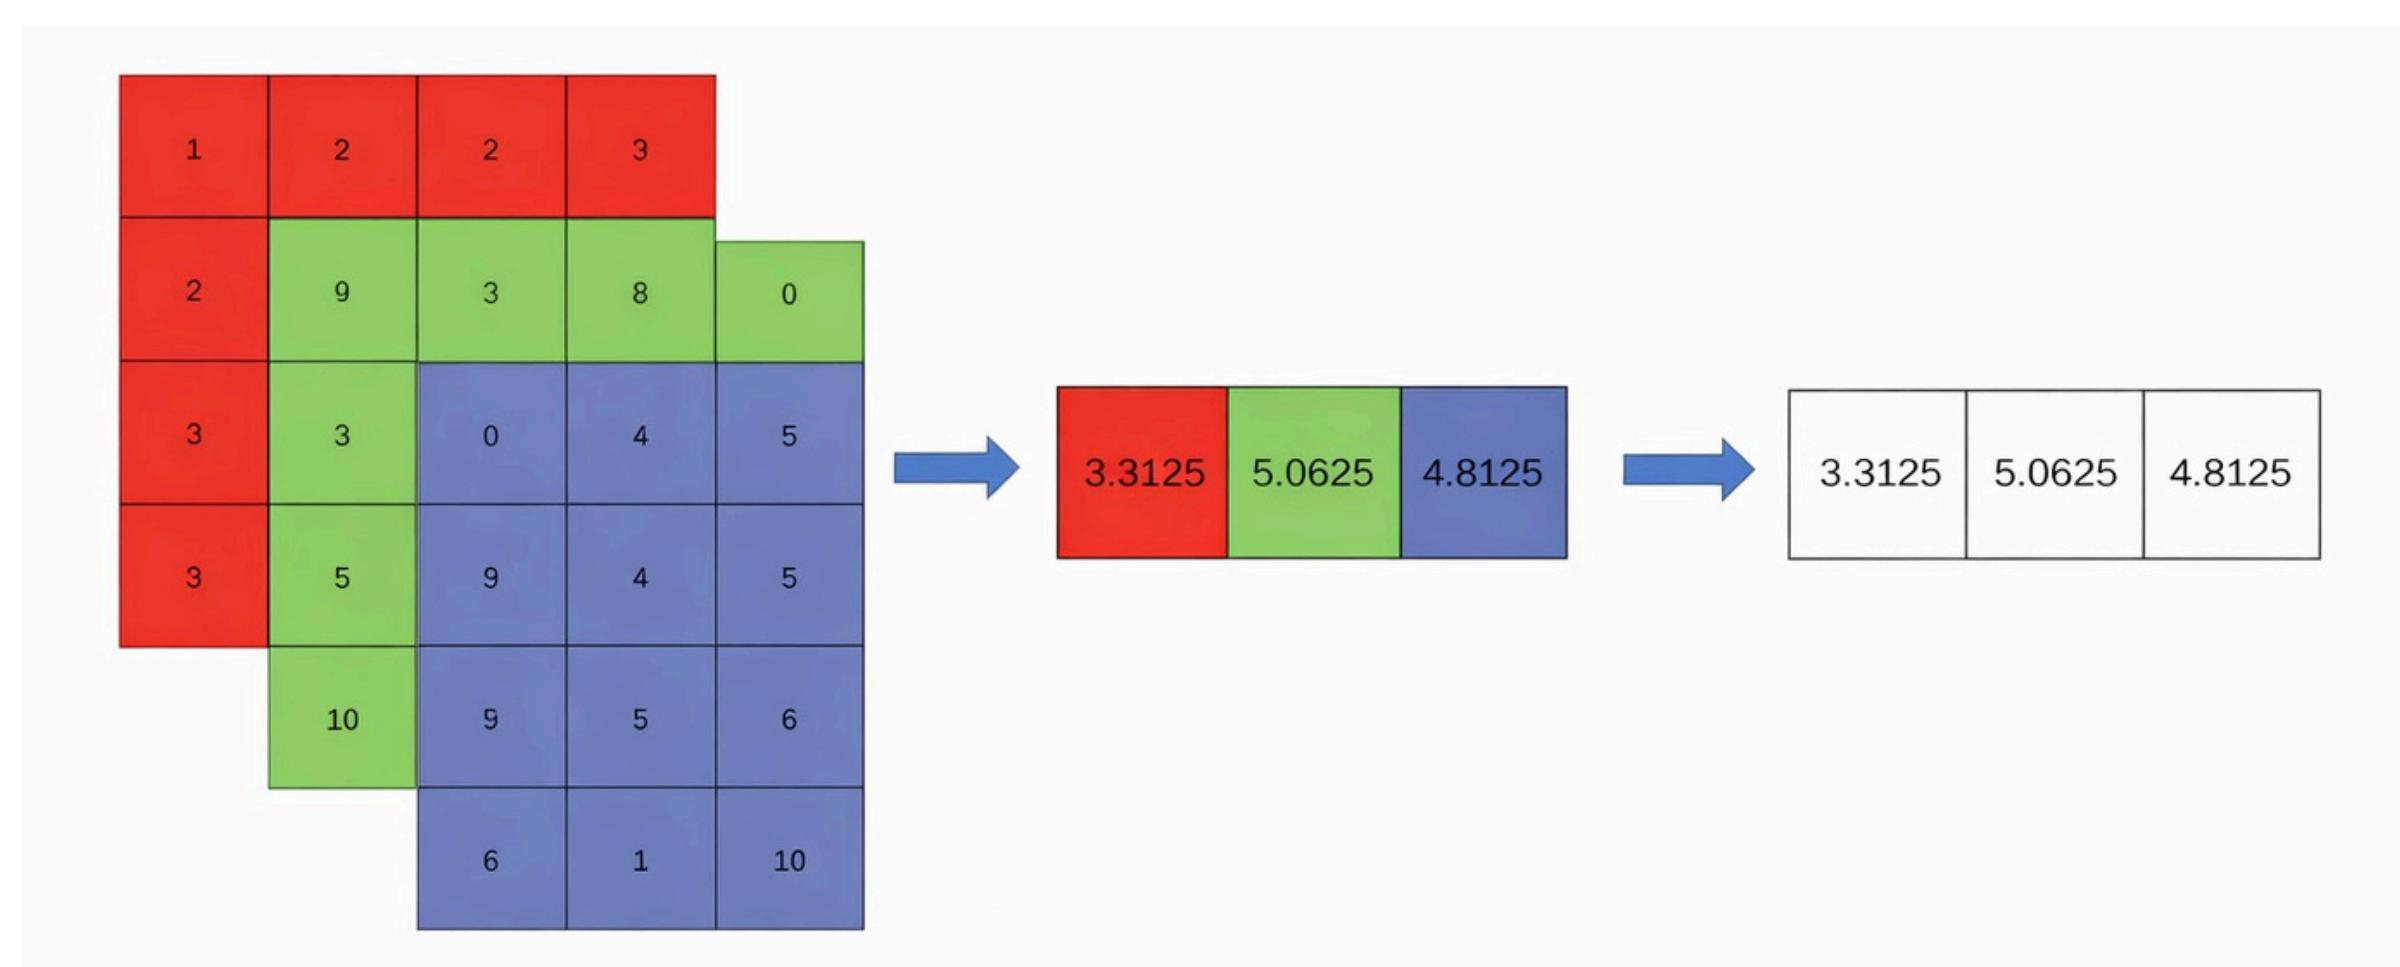
5.5	3
7.5	4.25

# Pooling Layer

---

- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

- Global Average Pooling

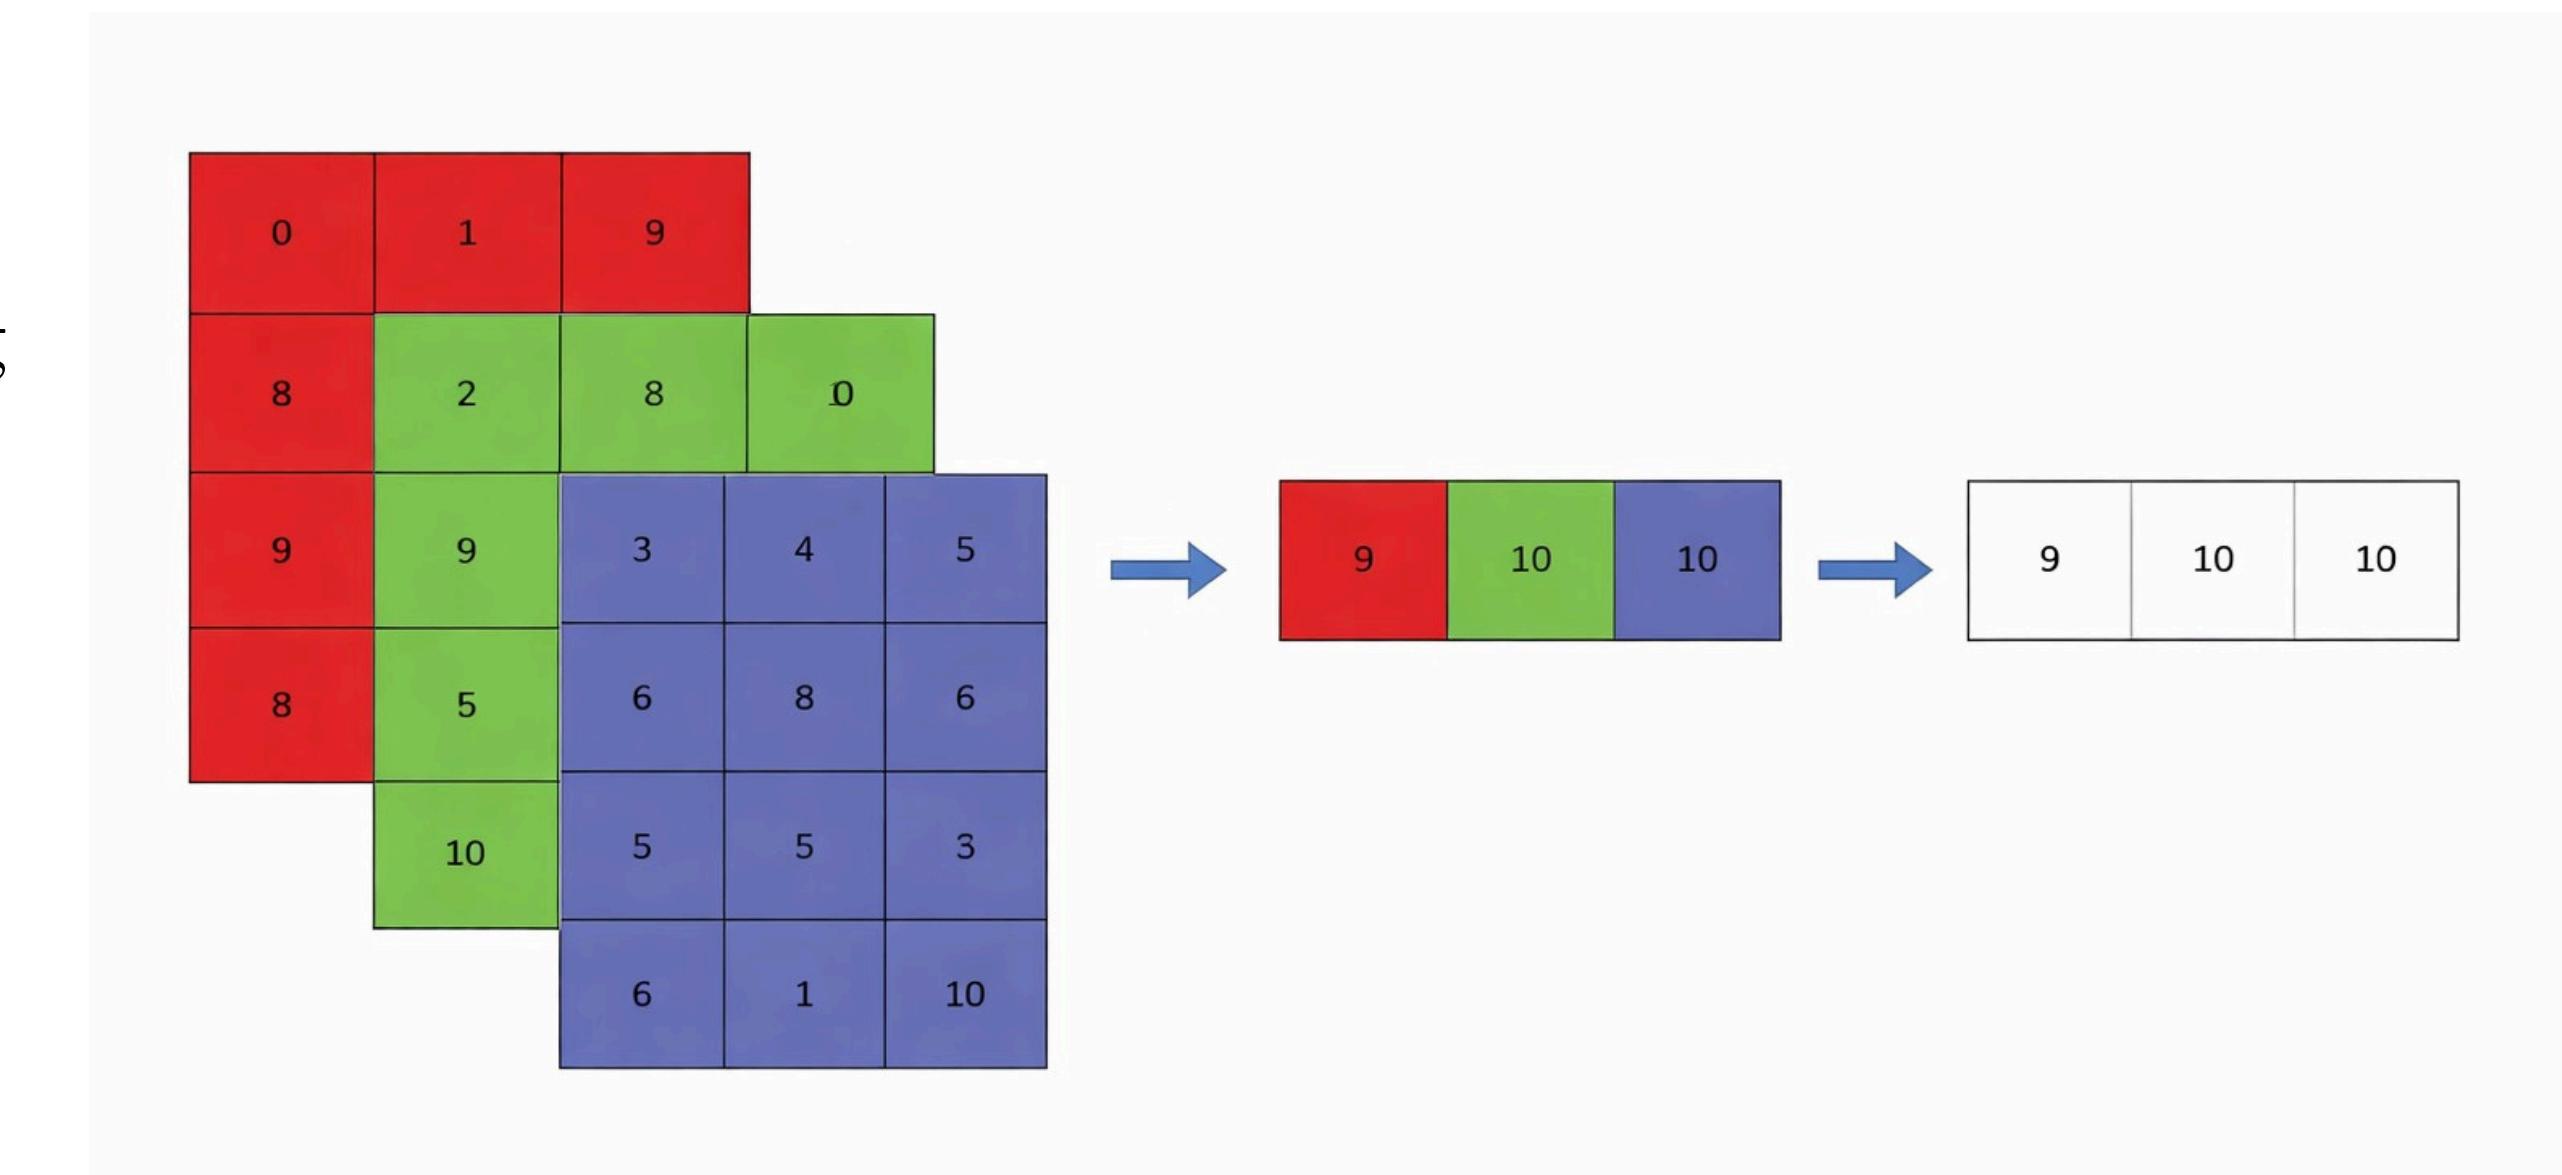


# Pooling Layer

---

- Pooling is used for dimensionality reduction in CNNs.  
It reduces the spatial size of feature maps while keeping important information.

## ➤ Global Max Pooling



# Convolutional Neural Networks

---

## ➤ CNN Output Size

$$n_{\text{out}} = \left\lfloor \frac{n_{\text{in}} + 2p - k}{s} \right\rfloor + 1$$

$n_{\text{in}}$  : number of input features

$n_{\text{out}}$  : number of output features

$k$  : convolution kernel size

$p$  : convolution padding size

$s$  : convolution stride

# Convolutional Neural Networks

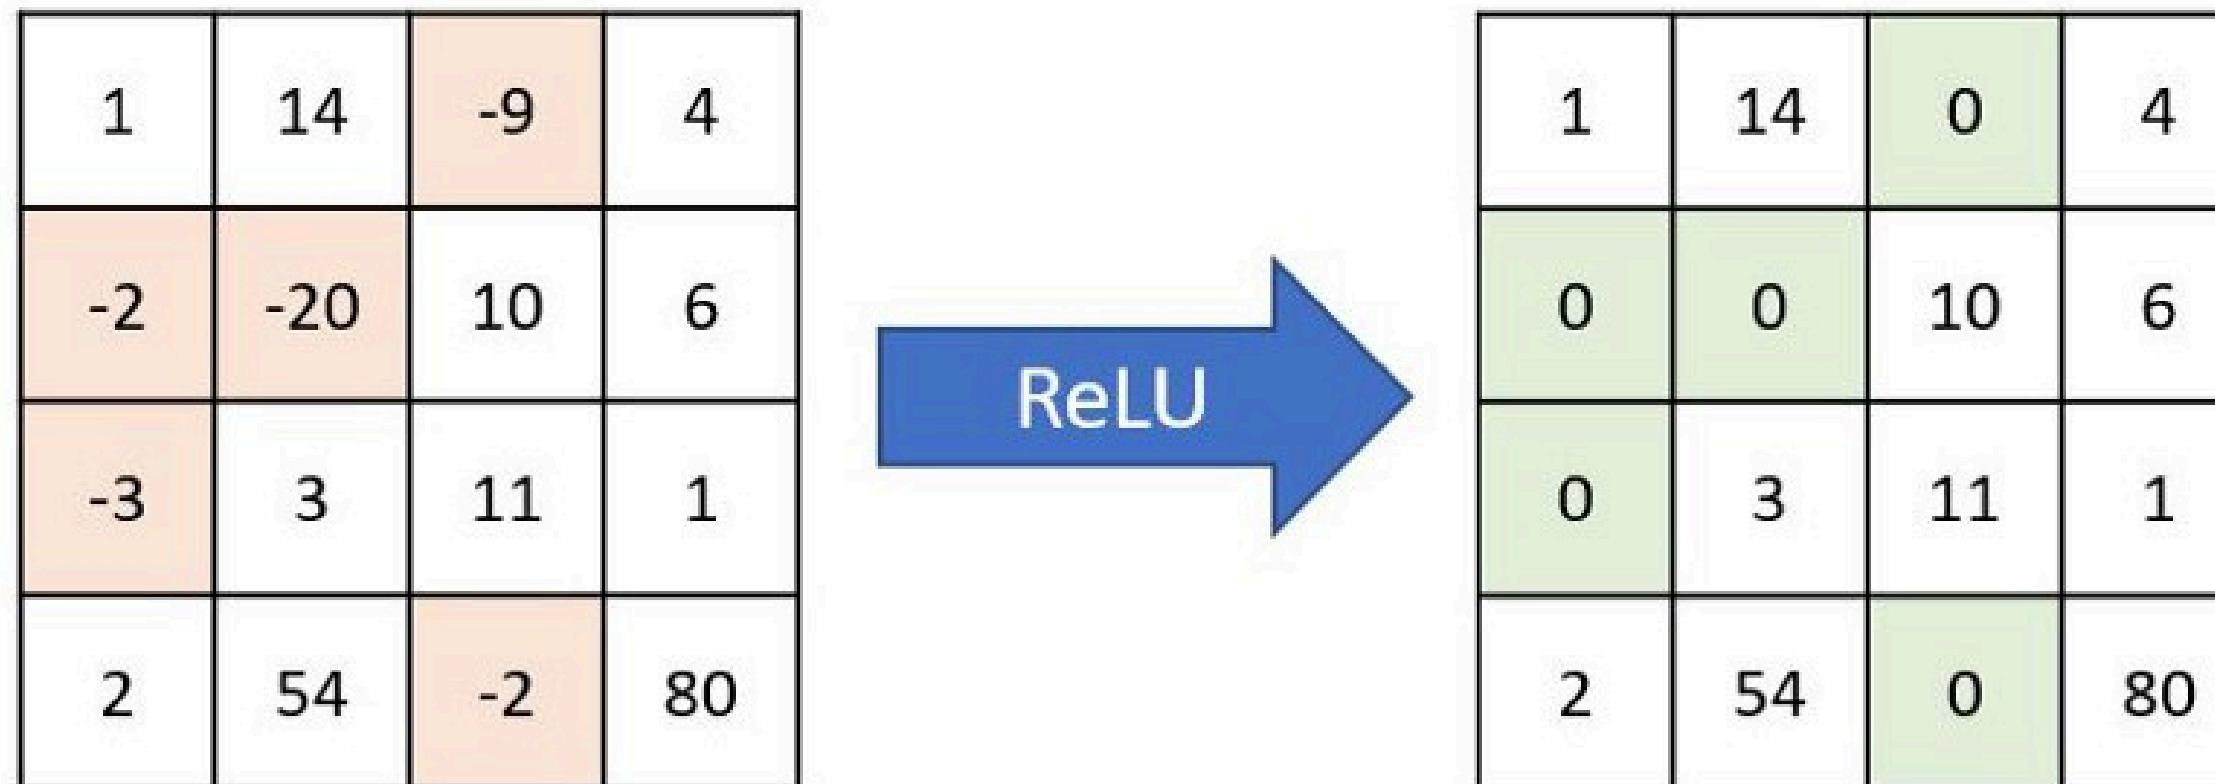
---

## > Activation

Just like Fully-Connected Neural Networks, we can apply an activation over the convolutional layer outputs

For example, the Rectified Linear Unit (ReLU) activation function is defined as

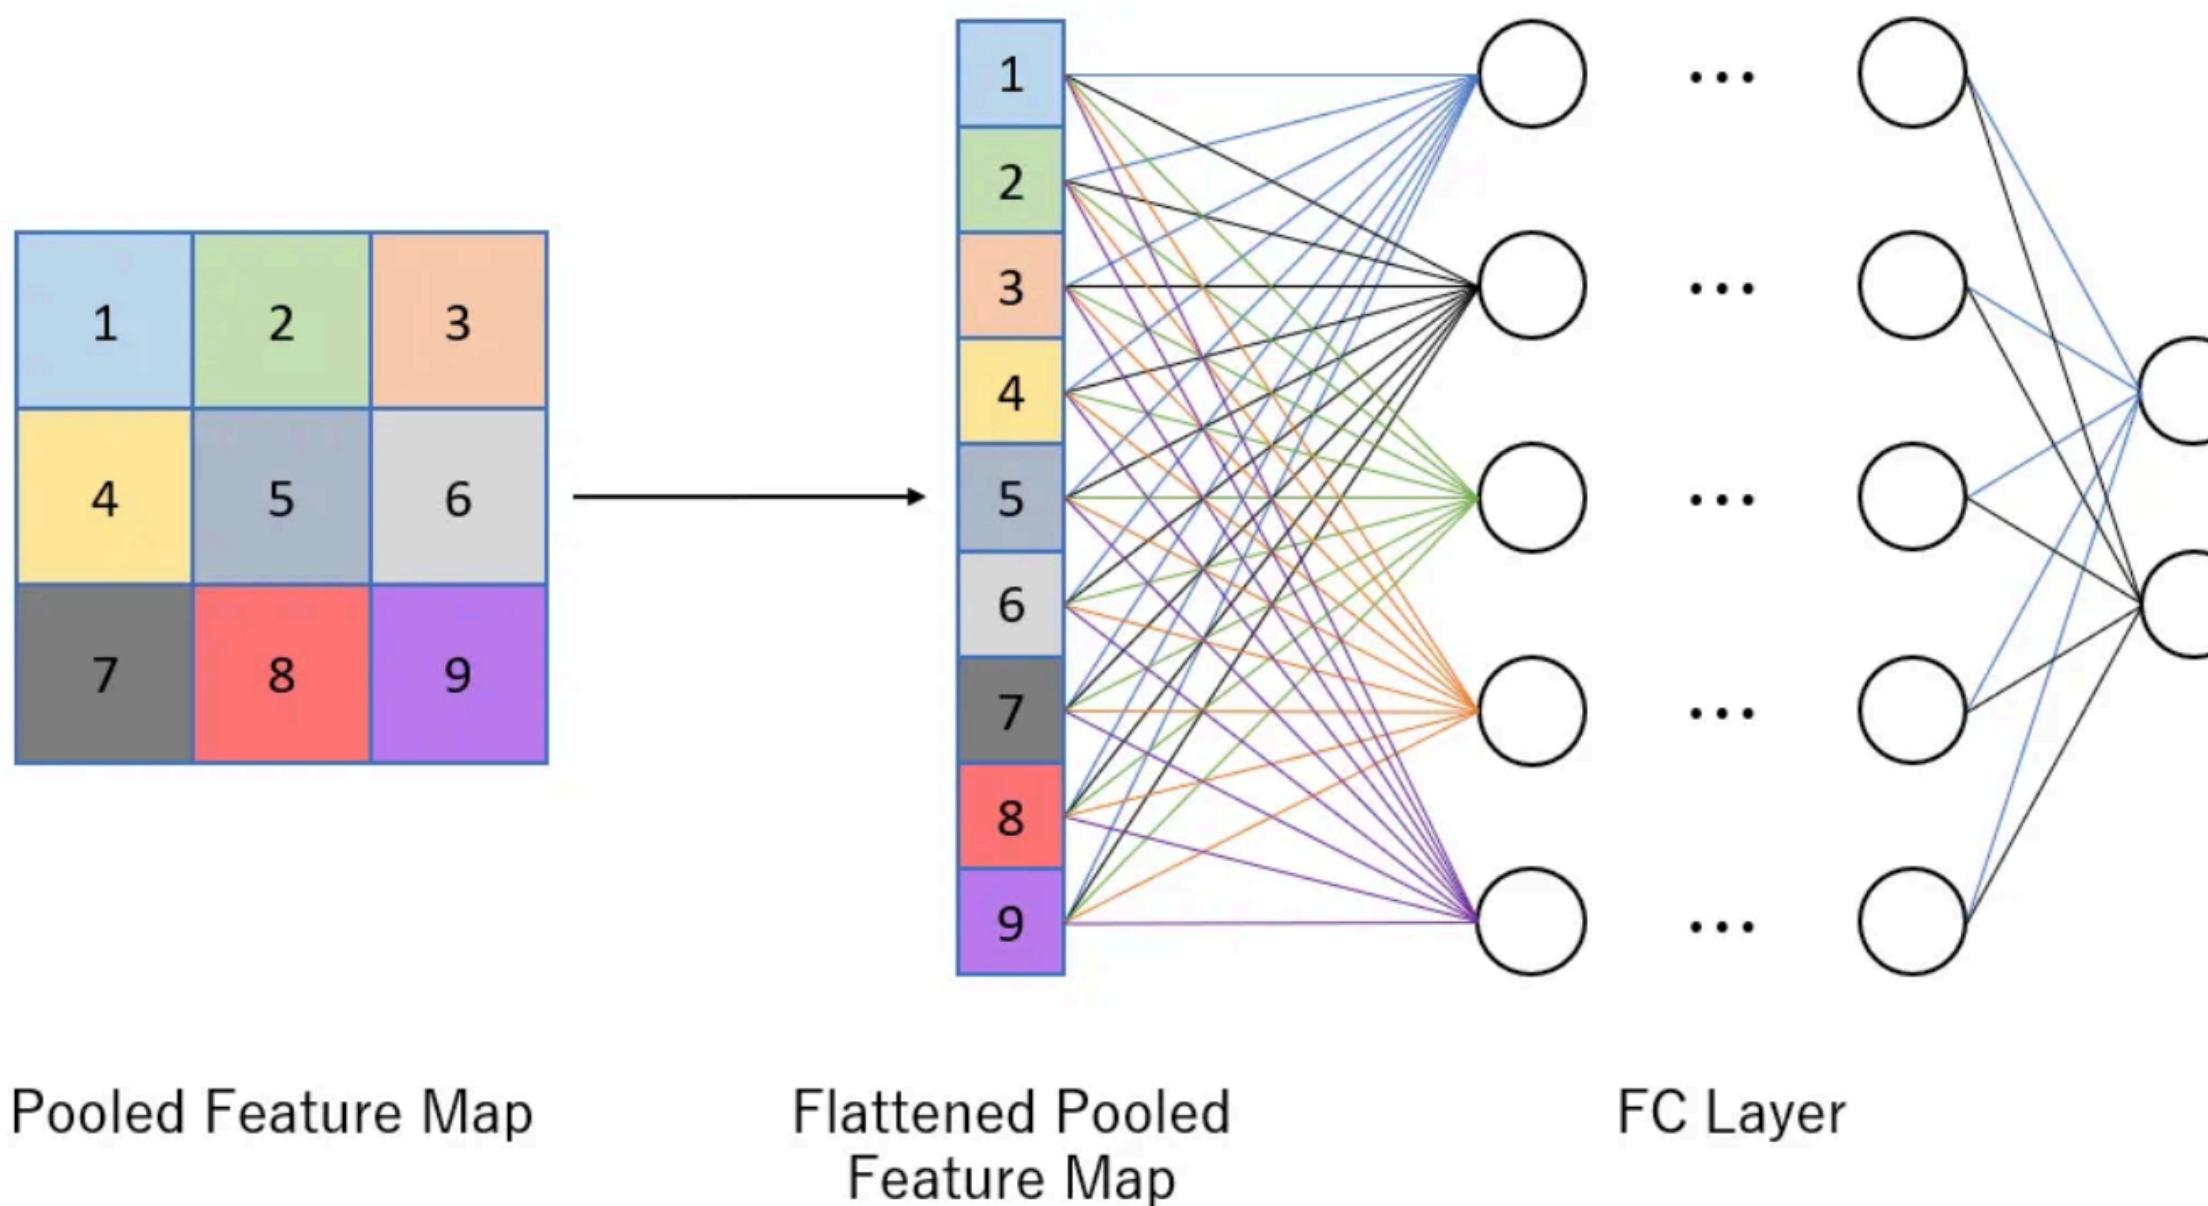
$$\sigma(x) = \max(0, x)$$



# Fully Connected layer

## > A Fully Connected layer

> It combines the features learned by earlier convolutional and pooling layers to perform high-level reasoning or classification.



# SOTA CNN Architectures

# SOTA CNN Architectures

---

- The most extensive data for Image Classification
- 3 RGB channels from 0 to 255
- 14,197,122 images
- 1000 classes



# SOTA CNN Architectures

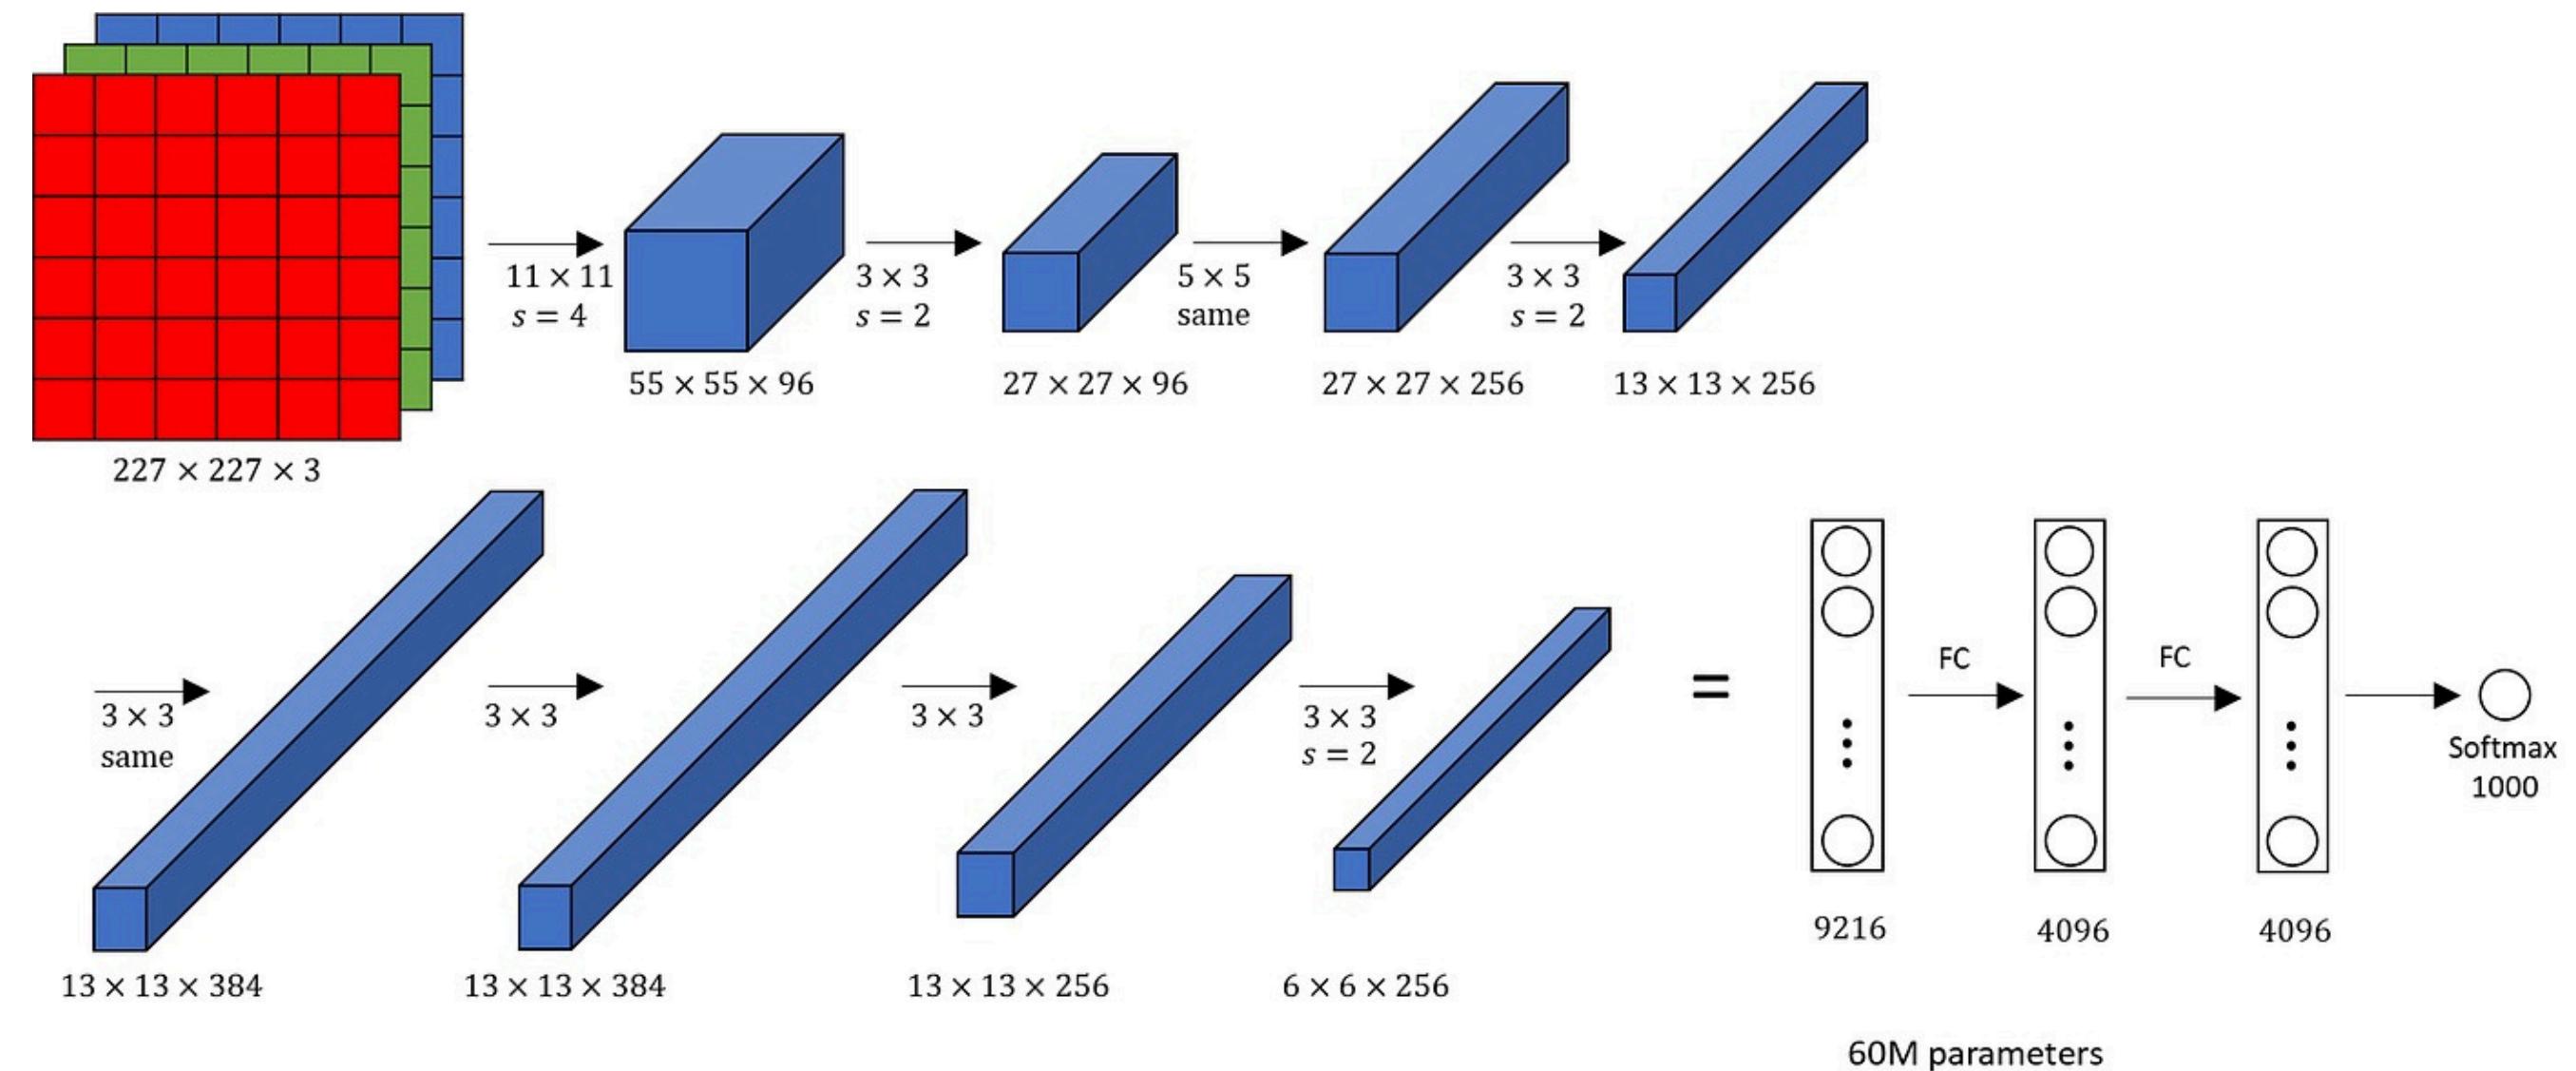
---

- Over time, researchers built advanced CNN architectures to improve performance and efficiency. These architectures introduced key innovations:
  - **AlexNet** [Krizhevsky et al. 2012]: The first CNN to achieve breakthrough performance on image classification.
  - **VGGNet** [Simonyan and Zisserman, 2014]: Used very deep networks (up to 19 layers).
  - **InceptionNet** (GoogLeNet) [Szegedy et al., 2014]: Used multiple filter sizes per layer (Inception modules).
  - **ResNet** [He et al., 2015]: Introduced skip connections for training very deep networks.
  - **EfficientNet** [Tan and Le, 2019]: Found a scaling method that simultaneously scales a CNN's depth, width, and resolution optimally using a single scaling coefficient.

# SOTA CNN Architectures

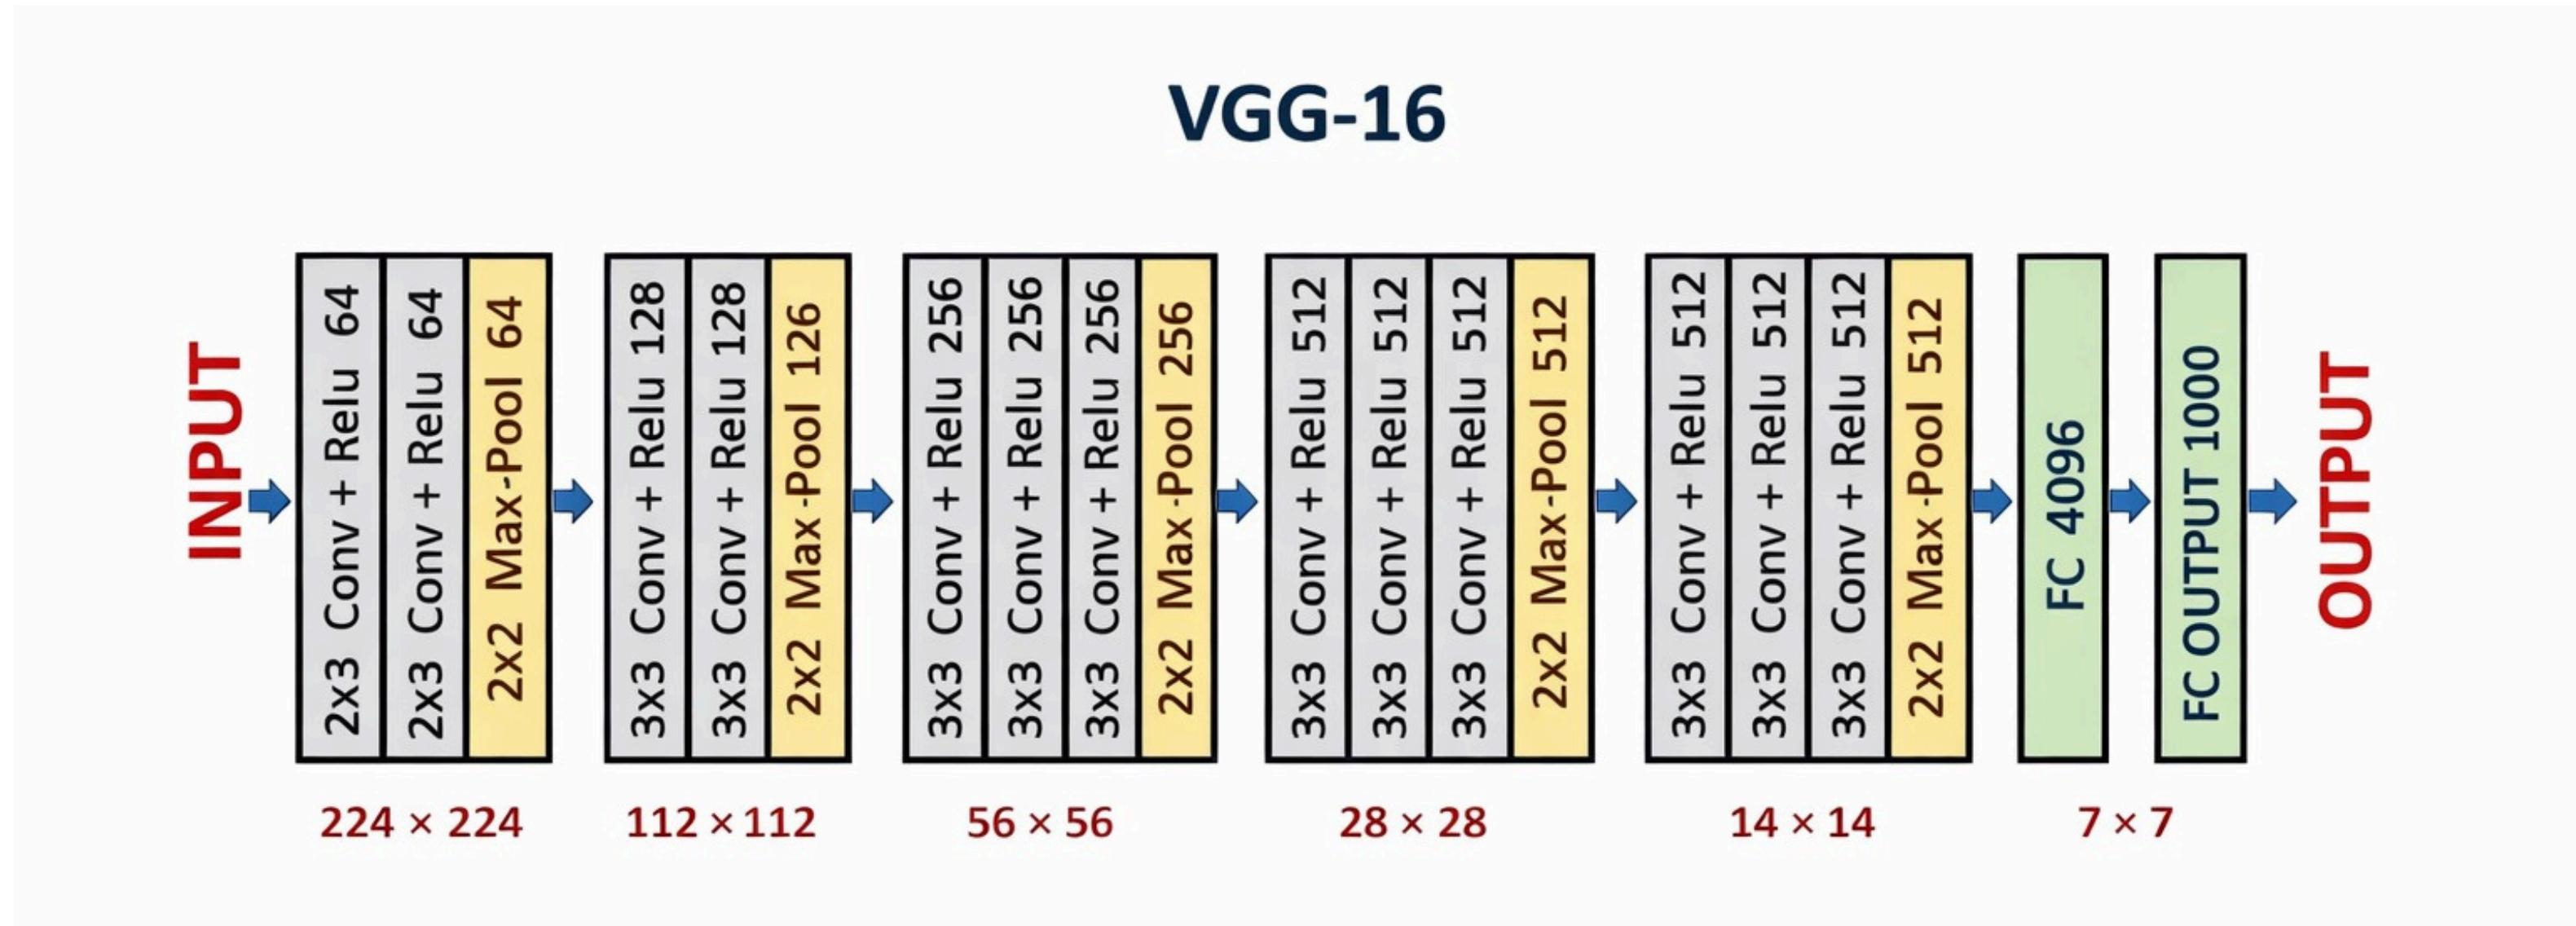
- **AlexNet** is a deep convolutional neural network that revolutionized computer vision by winning the ImageNet 2012 competition with a large margin.

- 5 Convolutional layers
- 3 Fully Connected layers
- ReLU after every layer
- Max Pooling
- Dropout



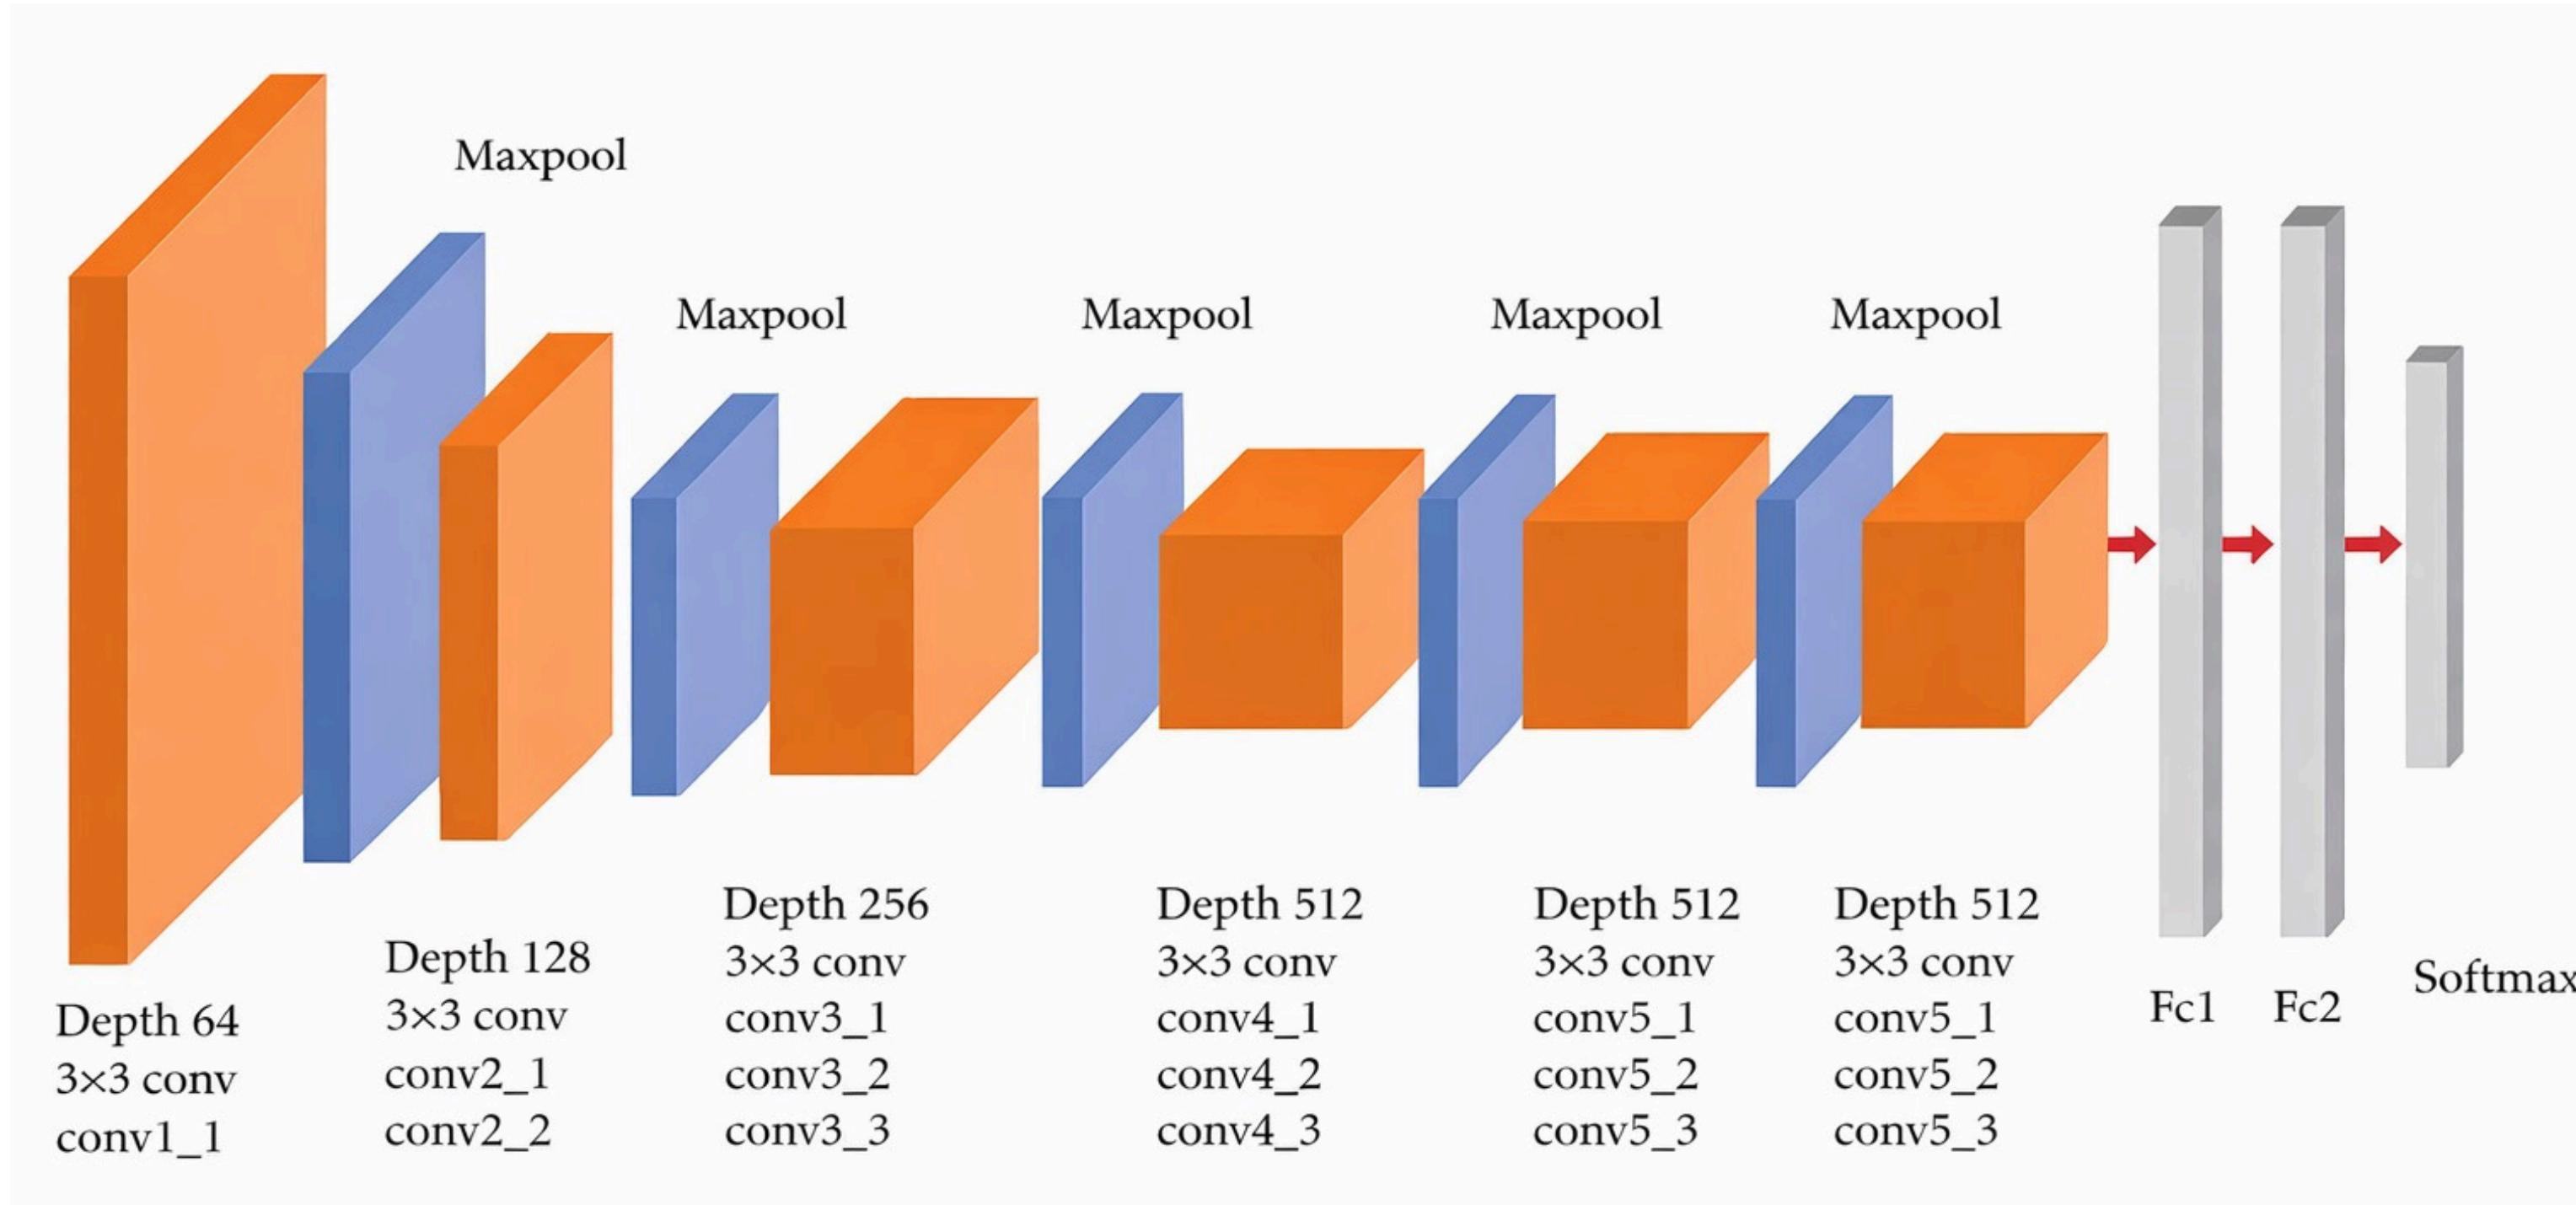
# SOTA CNN Architectures

- **VGG16** image highlight 5 convolutional blocks and 16 weight layers



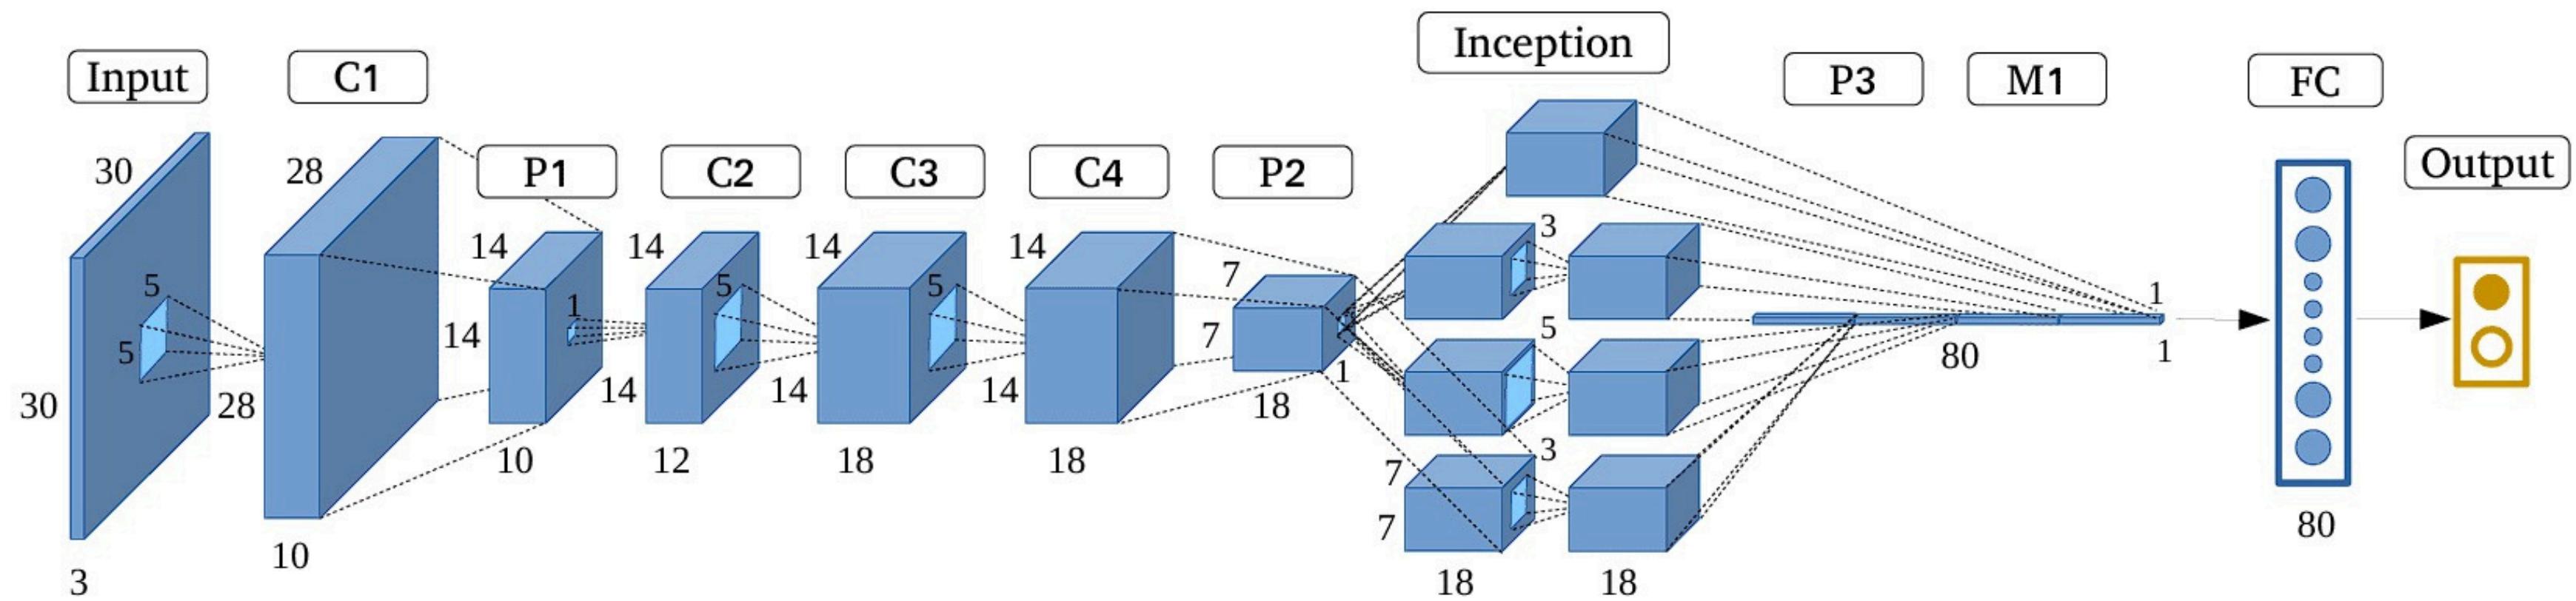
# SOTA CNN Architectures

➤ VGG19 image highlight deeper blocks with extra  $3 \times 3$  conv layers



# SOTA CNN Architectures

➤ InceptionNet achieves a depth of 22 layers while maintaining efficient computation.



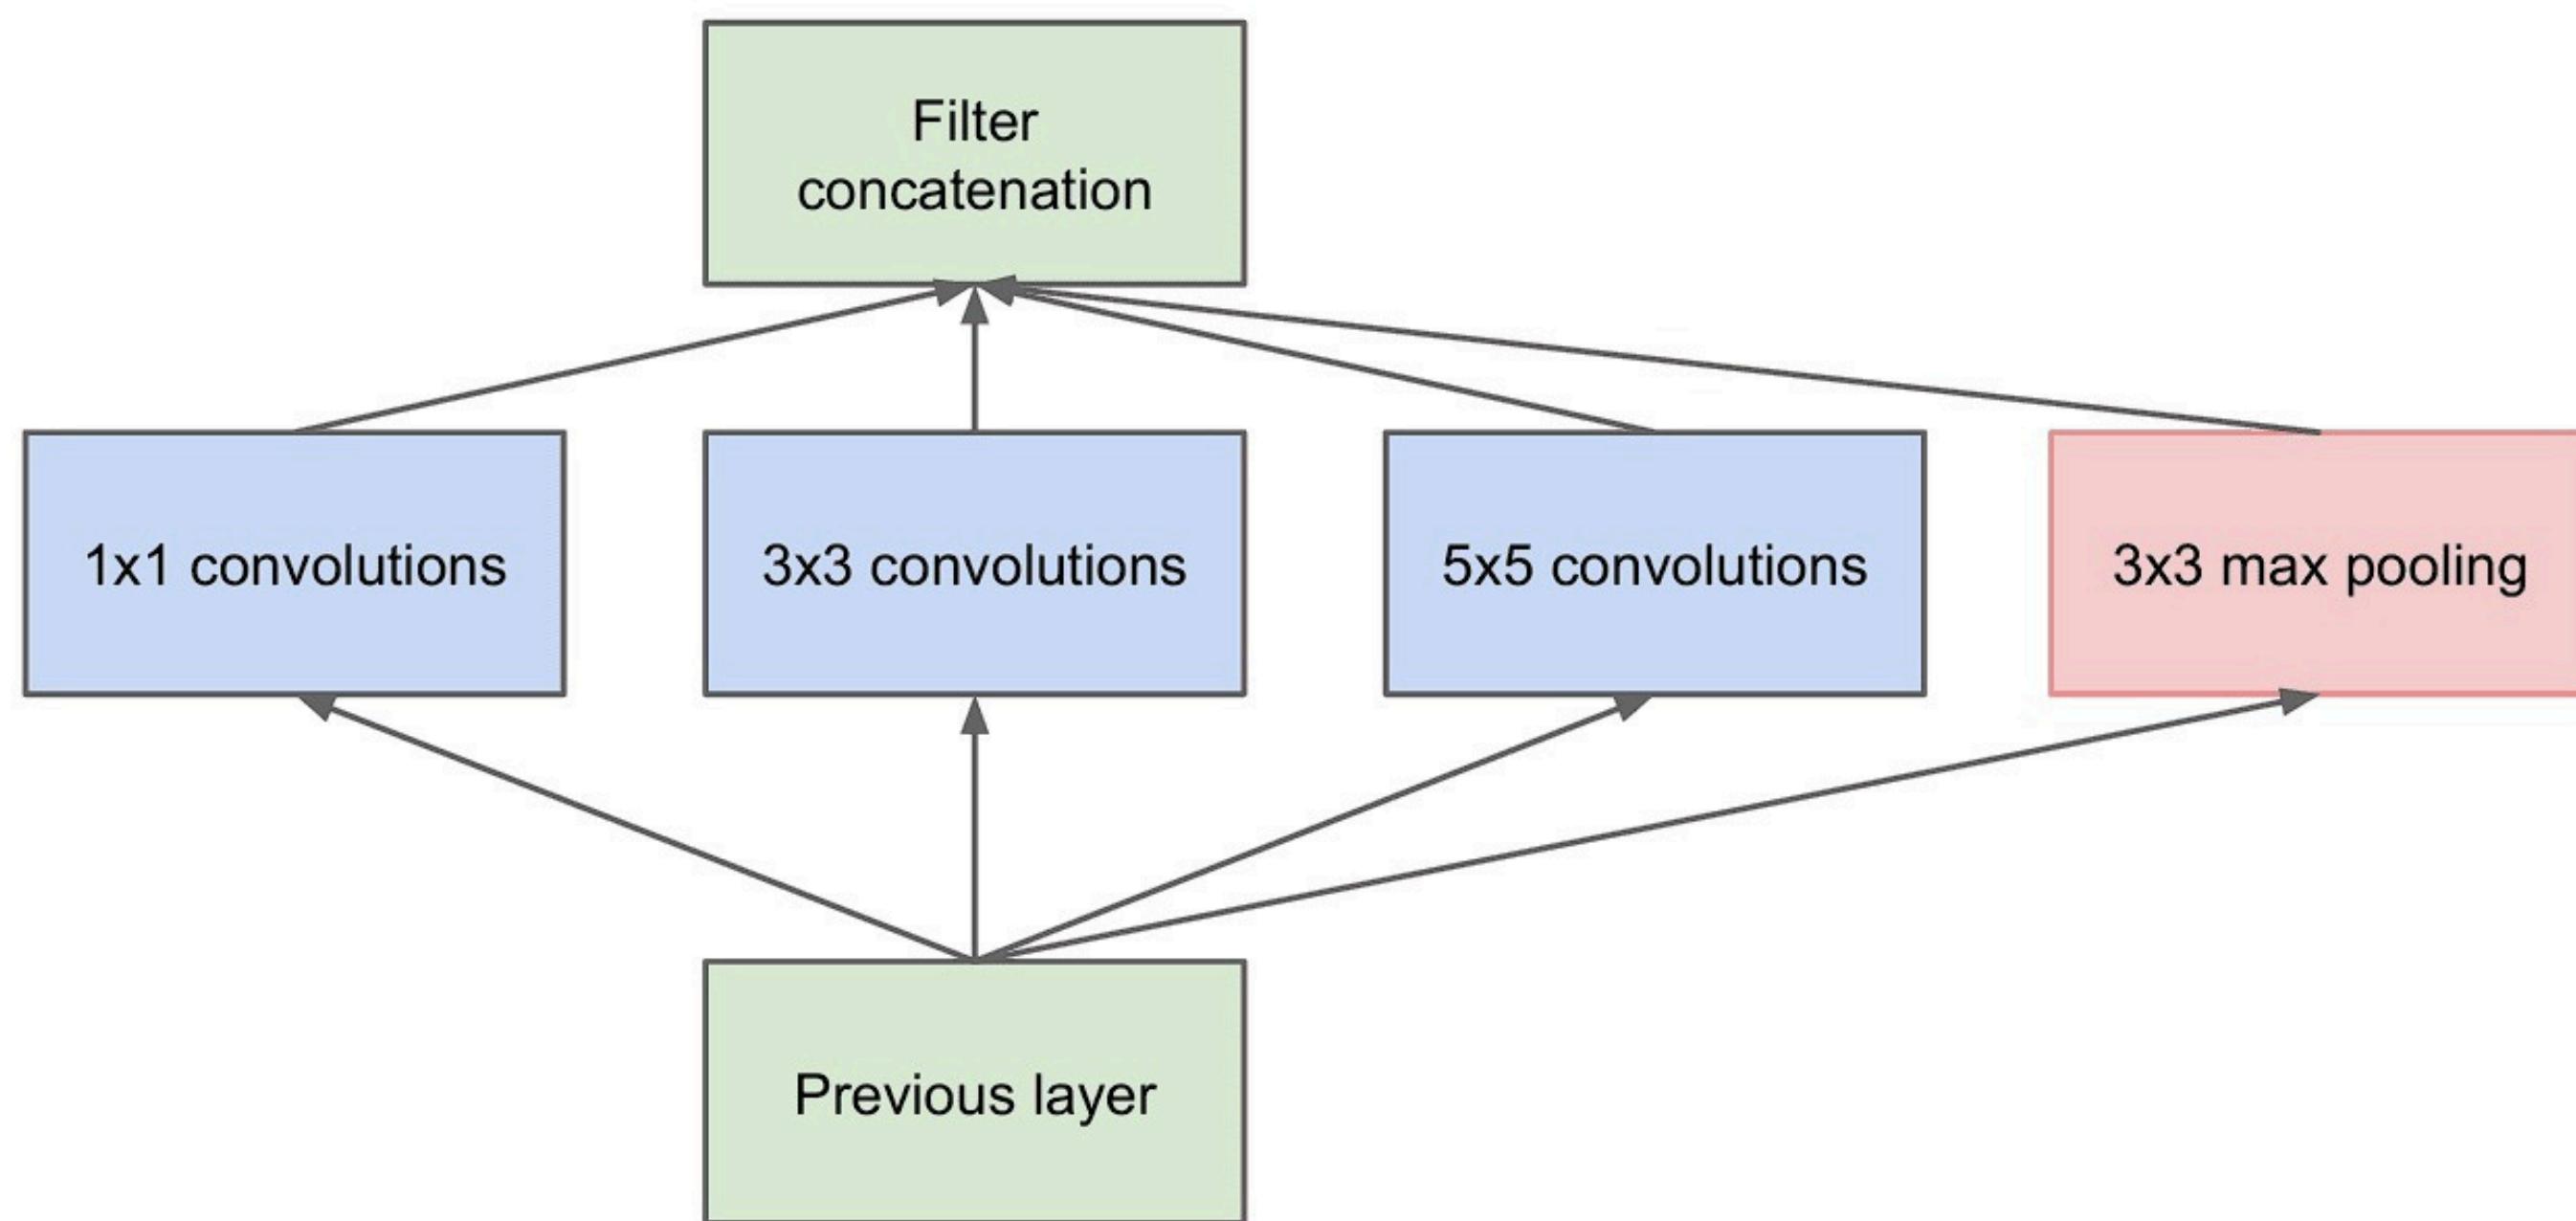
# SOTA CNN Architectures

type	patch size/stride	output size	depth	#1x1	#3x3 reduce	#3x3	#5x5 reduce	#5x5	pool proj	params	ops
<b>convolution</b>	7x7/2	112x112x64	1							2.7K	34M
<b>max pool</b>	3x3/2	56x56x64	0								
<b>convolution</b>	3x3/1	56x56x192	2		64	192				112K	360M
<b>max pool</b>	3x3/2	28x28x192	0								
<b>inception (3a)</b>		28x28x256	2	64	96	128	16	32	32	159K	128M
<b>inception (3b)</b>		28x28x480	2	128	128	192	32	96	64	380K	304M
<b>max pool</b>	3x3/2	14x14x480	0								
<b>inception (4a)</b>		14x14x512	2	192	96	208	16	48	64	364K	73M
<b>inception (4b)</b>		14x14x512	2	160	112	224	24	64	64	437K	88M
<b>inception (4c)</b>		14x14x512	2	128	128	256	24	64	64	463K	100M
<b>inception (4d)</b>		14x14x528	2	112	144	288	32	64	64	580K	119M
<b>inception (4e)</b>		14x14x832	2	256	160	320	32	128	128	840K	170M
<b>max pool</b>	3x3/2	7x7x832	0								
<b>inception (5a)</b>		7x7x832	2	256	160	320	32	128	128	1072K	54M
<b>inception (5b)</b>		7x7x1024	2	384	192	384	48	128	128	1388K	71M
<b>avg pool</b>	7x7/1	1x1x1024	0								
<b>dropout (40%)</b>		1x1x1024	0								
<b>linear</b>		1x1x1000	1							1000K	1M
<b>softmax</b>		1x1x1000	0								

# SOTA CNN Architectures

---

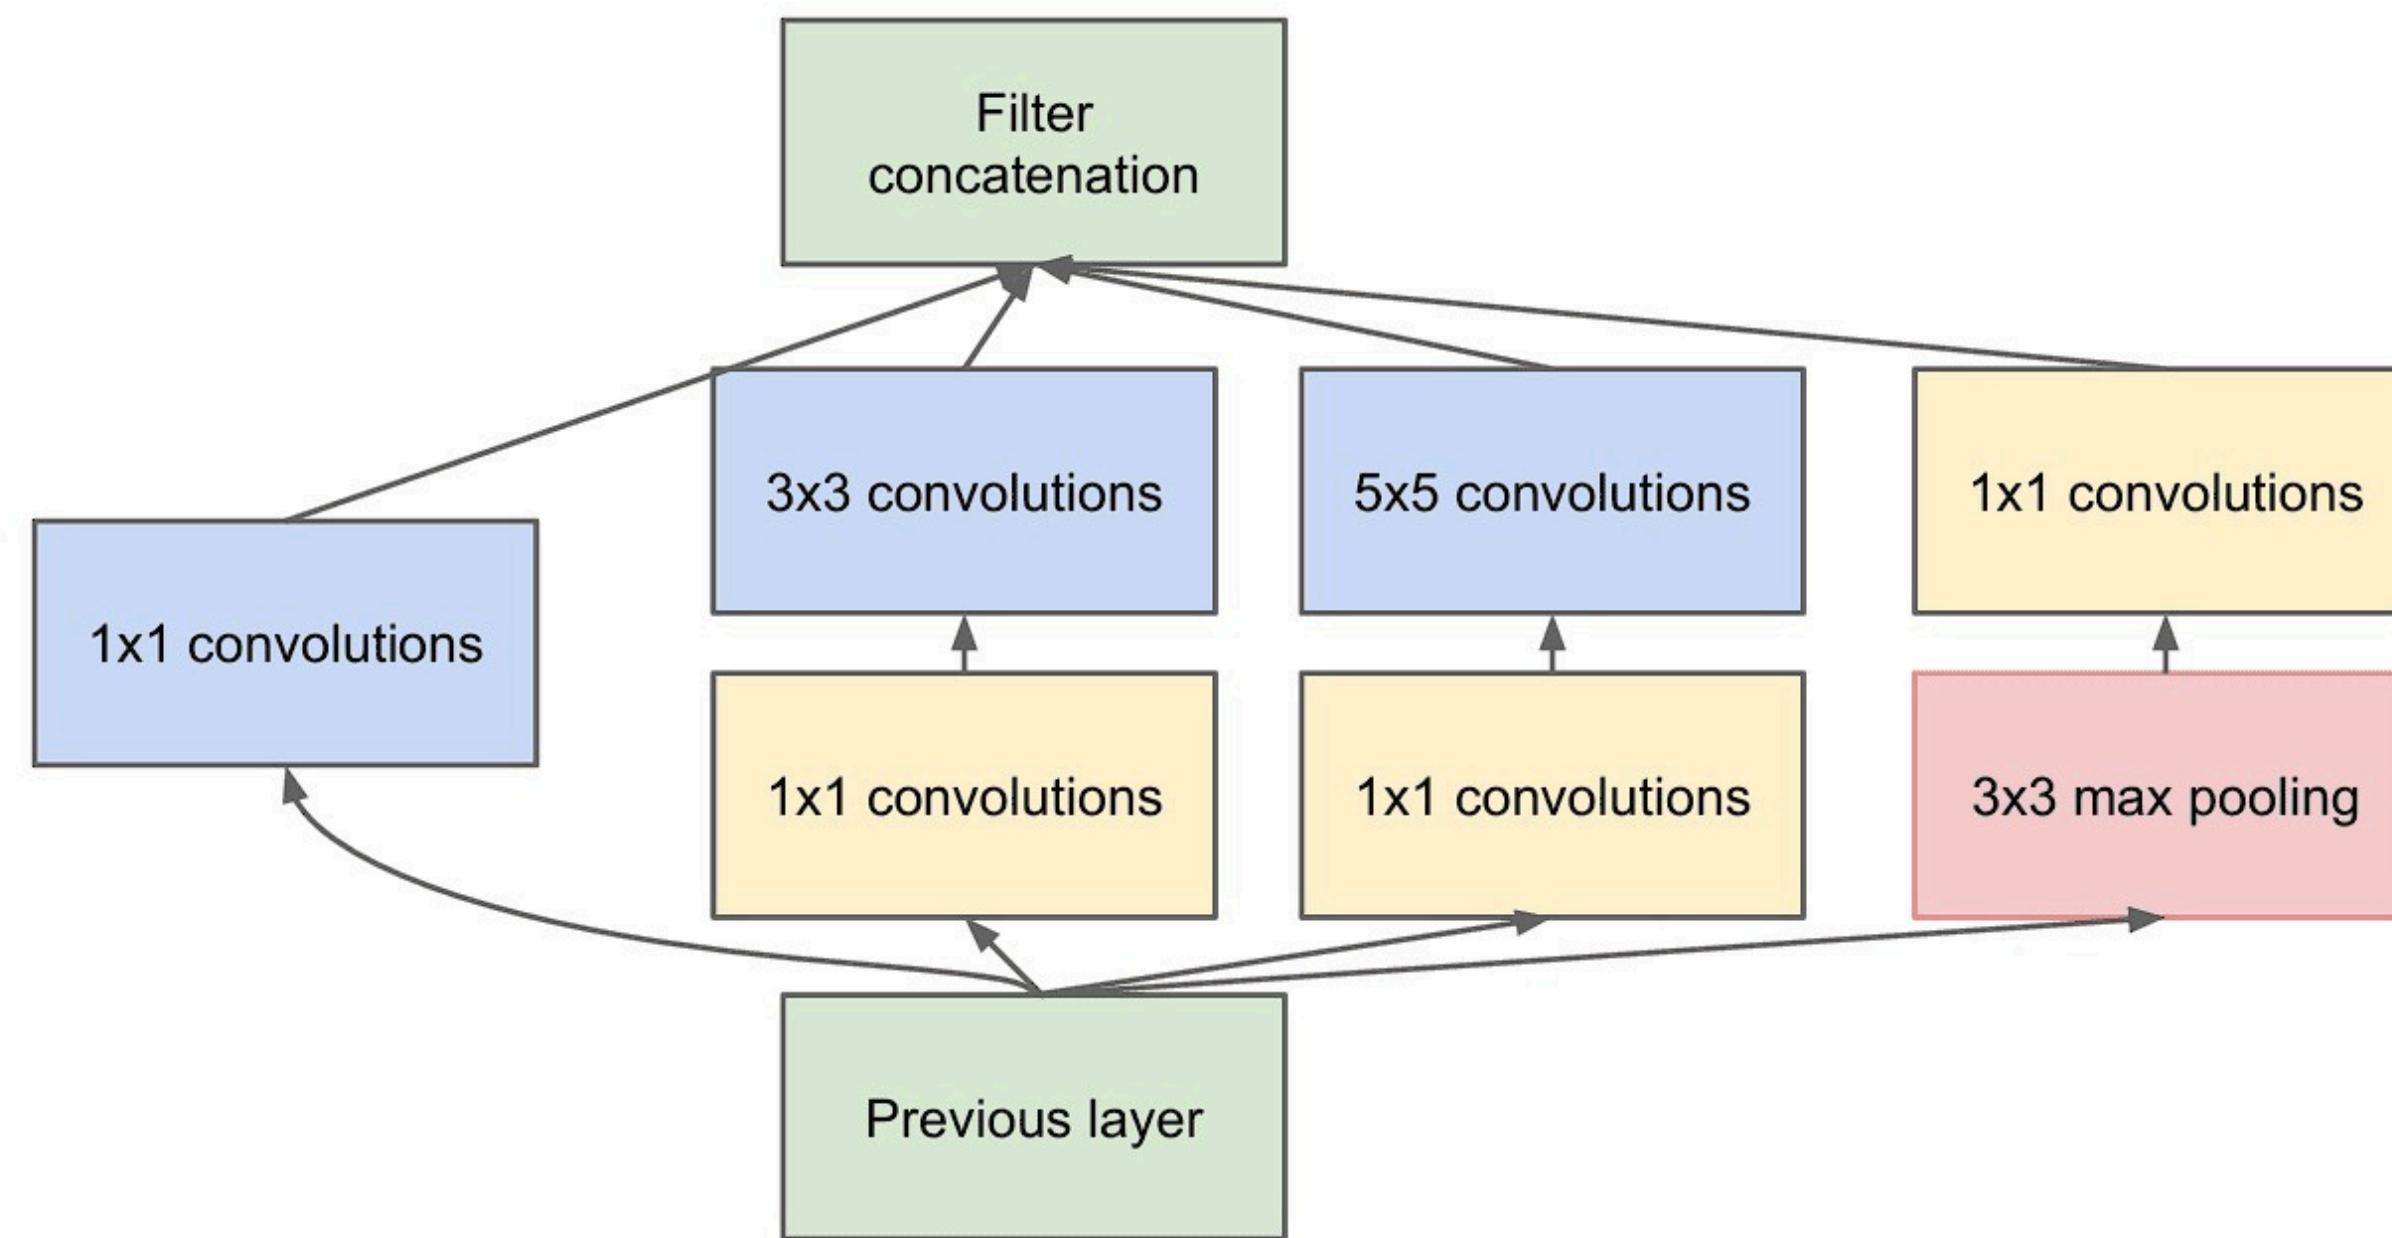
- **Inception module:** Uses multiple filter sizes ( $1\times 1$ ,  $3\times 3$ ,  $5\times 5$ ), in parallel, to capture different features, then combines their outputs.



# SOTA CNN Architectures

---

- **Inception module:** Uses multiple filter sizes ( $1 \times 1$ ,  $3 \times 3$ ,  $5 \times 5$ ), in parallel, to capture different features, then combines their outputs.



# SOTA CNN Architectures

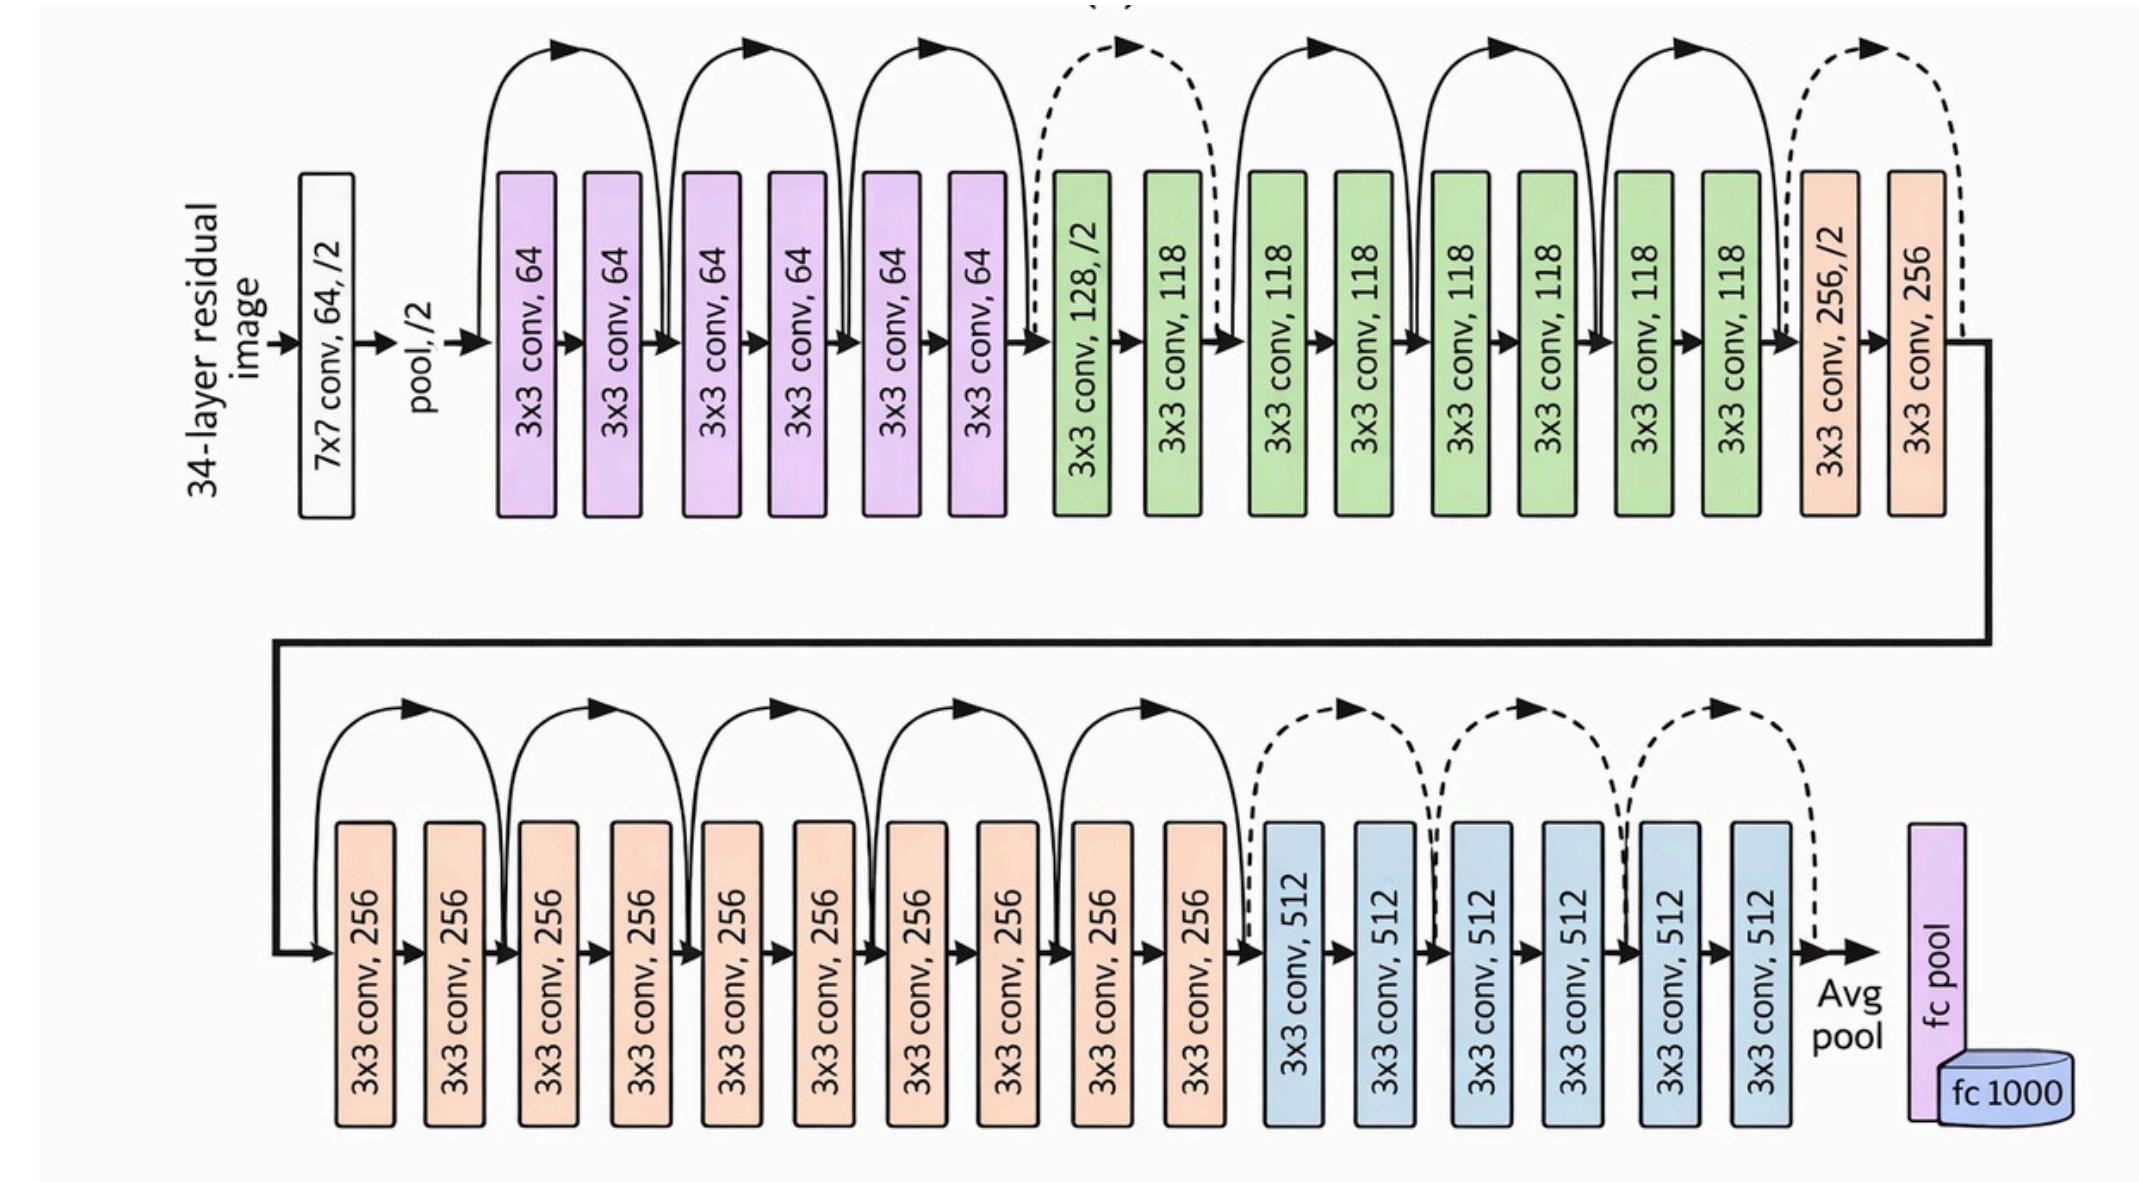
---

- Deeper Is Not Always Better
- Simply increasing network depth does not guarantee higher accuracy and can even degrade performance.
- In very deep networks, gradients can become extremely small during backpropagation. This slows learning or completely prevents effective weight updates in early layers.

# SOTA CNN Architectures

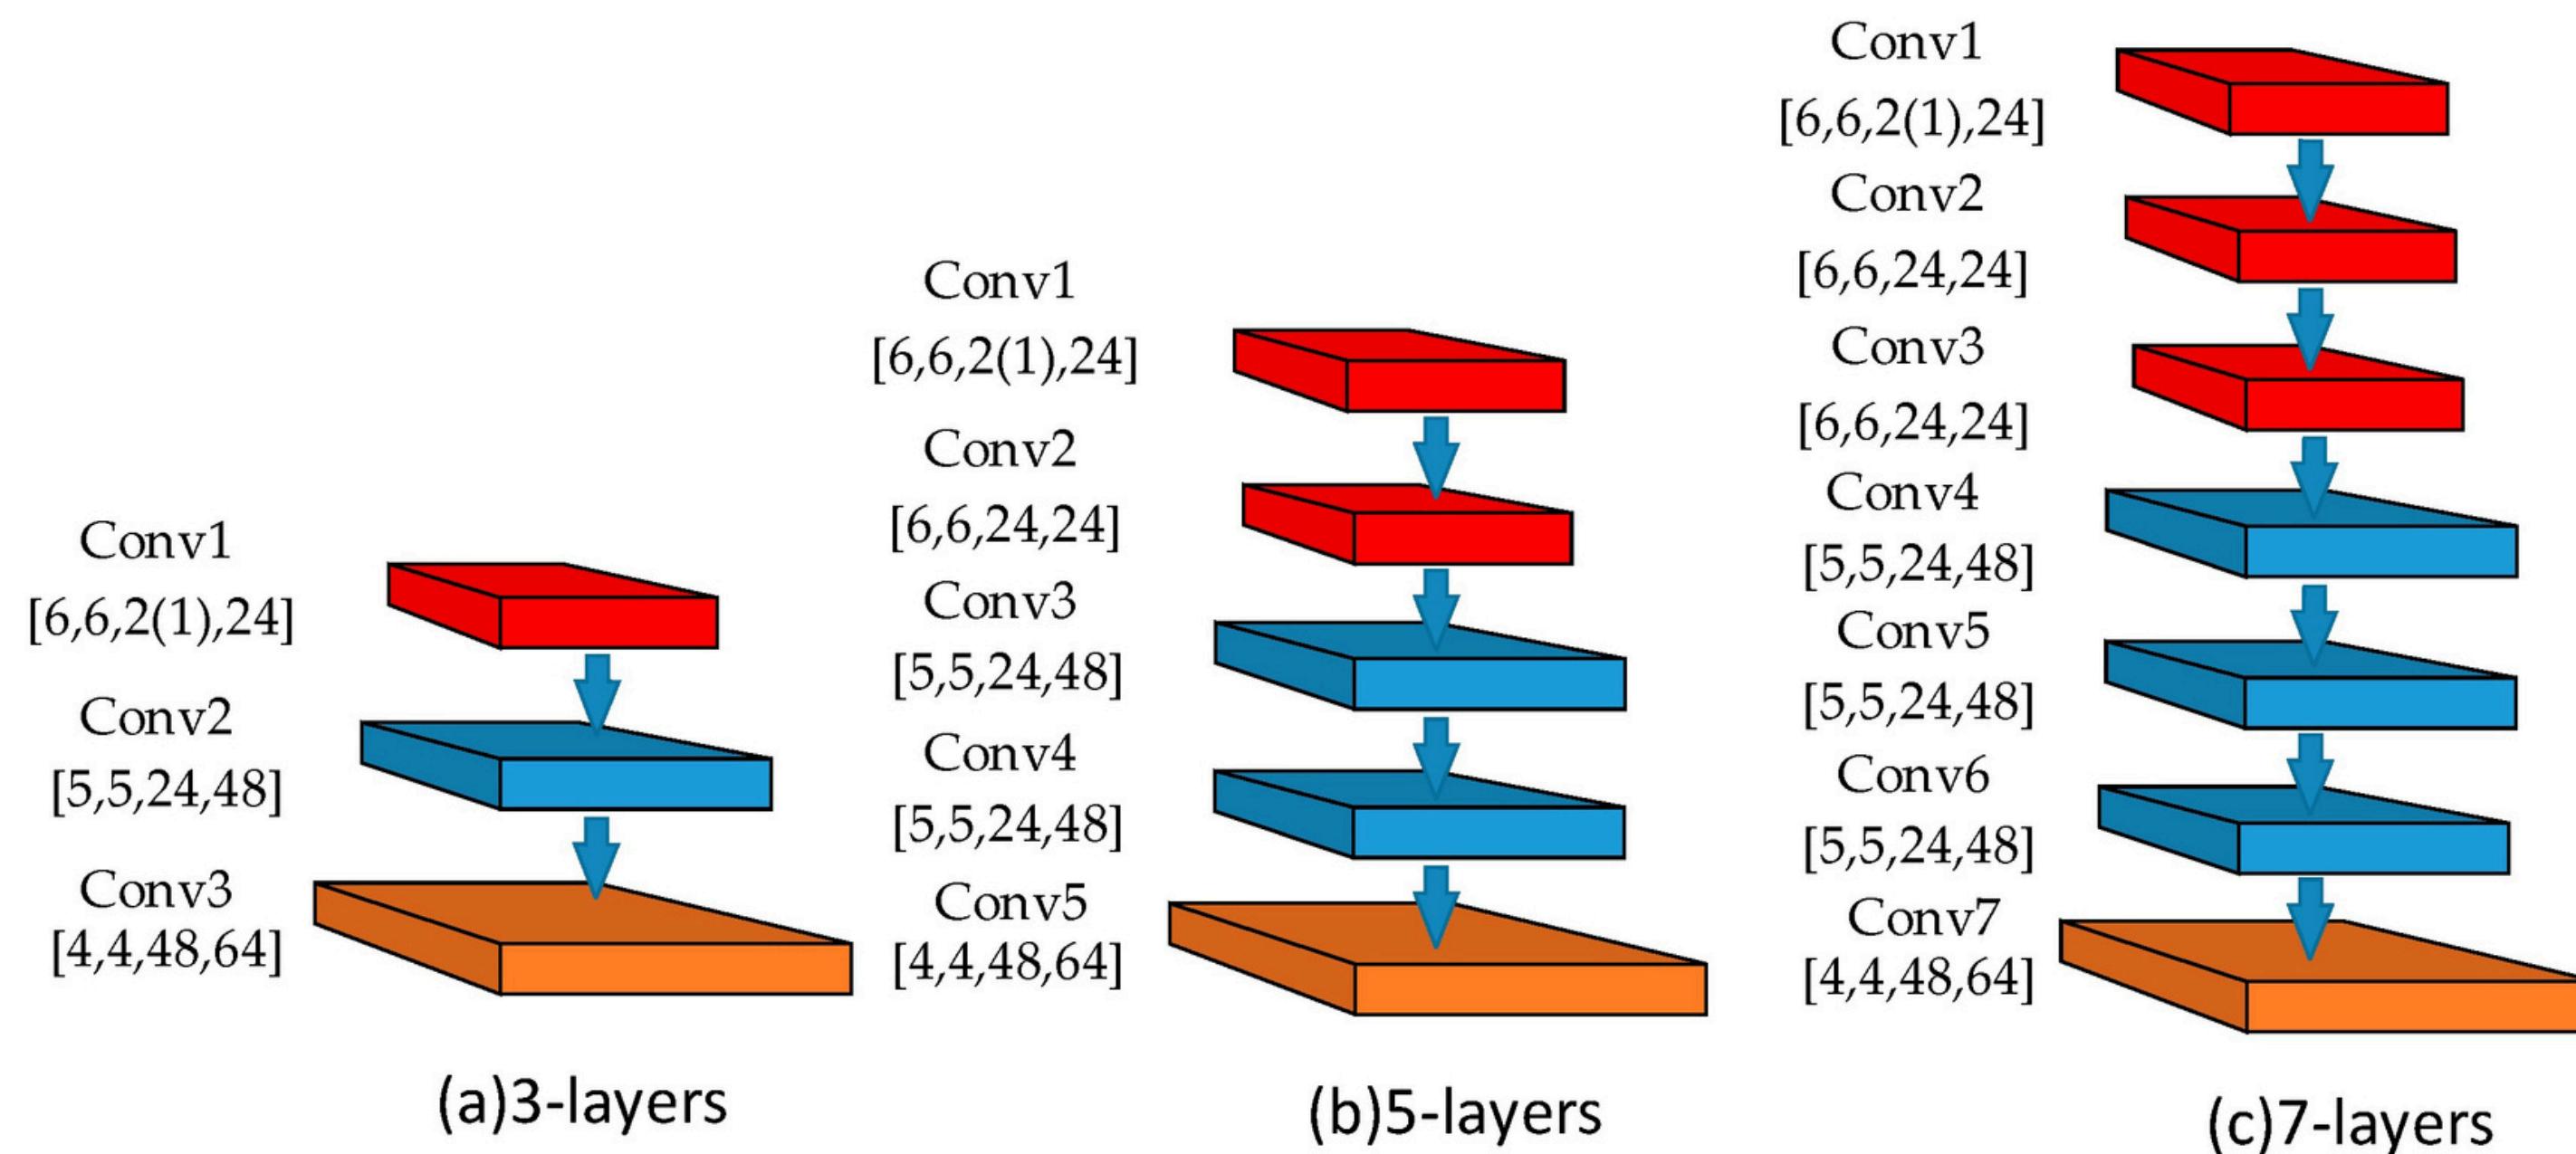
## > ResNet

- Very deep networks using residual connections
- 152-layer model for ImageNet
- Residual: A shortcut connection that helps the network pass information through layers more easily.



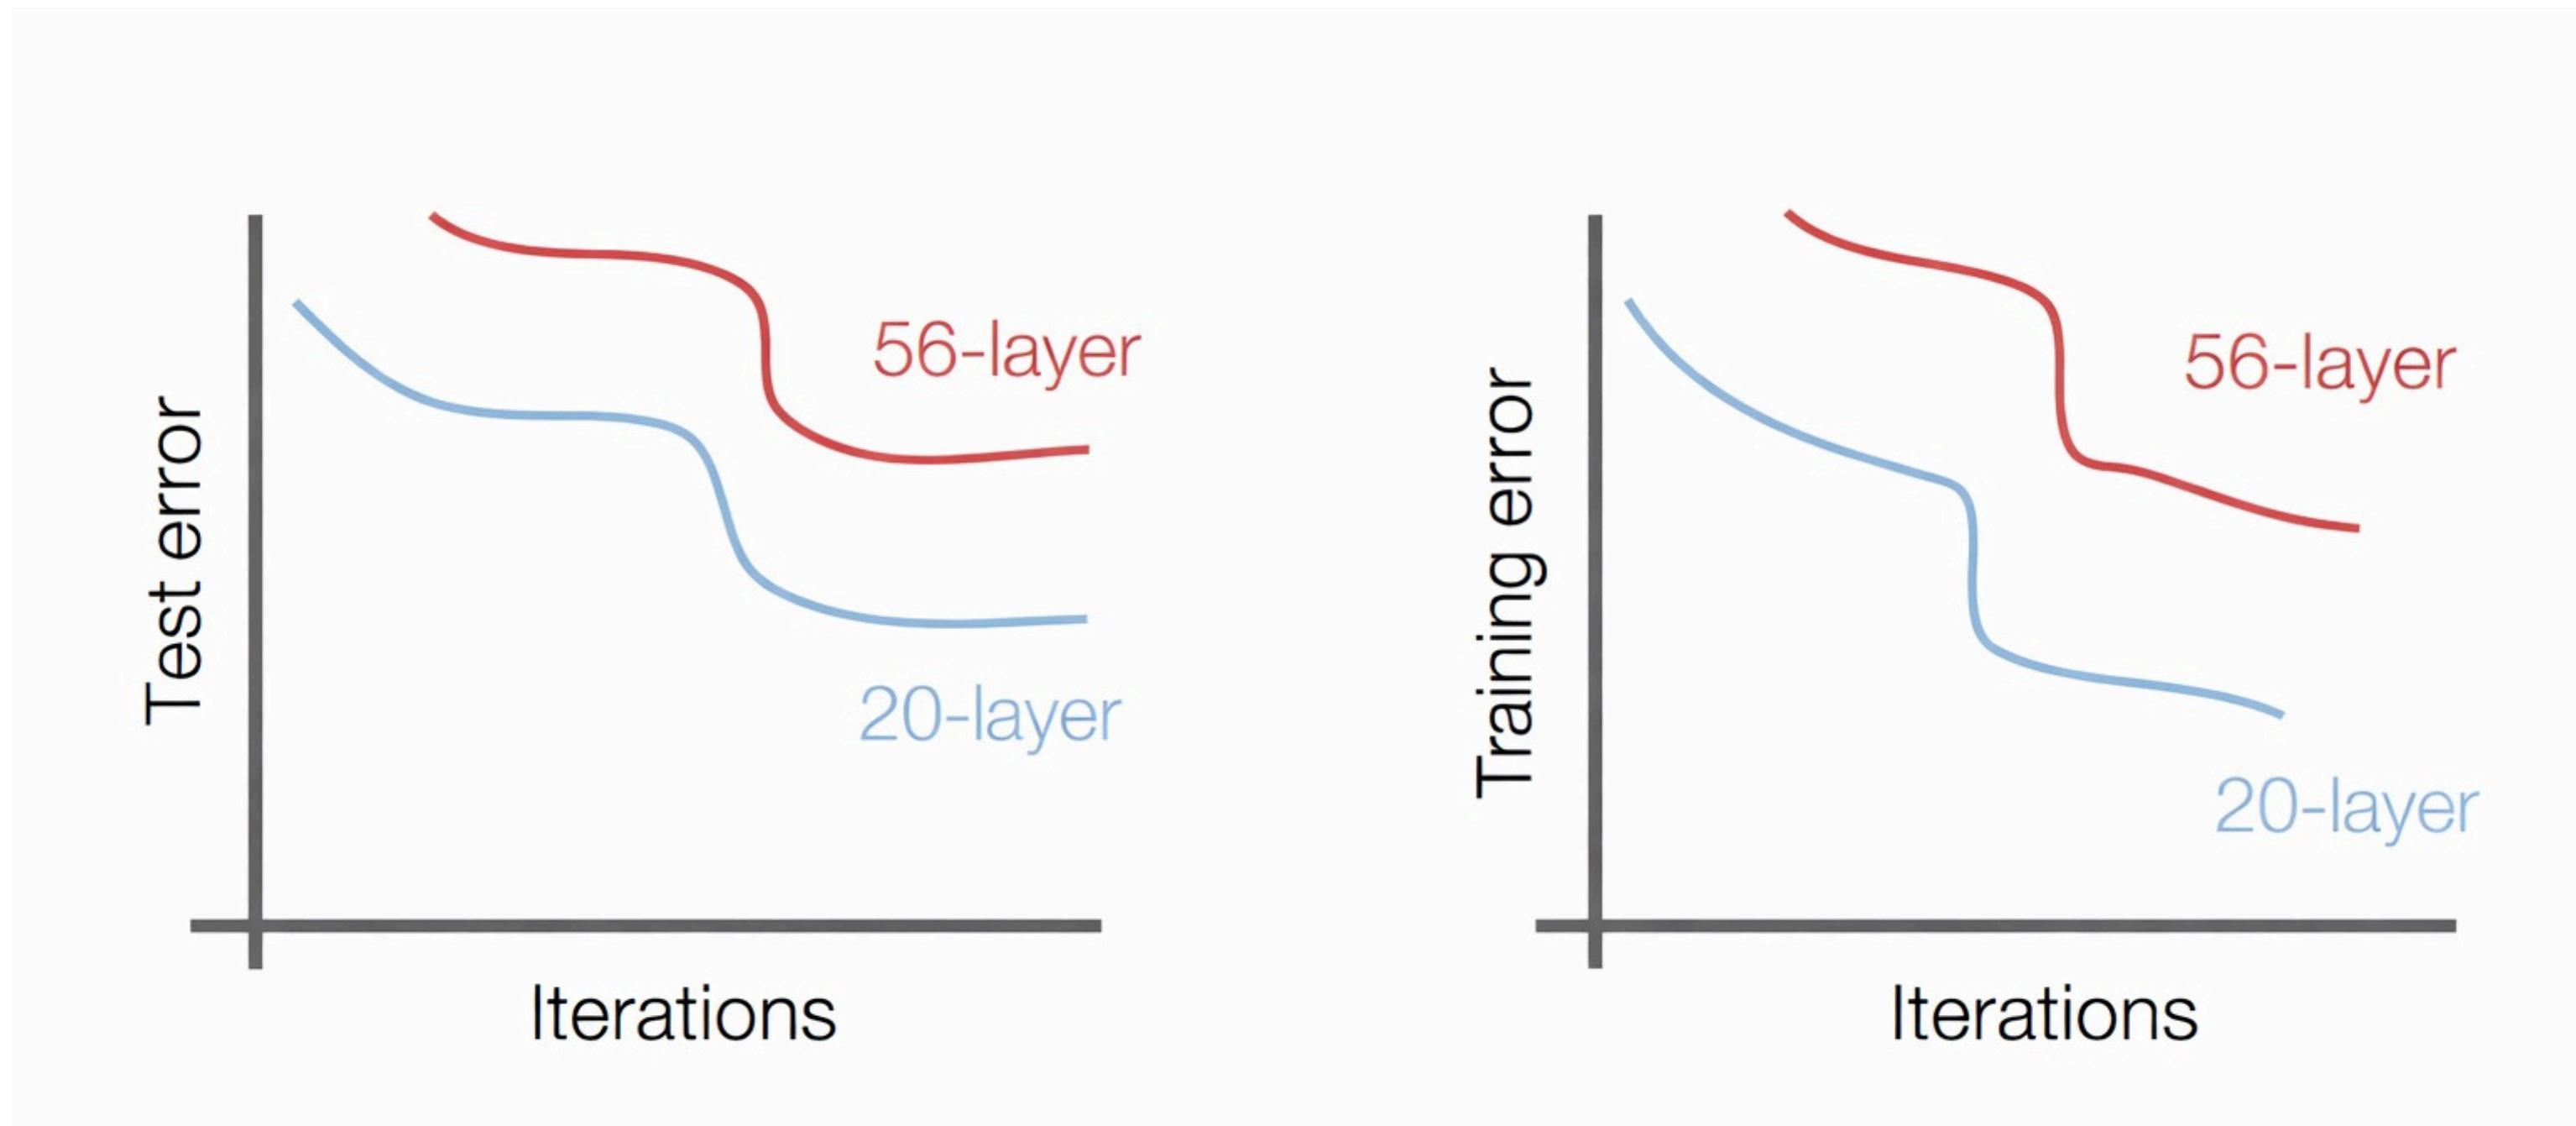
# SOTA CNN Architectures

- What happens when we continue stacking deeper layers on a "plain" Convolutional neural network?



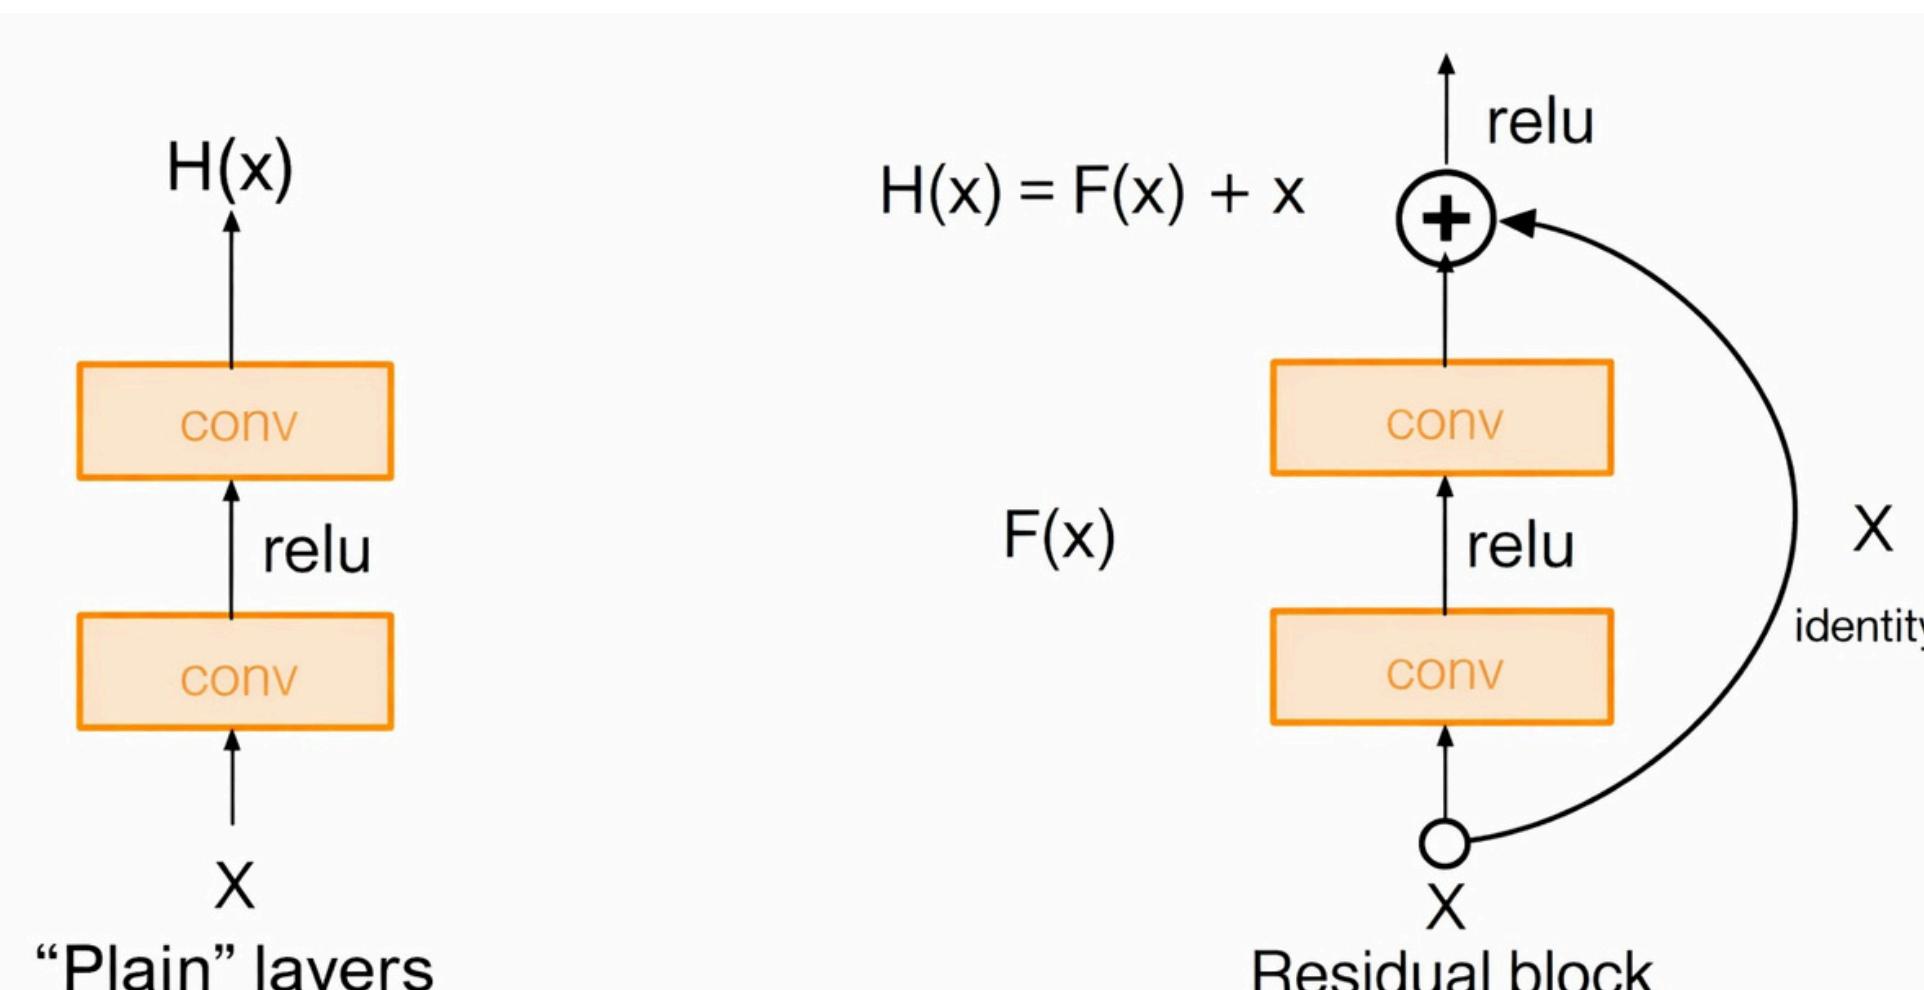
# SOTA CNN Architectures

What happens when we continue stacking deeper layers on a "plain" Convolutional neural network?



# SOTA CNN Architectures

- Deep models have more representation power (more parameters) than shallower models.
- The problem is an optimization problem, deeper models are harder to optimize
- Use network layers to fit a residual mapping instead of directly trying to fit a desired underlying mapping



# SOTA CNN Architectures

---

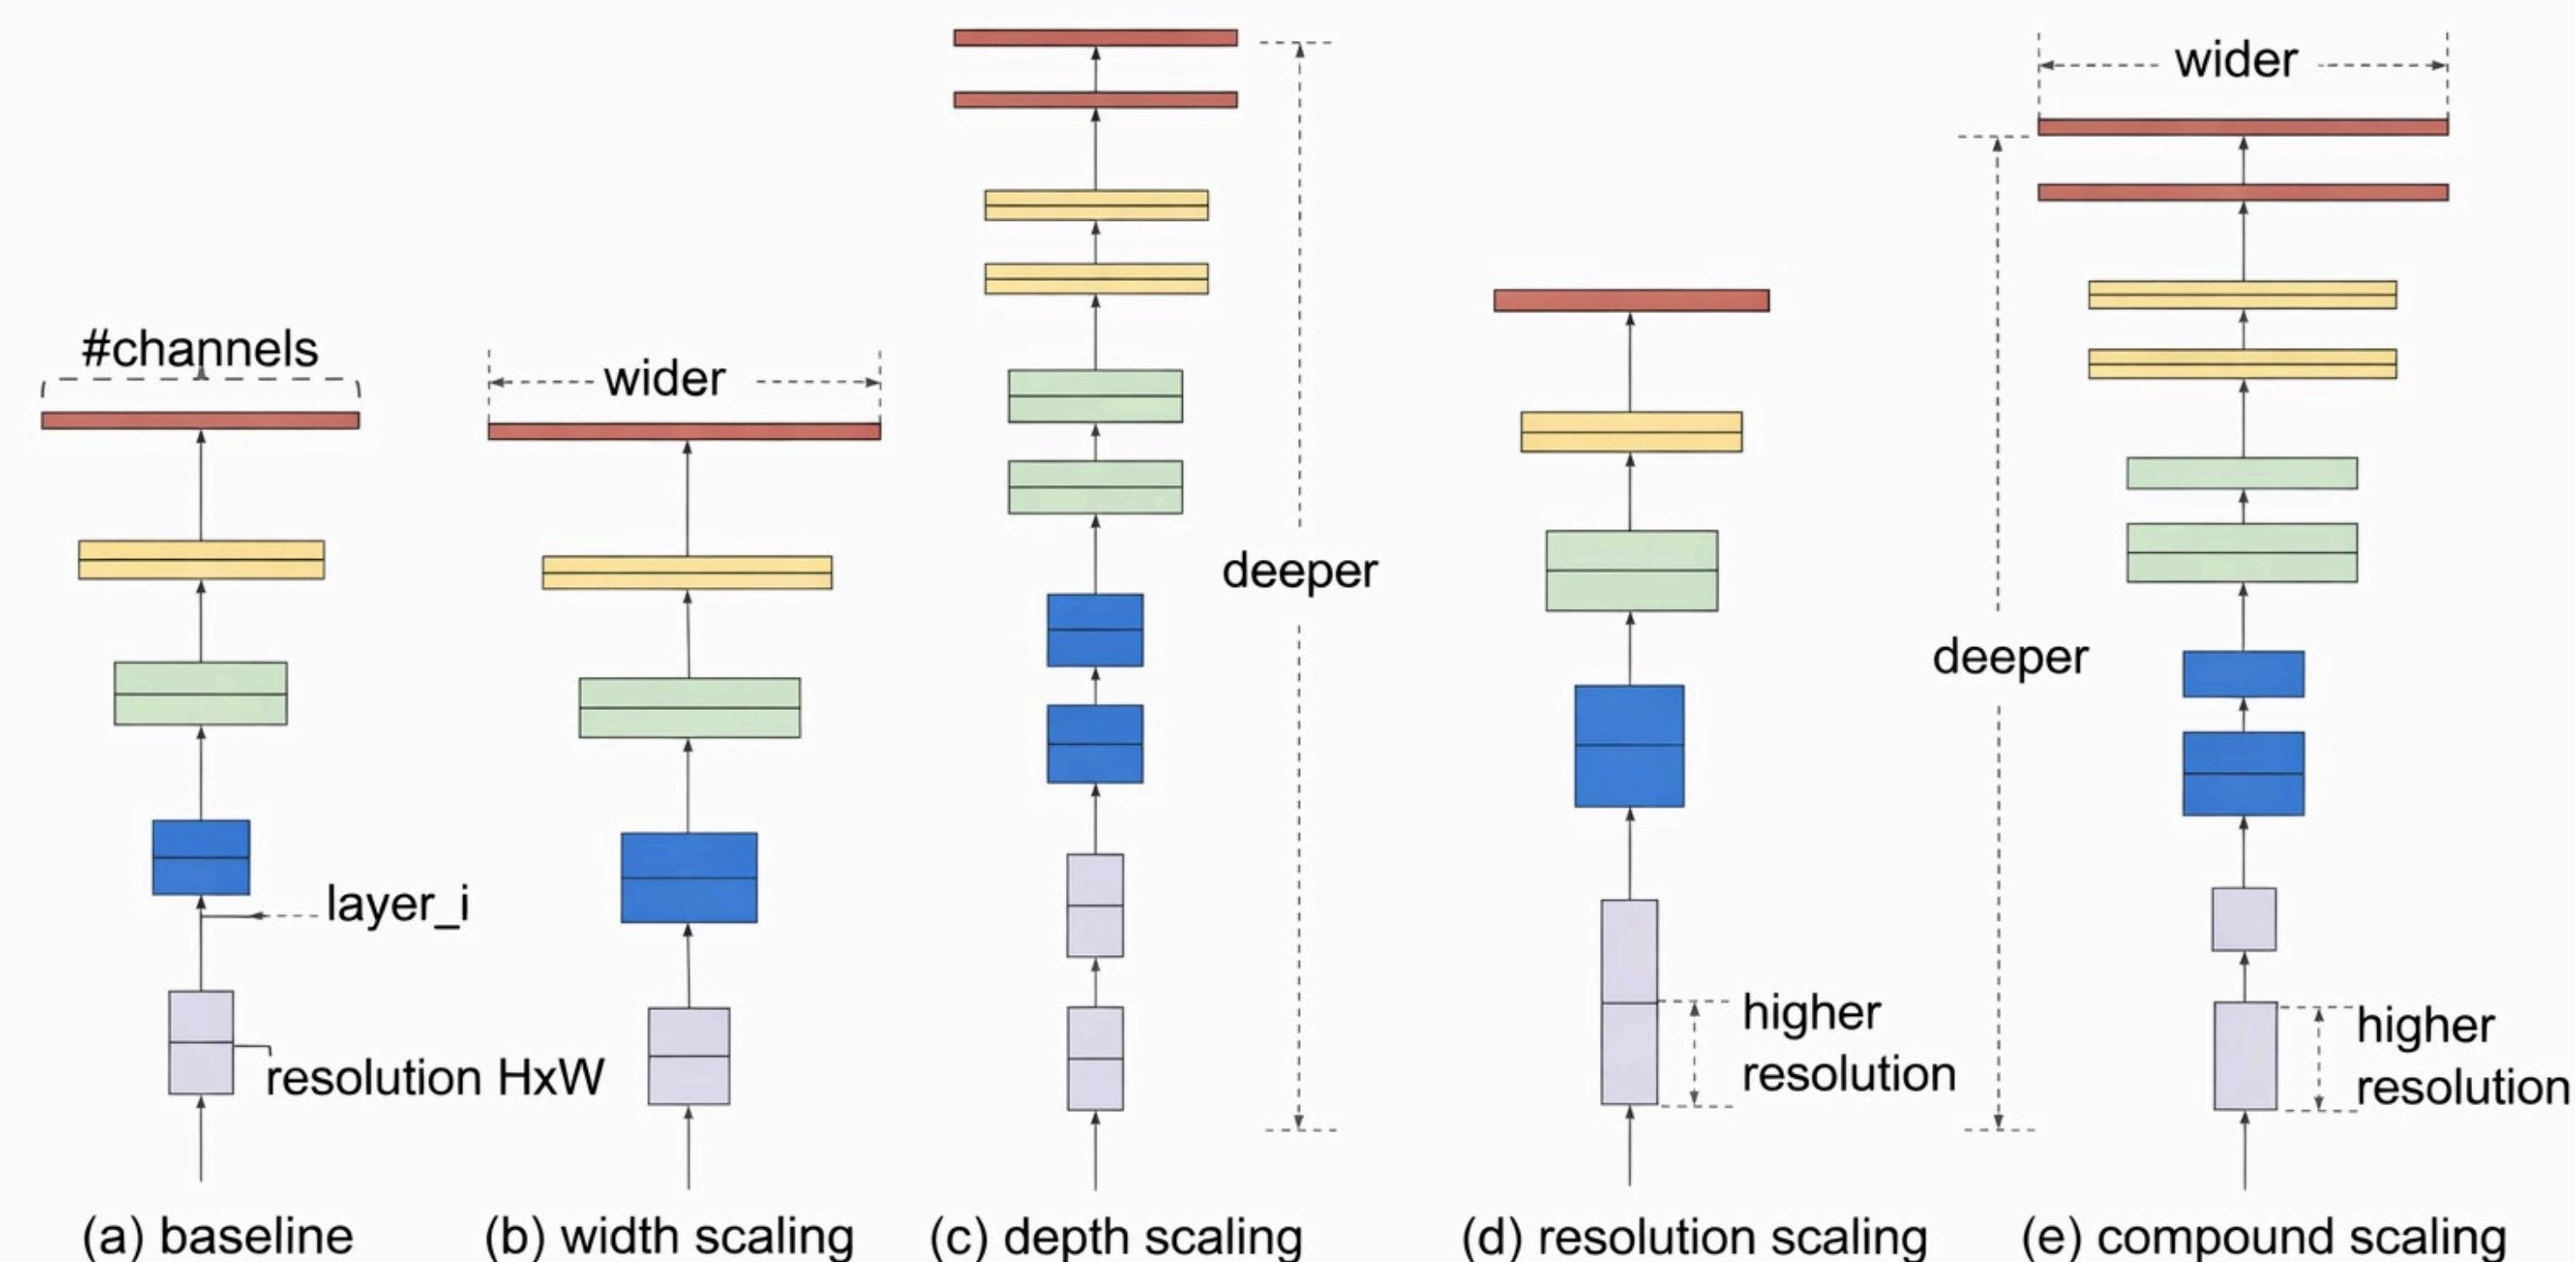
› **EfficientNet** introduced a compound scaling method to systematically find the optimal balance among depth, width, and resolution.

- Increase the number of layers (depth)
- Increase the number of neurons in each layer (width)
- Use higher-resolution images (resolution)

*But finding the right balance of these three was largely based on trial and error.*

# SOTA CNN Architectures

## > EfficientNet



# SOTA CNN Architectures

---

## ➤ EfficientNet

Compound Scaling uses a single scaling coefficient ( $\phi$ ) to control:

$\alpha^\phi$  : Depth,     $\beta^\phi$  : Width,     $\gamma^\phi$  : Resolution

Goal: find  $\alpha, \beta, \gamma$  that balance accuracy and efficiency, then scale up optimally by increasing  $\phi$

# SOTA CNN Architectures

---

## ➤ EfficientNet

Optimizes depth, width, and resolution using this constraint:

$$\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$$

Increasing depth ( $\alpha$ ) increases FLOPs linearly

Increasing width ( $\beta$ ) increases FLOPs quadratically ( $\beta^2$ )

Increasing resolution ( $\gamma$ ) increases FLOPs quadratically ( $\gamma^2$ )

To double total FLOPs, the three factors must be balanced together.

# SOTA CNN Architectures

---

## ➤ EfficientNet

- The authors of EfficientNet searched for the best scaling factors on a small baseline model.
- They found:

$$\alpha = 1.2, \quad \beta = 1.1, \quad \gamma = 1.15$$

# SOTA CNN Architectures

---

$$\alpha = 1.2, \quad \beta = 1.1, \quad \gamma = 1.15$$

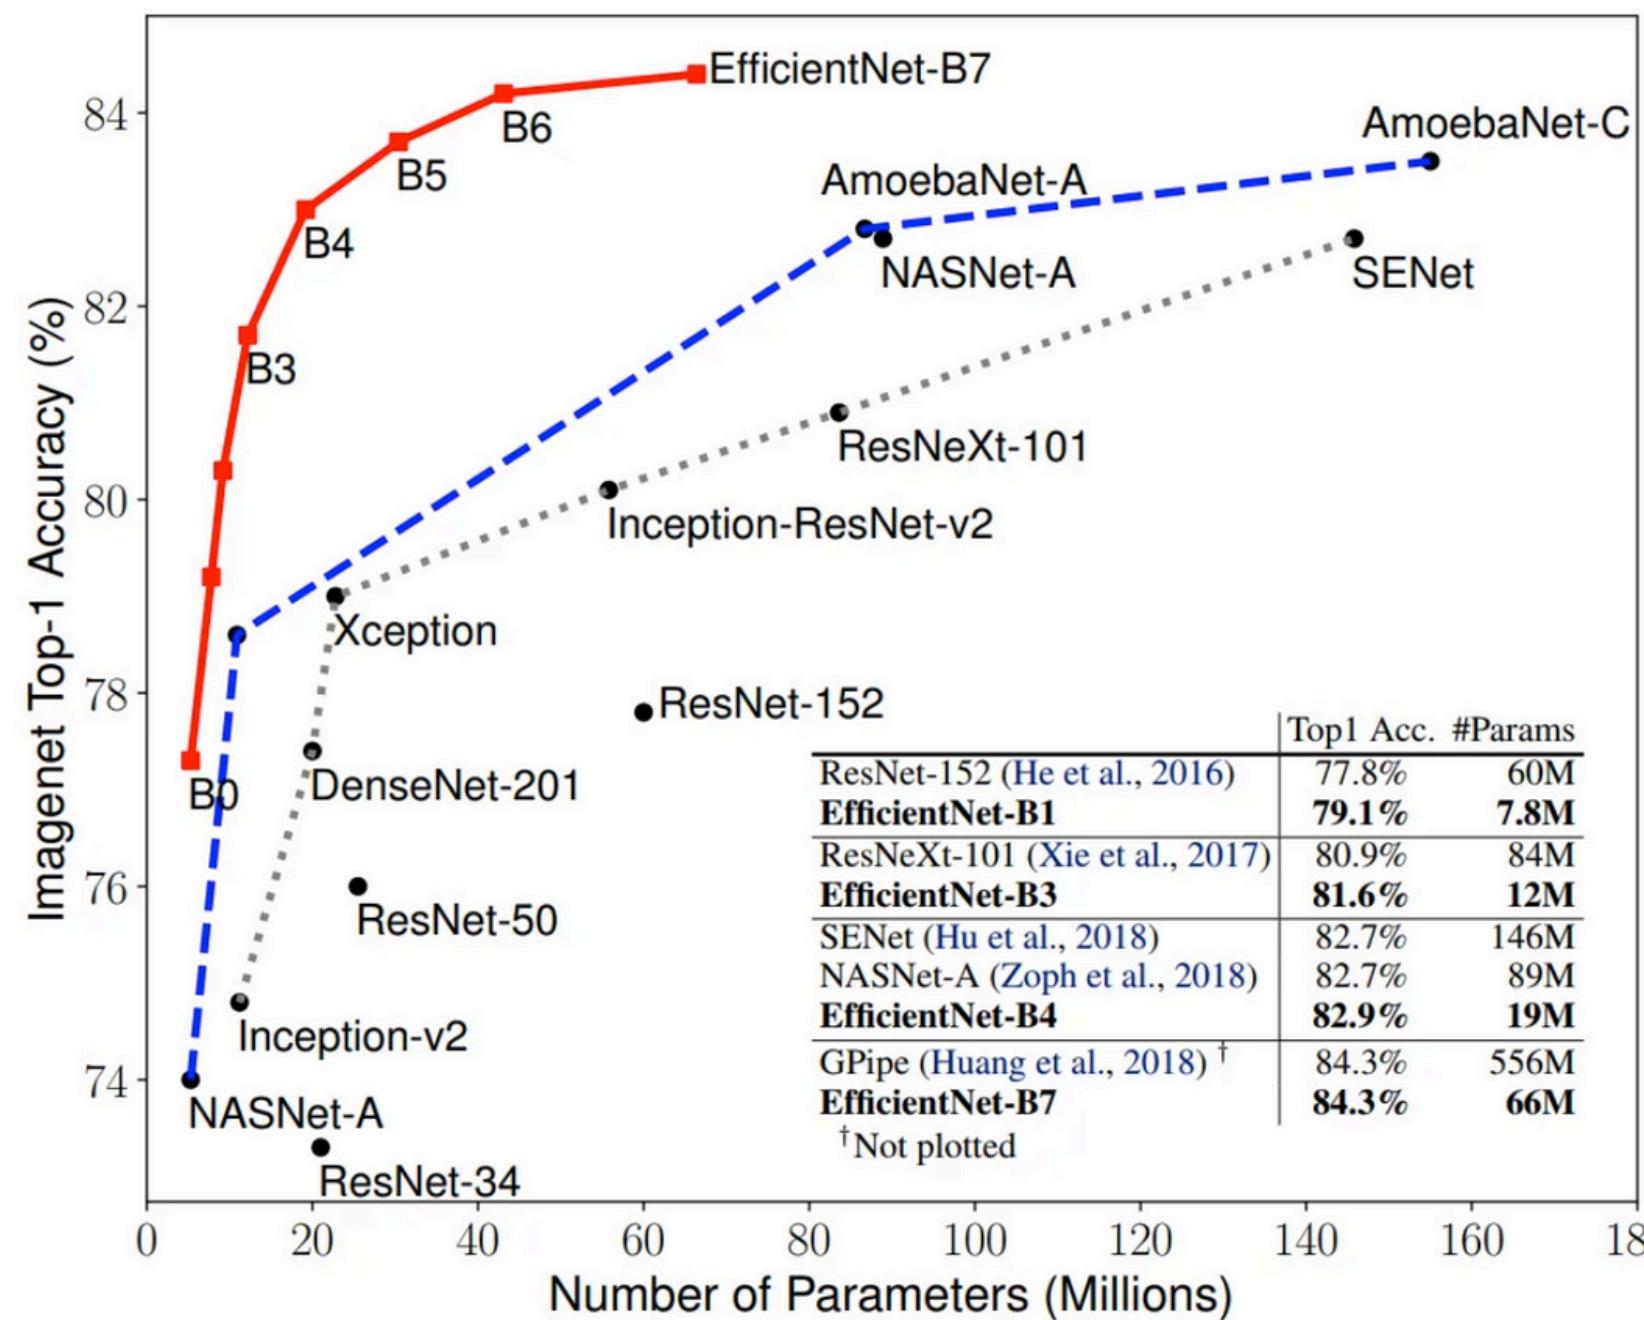
<b>Model</b>	$\phi$	<b>Depth</b> ( $\alpha^\phi$ )	<b>Width</b> ( $\beta^\phi$ )
B0	0	$1.2^0 = 1.0$	$1.1^0 = 1.0$
B1	1	$1.2^1 = 1.2$	$1.1^1 = 1.1$
B2	2	$1.2^2 = 1.44$	$1.1^2 = 1.21$
B3	3	$1.2^3 = 1.73$	$1.1^3 = 1.33$
B4	4	$1.2^4 = 2.07$	$1.1^4 = 1.46$
B5	5	$1.2^5 = 2.49$	$1.1^5 = 1.61$
B6	6	$1.2^6 = 2.99$	$1.1^6 = 1.77$
B7	7	$1.2^7 = 3.58$	$1.1^7 = 1.94$

Table 1: Scaling EfficientNet from B0 to B7

# SOTA CNN Architectures

## ➤ EfficientNet

- Achieve state-of-the-art accuracy with significantly fewer parameters and FLOPs.

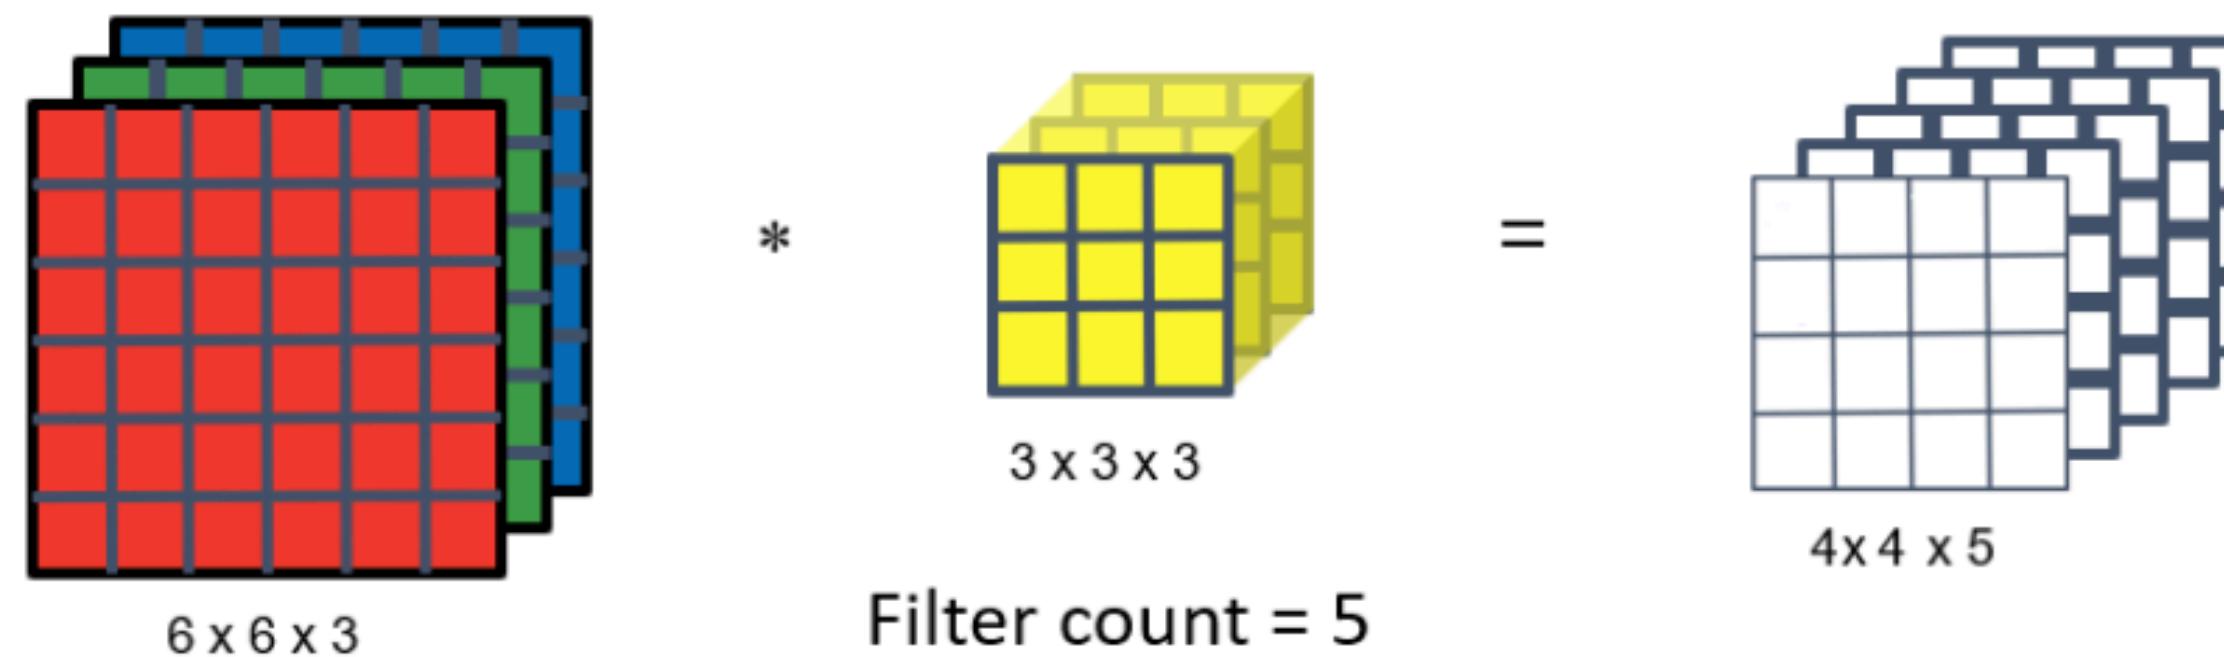


# SOTA CNN Architectures

---

## ➤ MobileNets

- Small-sized models are crucial for mobile and embedded devices.
- MobileNets reduce computational cost and memory usage while maintaining good accuracy.
- Use depthwise-separable convolutions to significantly reduce computation compared to standard convolutions.

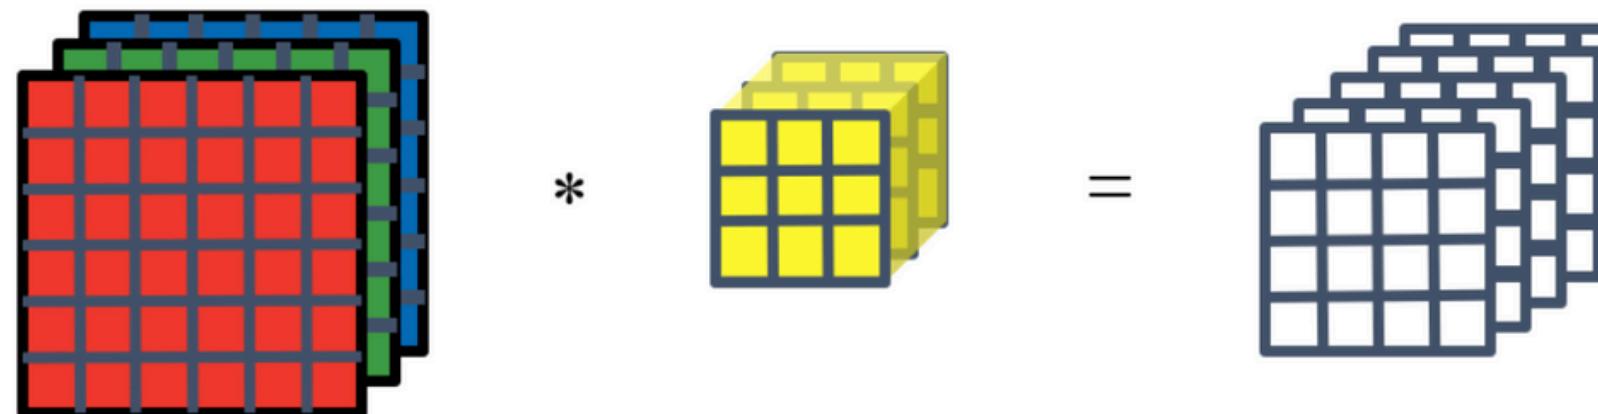


# SOTA CNN Architectures

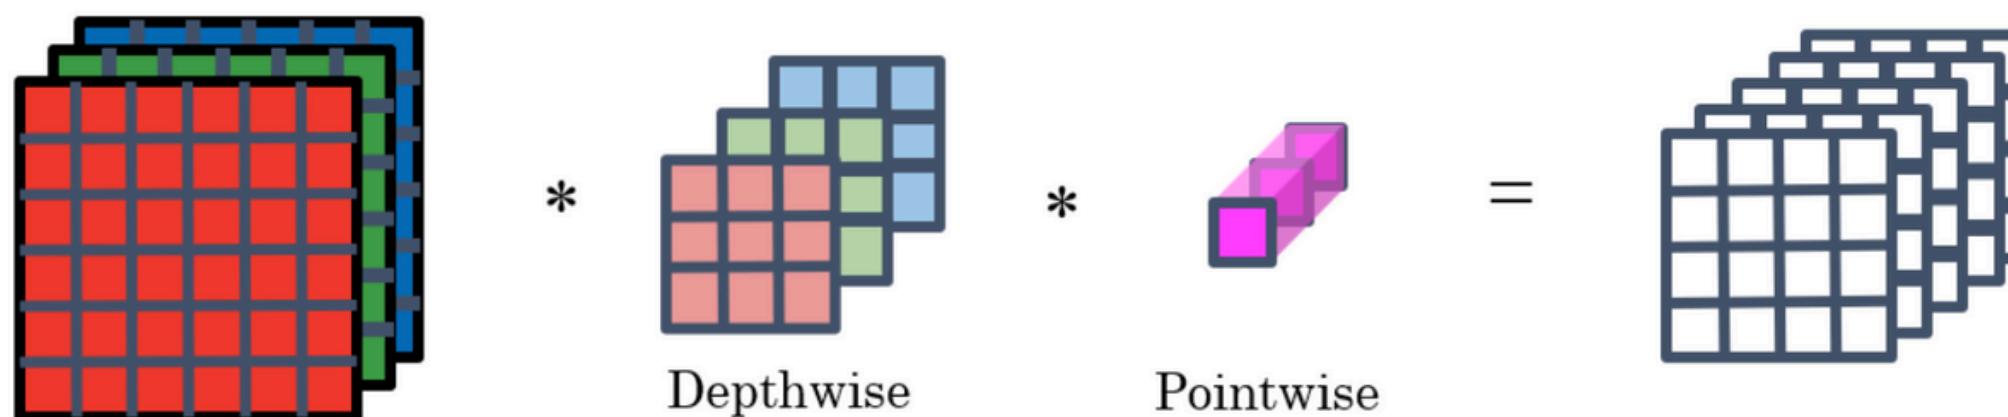
## > MobileNets

- Computational cost of standard convolution:
- Cost = # filter params  $\times$  # filter positions  $\times$  # filters
- Filters operate on all input channels, increasing computation significantly.

Normal Convolution



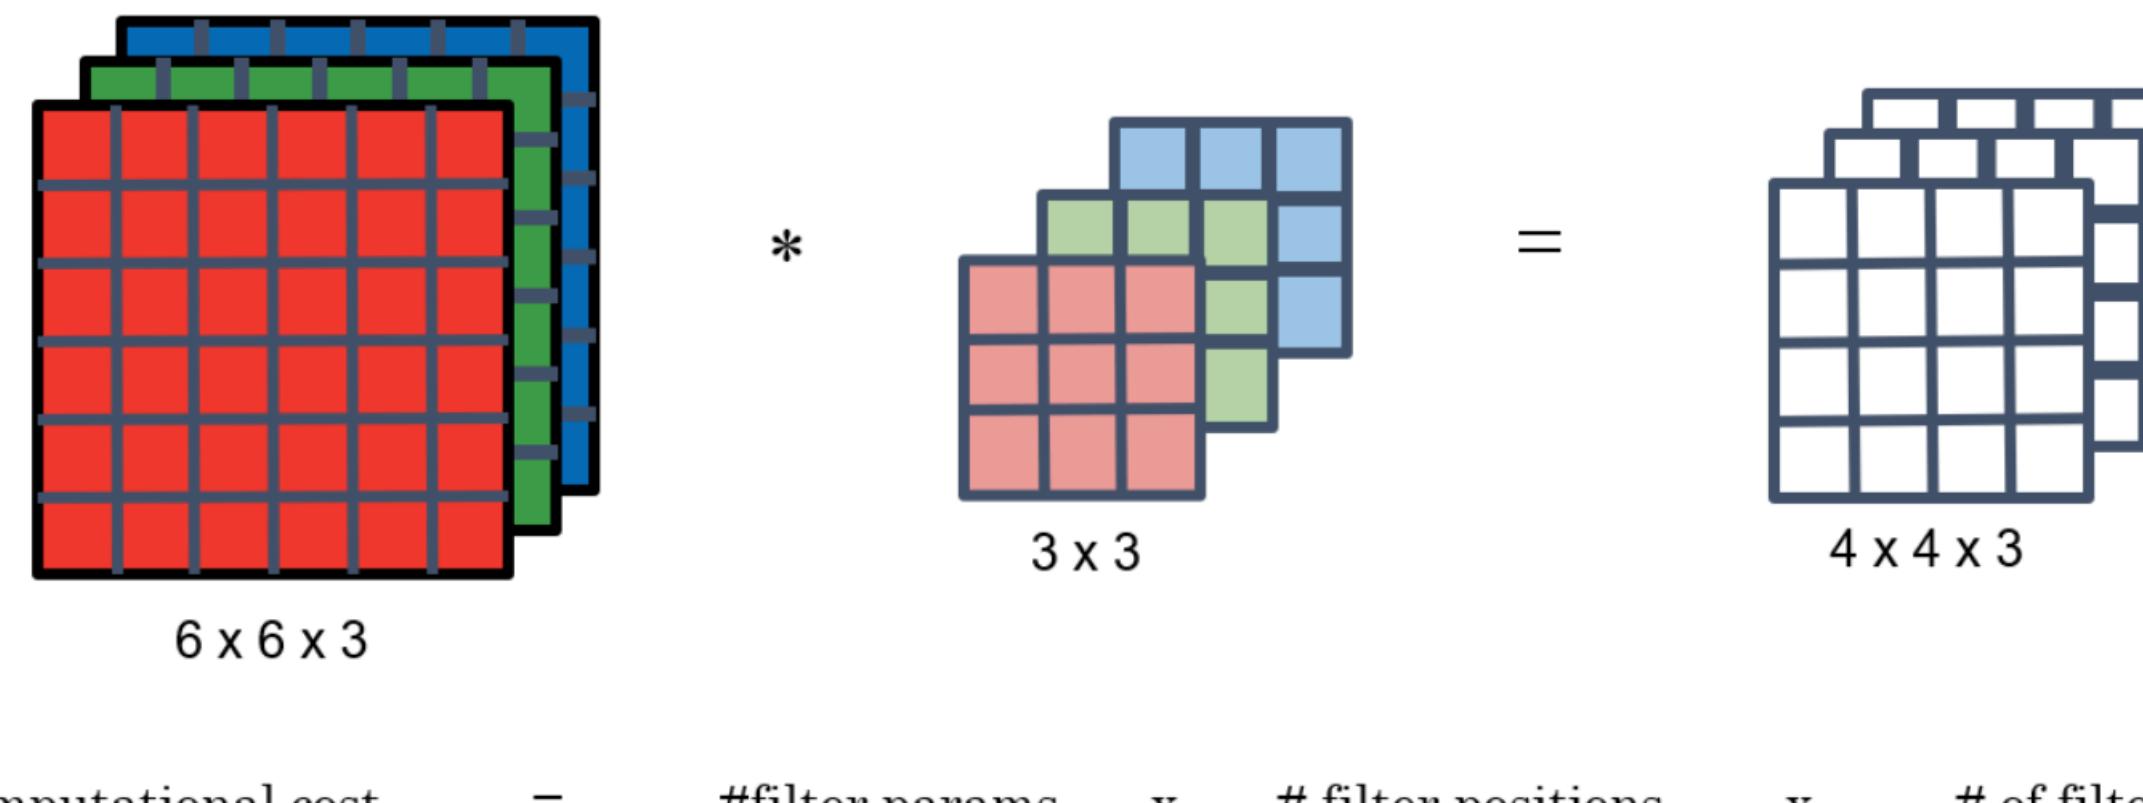
Depthwise Separable Convolution



# SOTA CNN Architectures

## > MobileNets

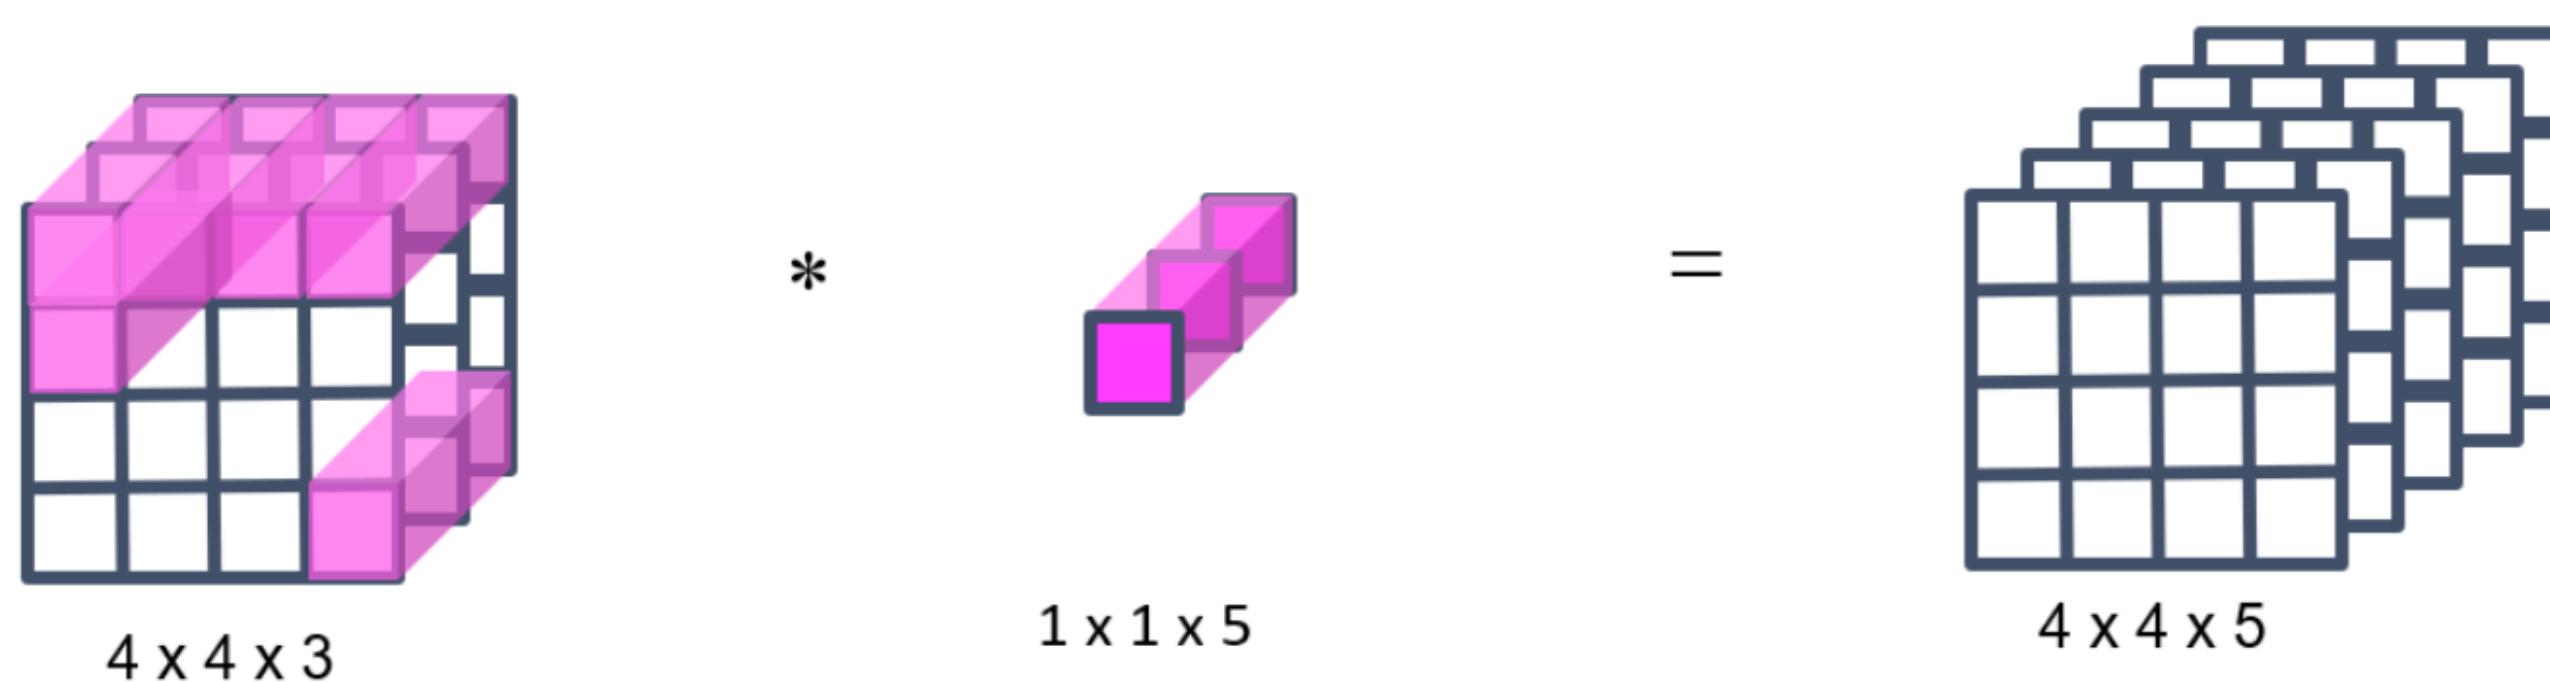
- Depthwise-Separable Convolutions
- Split standard convolution into two steps:
- Depthwise Convolution: Applies a single filter per input channel.
- Pointwise Convolution: Combines outputs from depthwise convolution.
- Key Benefit: Reduces computational cost significantly compared to standard convolution.



# SOTA CNN Architectures

## ➤ MobileNets

- Combines outputs from depthwise convolution using  $1 \times 1$  convolutions (mixes channels).

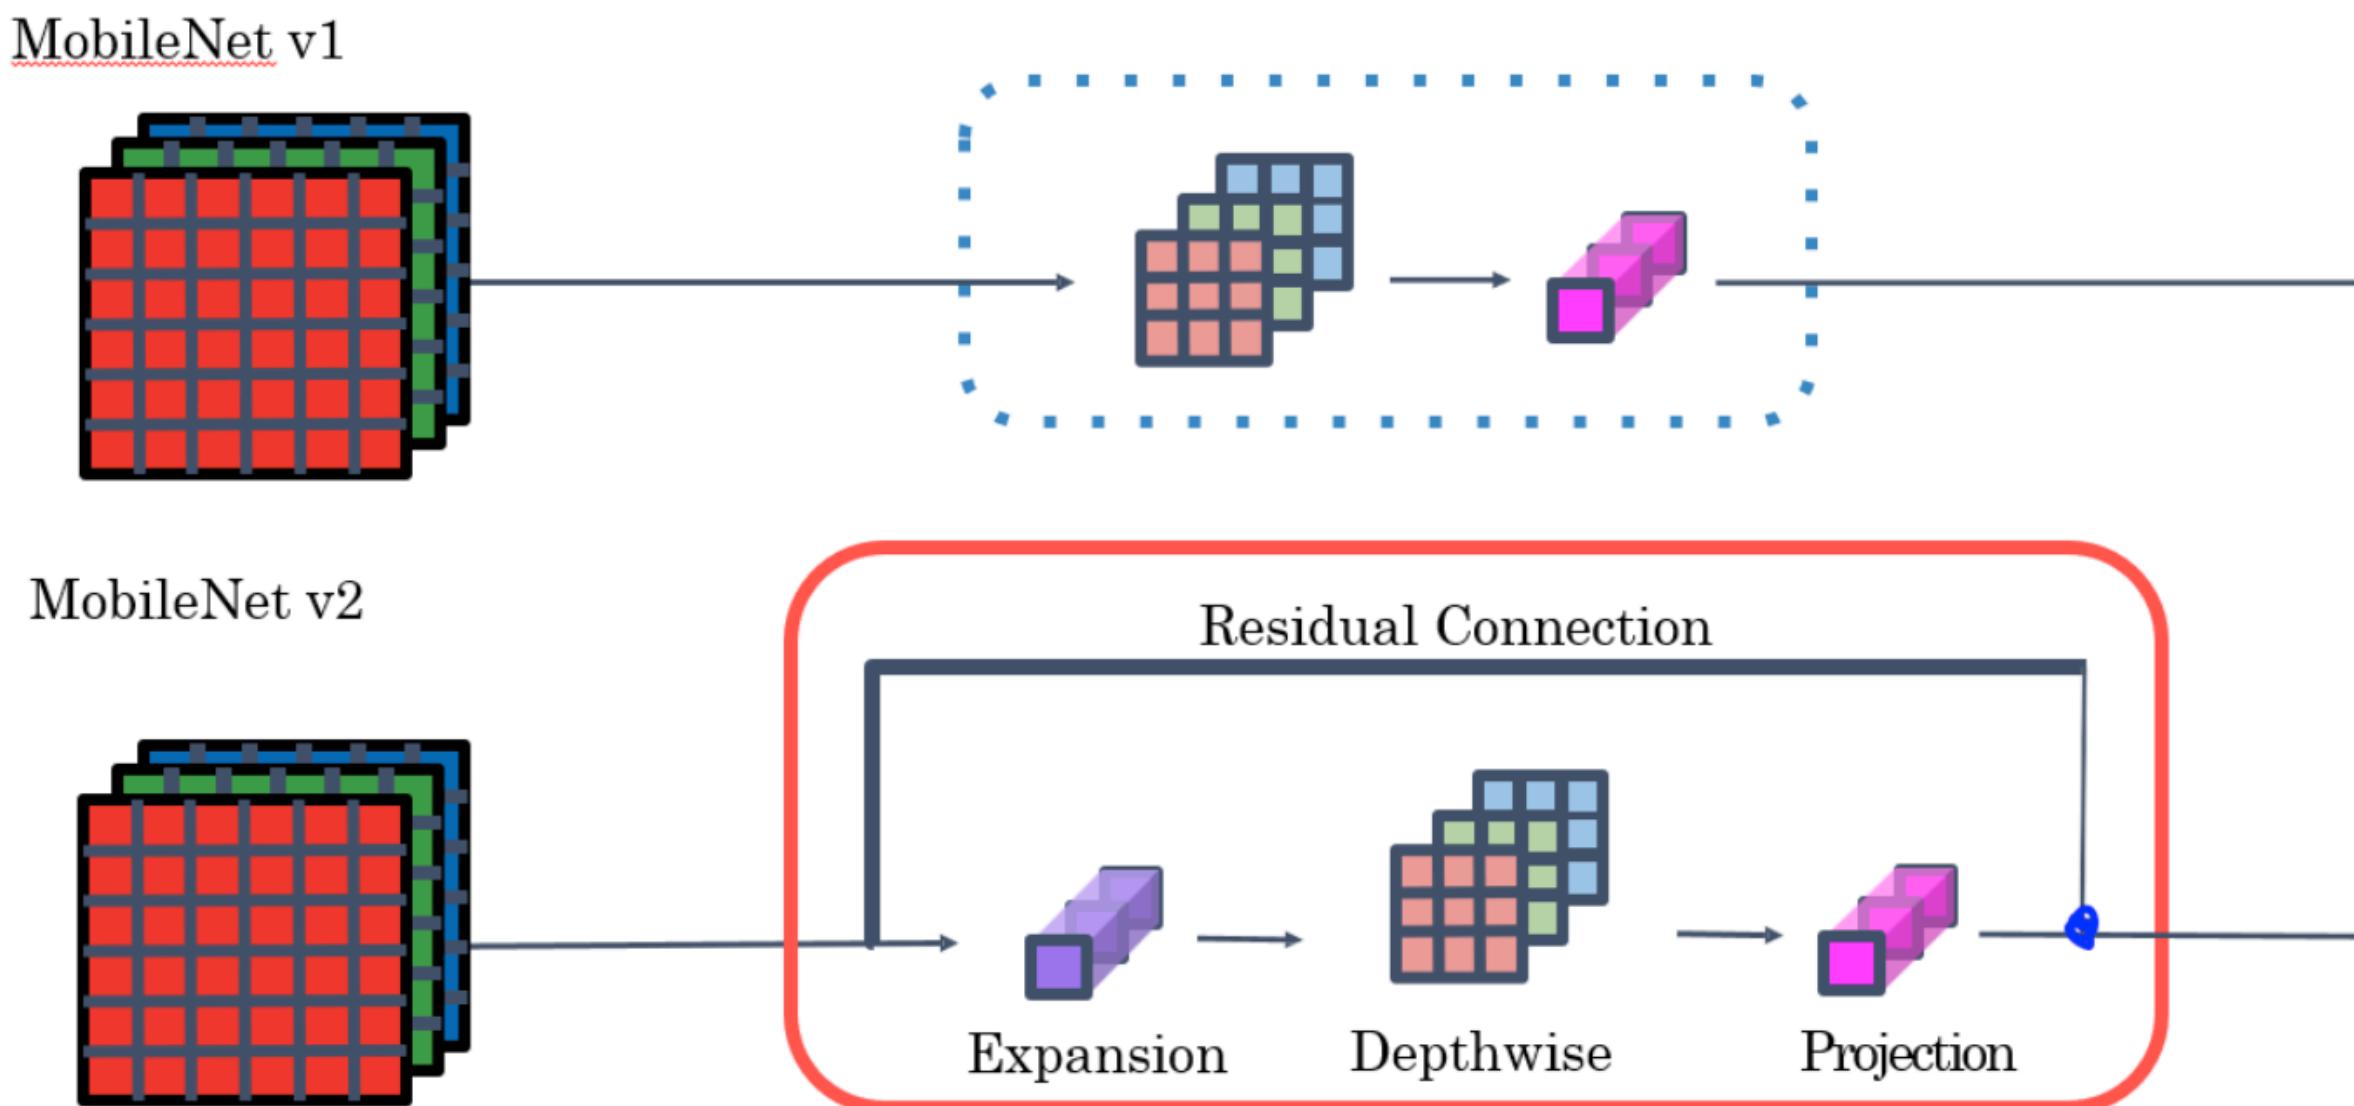


Computational cost = #filter params  $\times$  # filter positions  $\times$  # of filters

# SOTA CNN Architectures

## ➤ MobileNet v2:

- Adds residual connections.
- Expansion step: Expands input dimensions before depthwise convolution.
- Projection step: Reduces dimensions after processing.



A graphic icon consisting of two white speech bubbles with dark green outlines. The left bubble contains the letter 'Q' and the right bubble contains the letter 'A', representing a question and answer pair.

Q A

Thank You